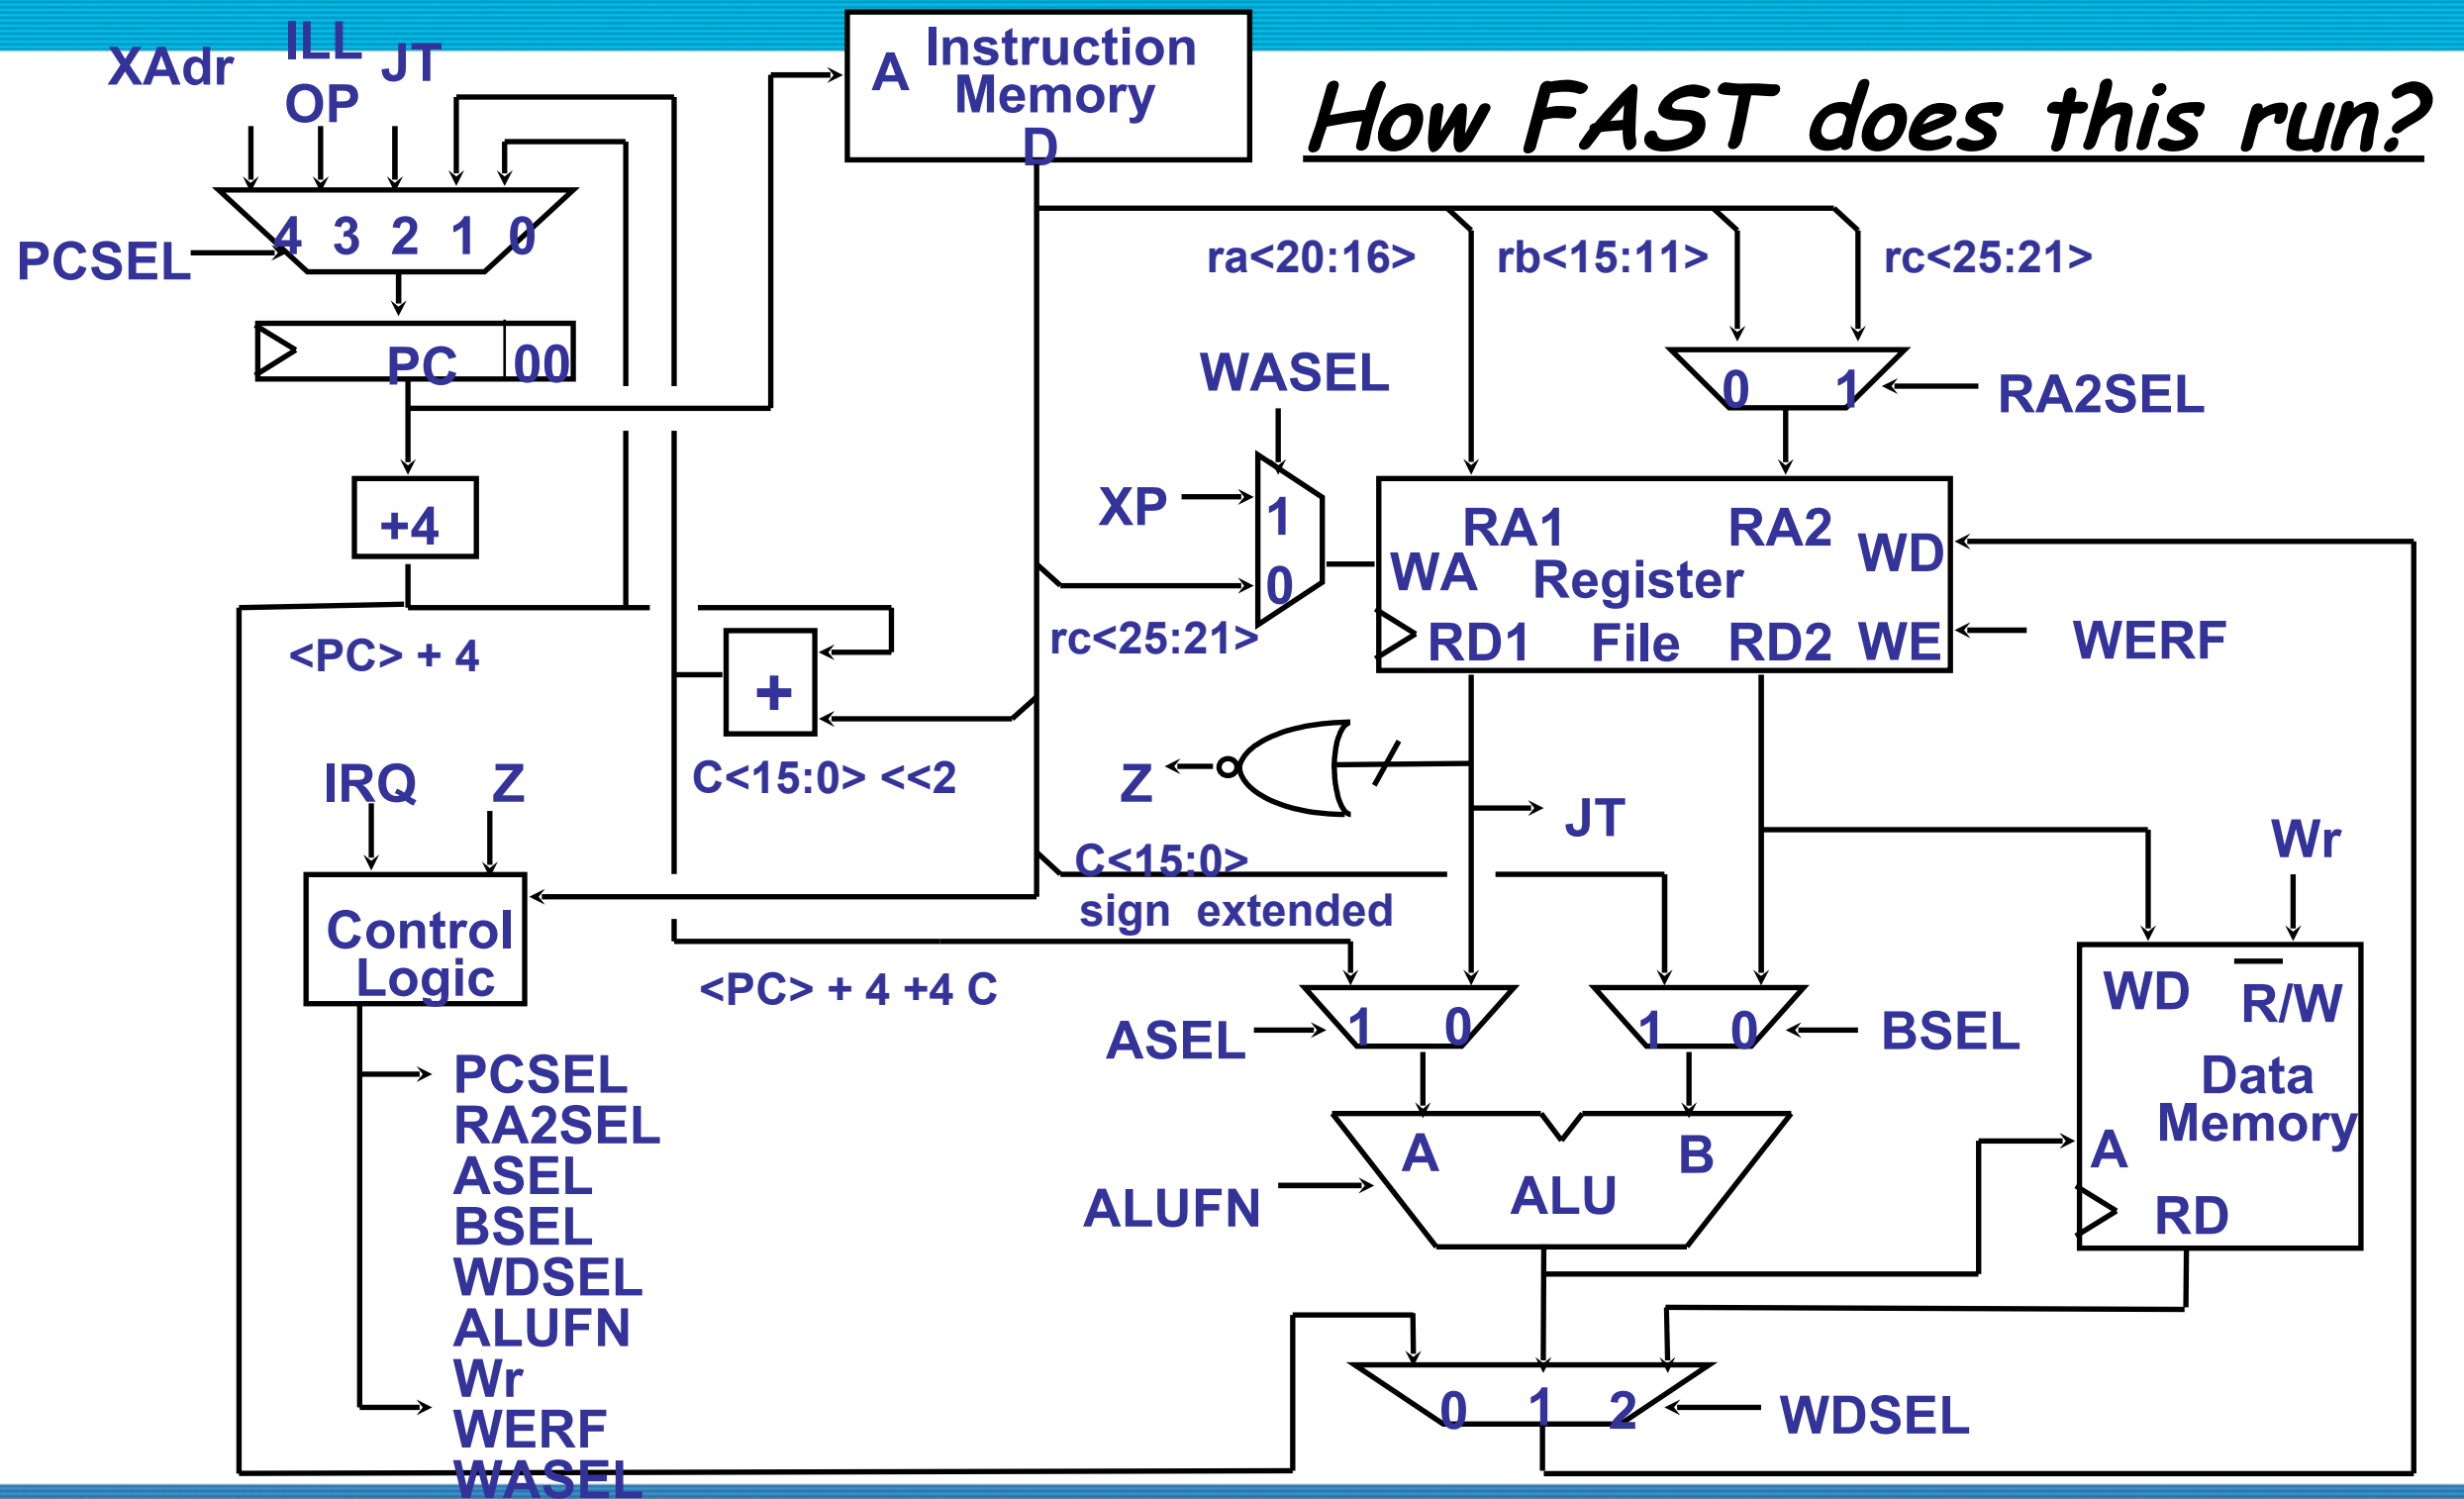



CPU Performance



0010101001010100011110100001100
10001100100001111001101010010101
11001010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010100101010010101



DISCS

CPU Performance

- ▶ **MIPS: Millions of Instructions Per Second**

- ▶ **Common measure of performance:**

$$\text{MIPS} = \frac{\text{Frequency in MHz}}{\text{Cycles Per Instruction}}$$


- ▶ **Clocks Per Instruction (CPI)**

- ▶ **Measure of *throughput*.**

- ▶ **CPI = Number of clock cycles /
Number of instructions finished in that time**

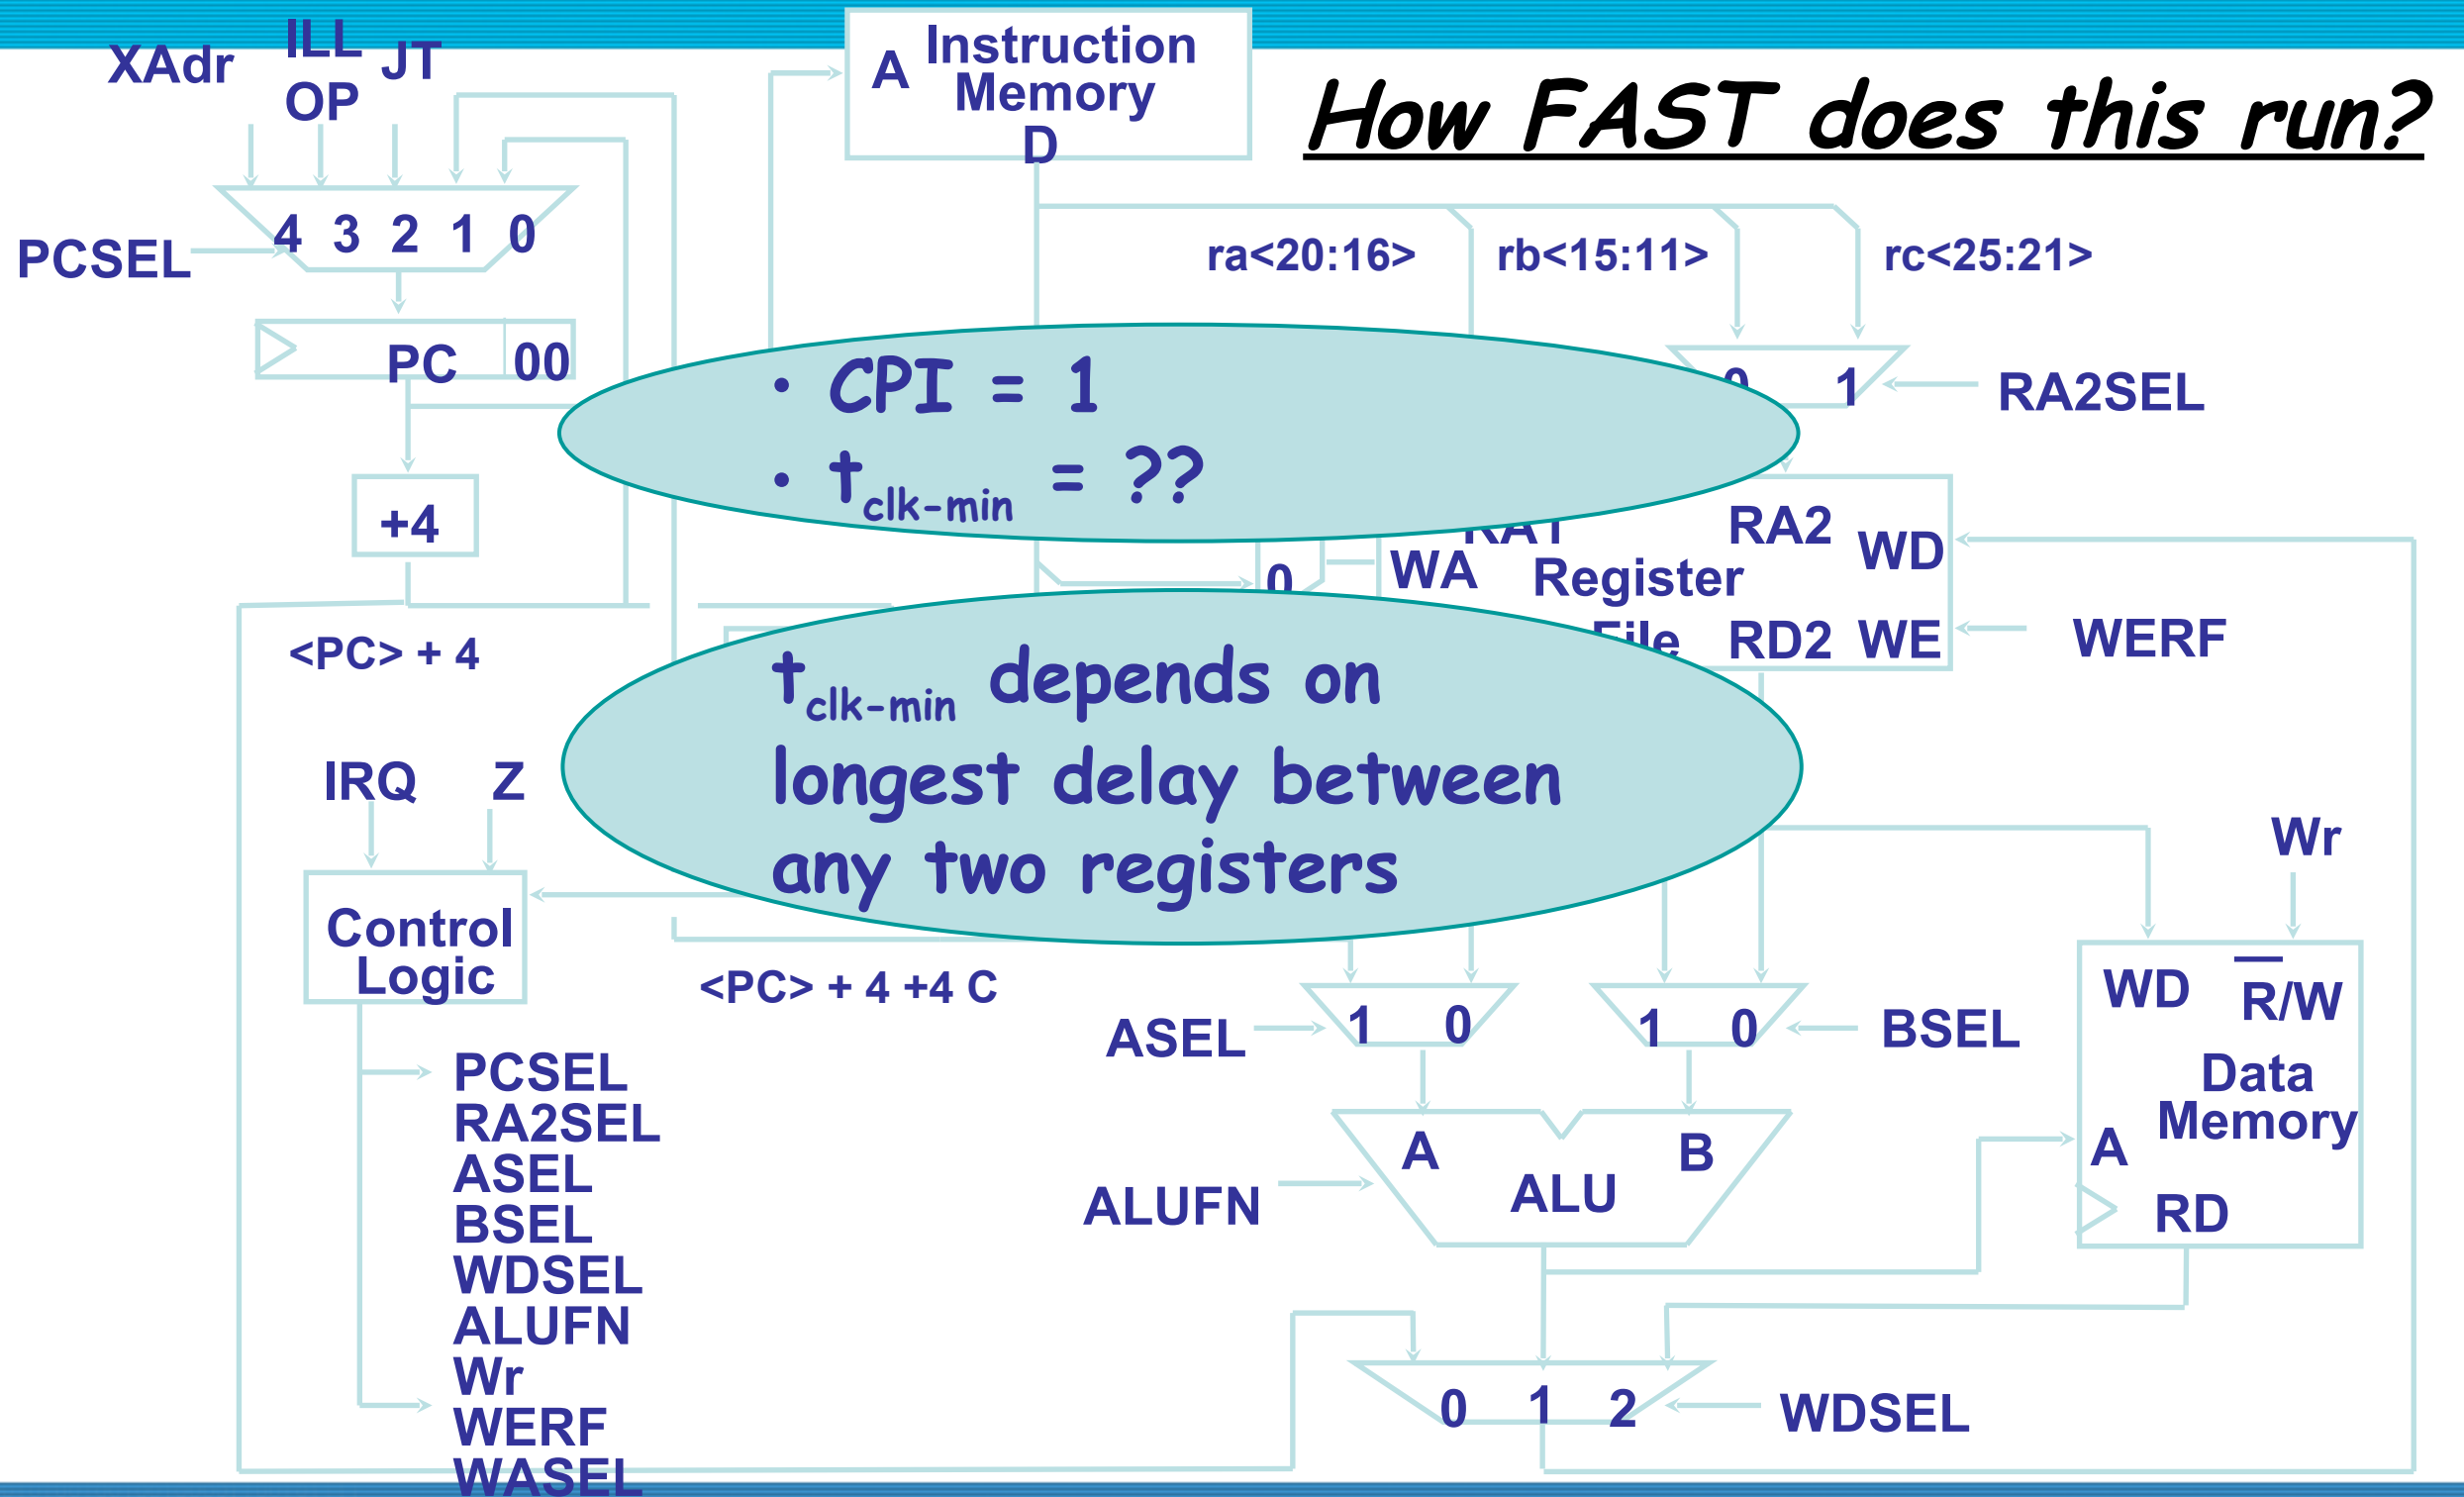
- ▶ **Frequency: $1 / t_{\text{clk-min}}$**

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010101010101010101010101010101




DISCS

CPU Performance



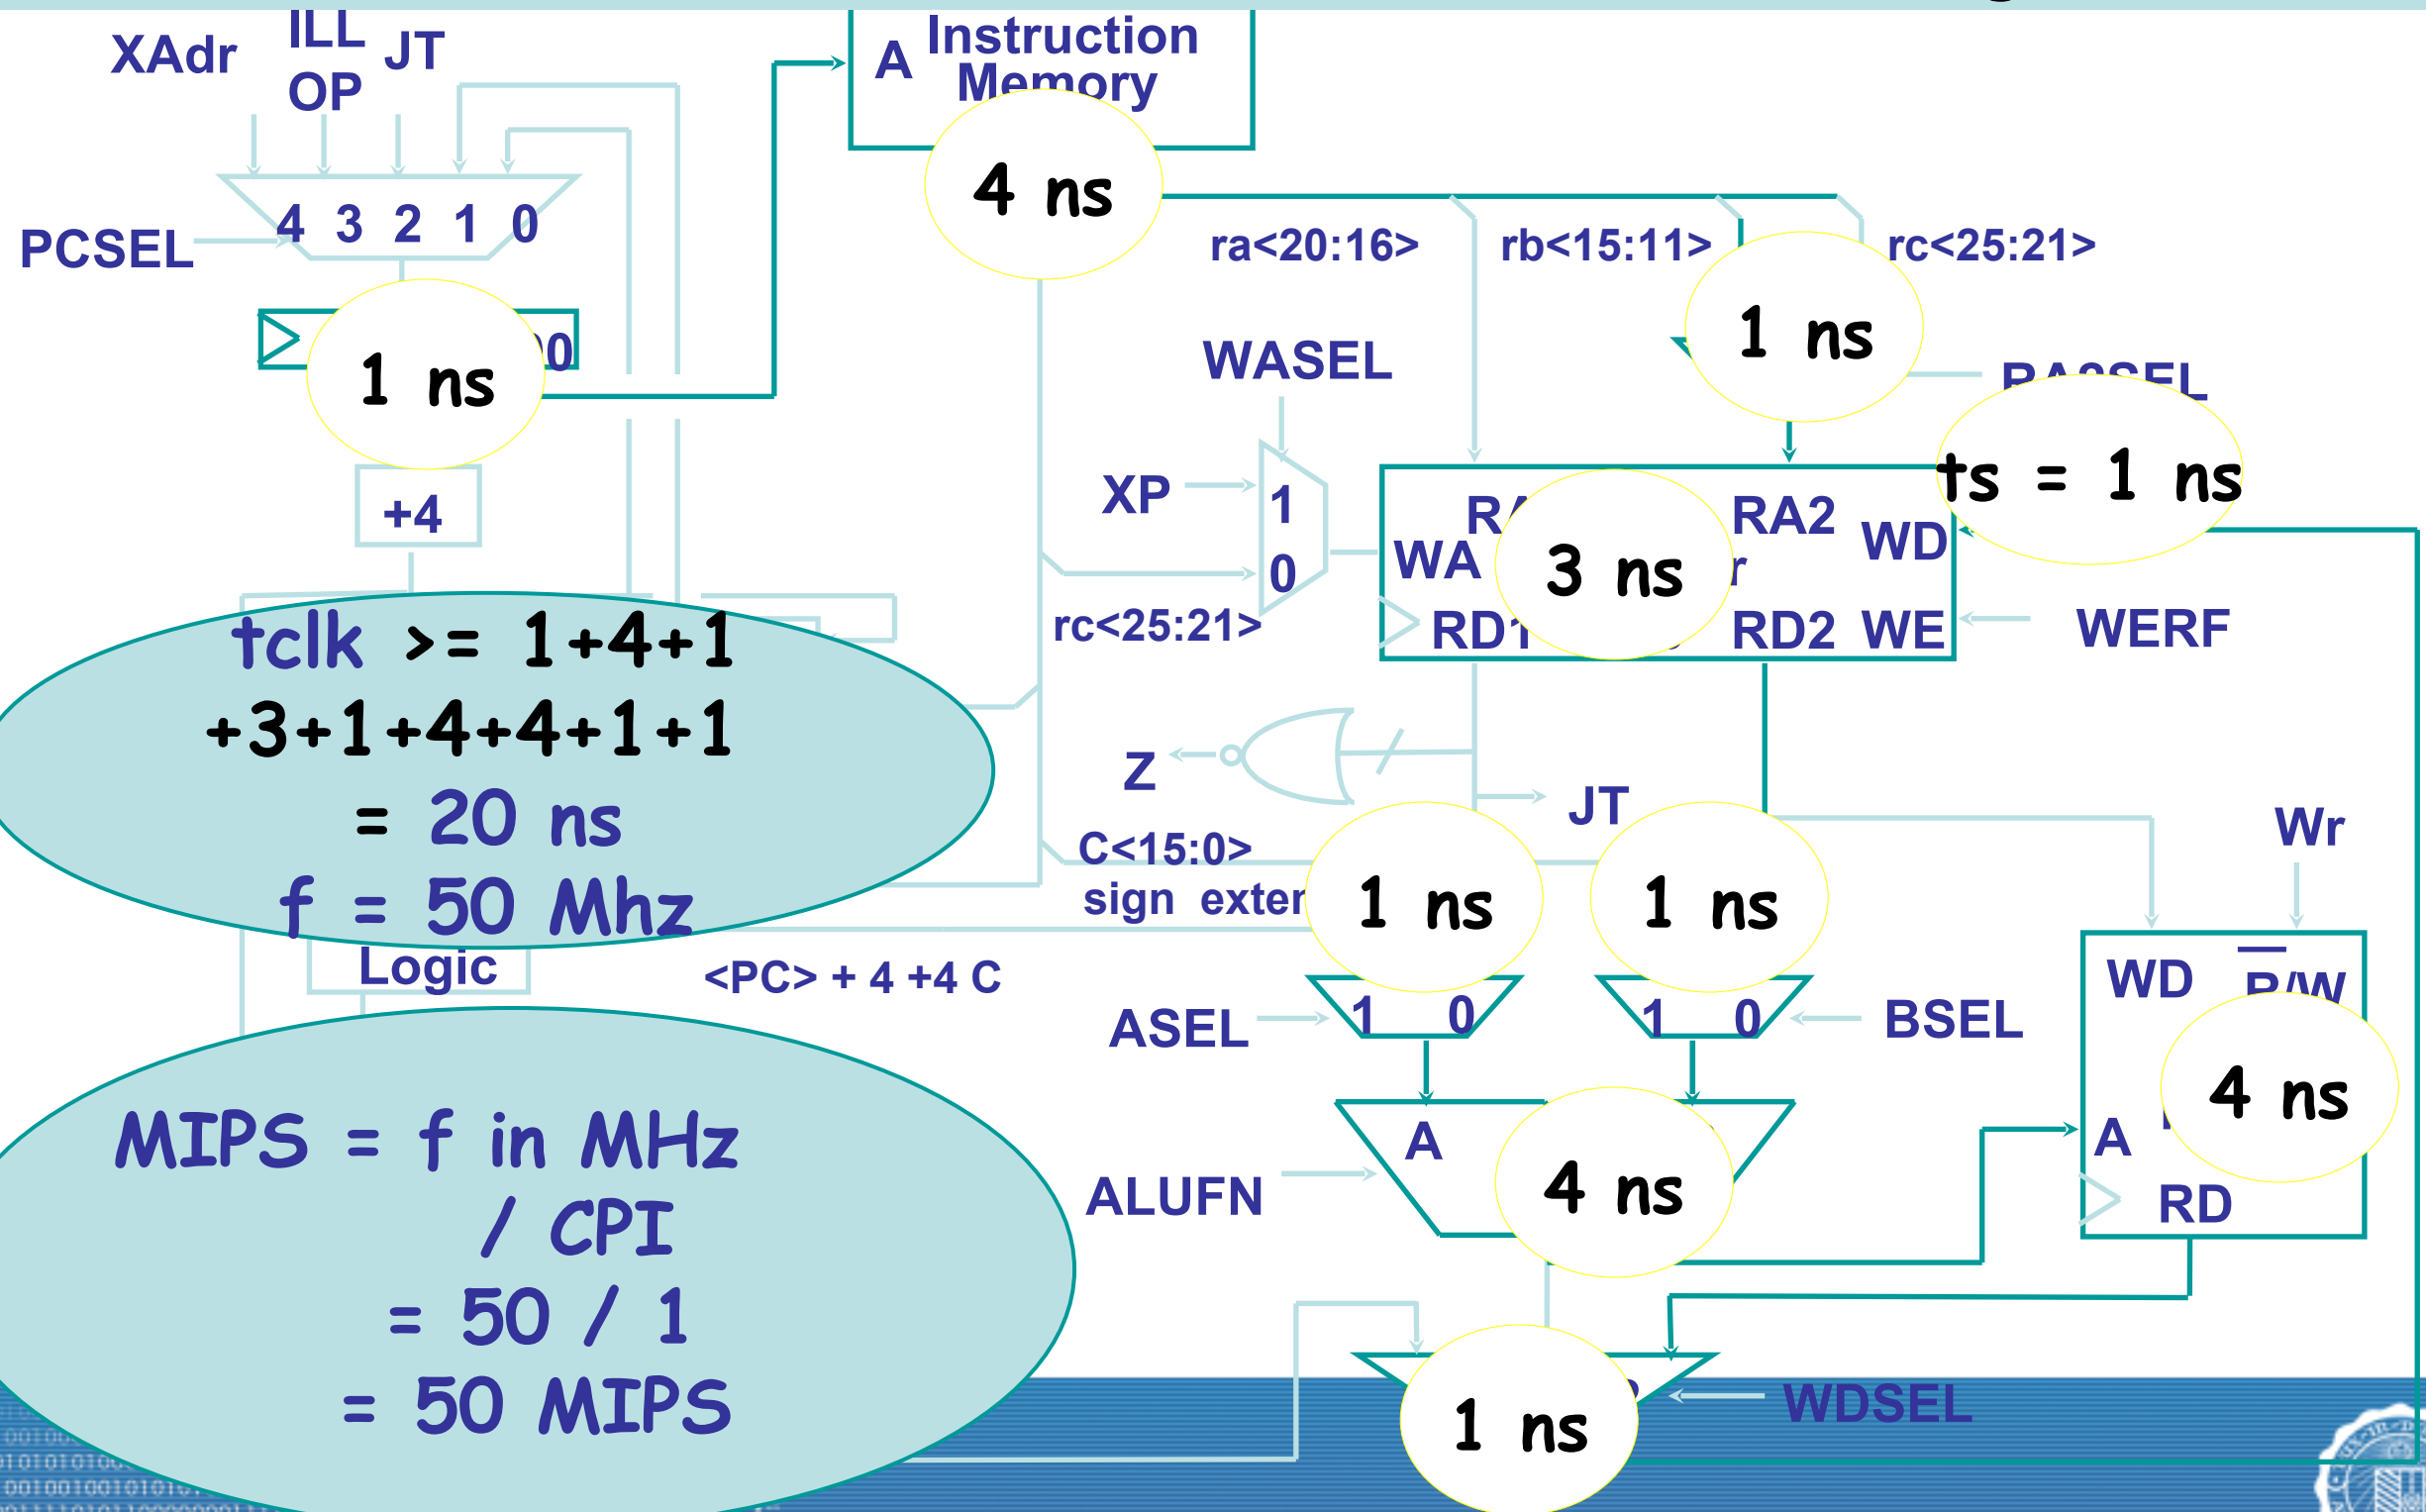
0010101001010100011110100001100
10001100100001111001101010010101
11001010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010101010101010101010101010101



DISCS

CPU Performance

Longest delay here is: PC → Imem → Iword → RA1, RA2 → RegFile → RD1, RD2 → ALU → EA → DMem → Dword → RegFileWD




CPU Performance

$$\text{MIPS} = \frac{\text{Frequency in MHz}}{\text{Clocks Per Instruction (CPI)}}$$

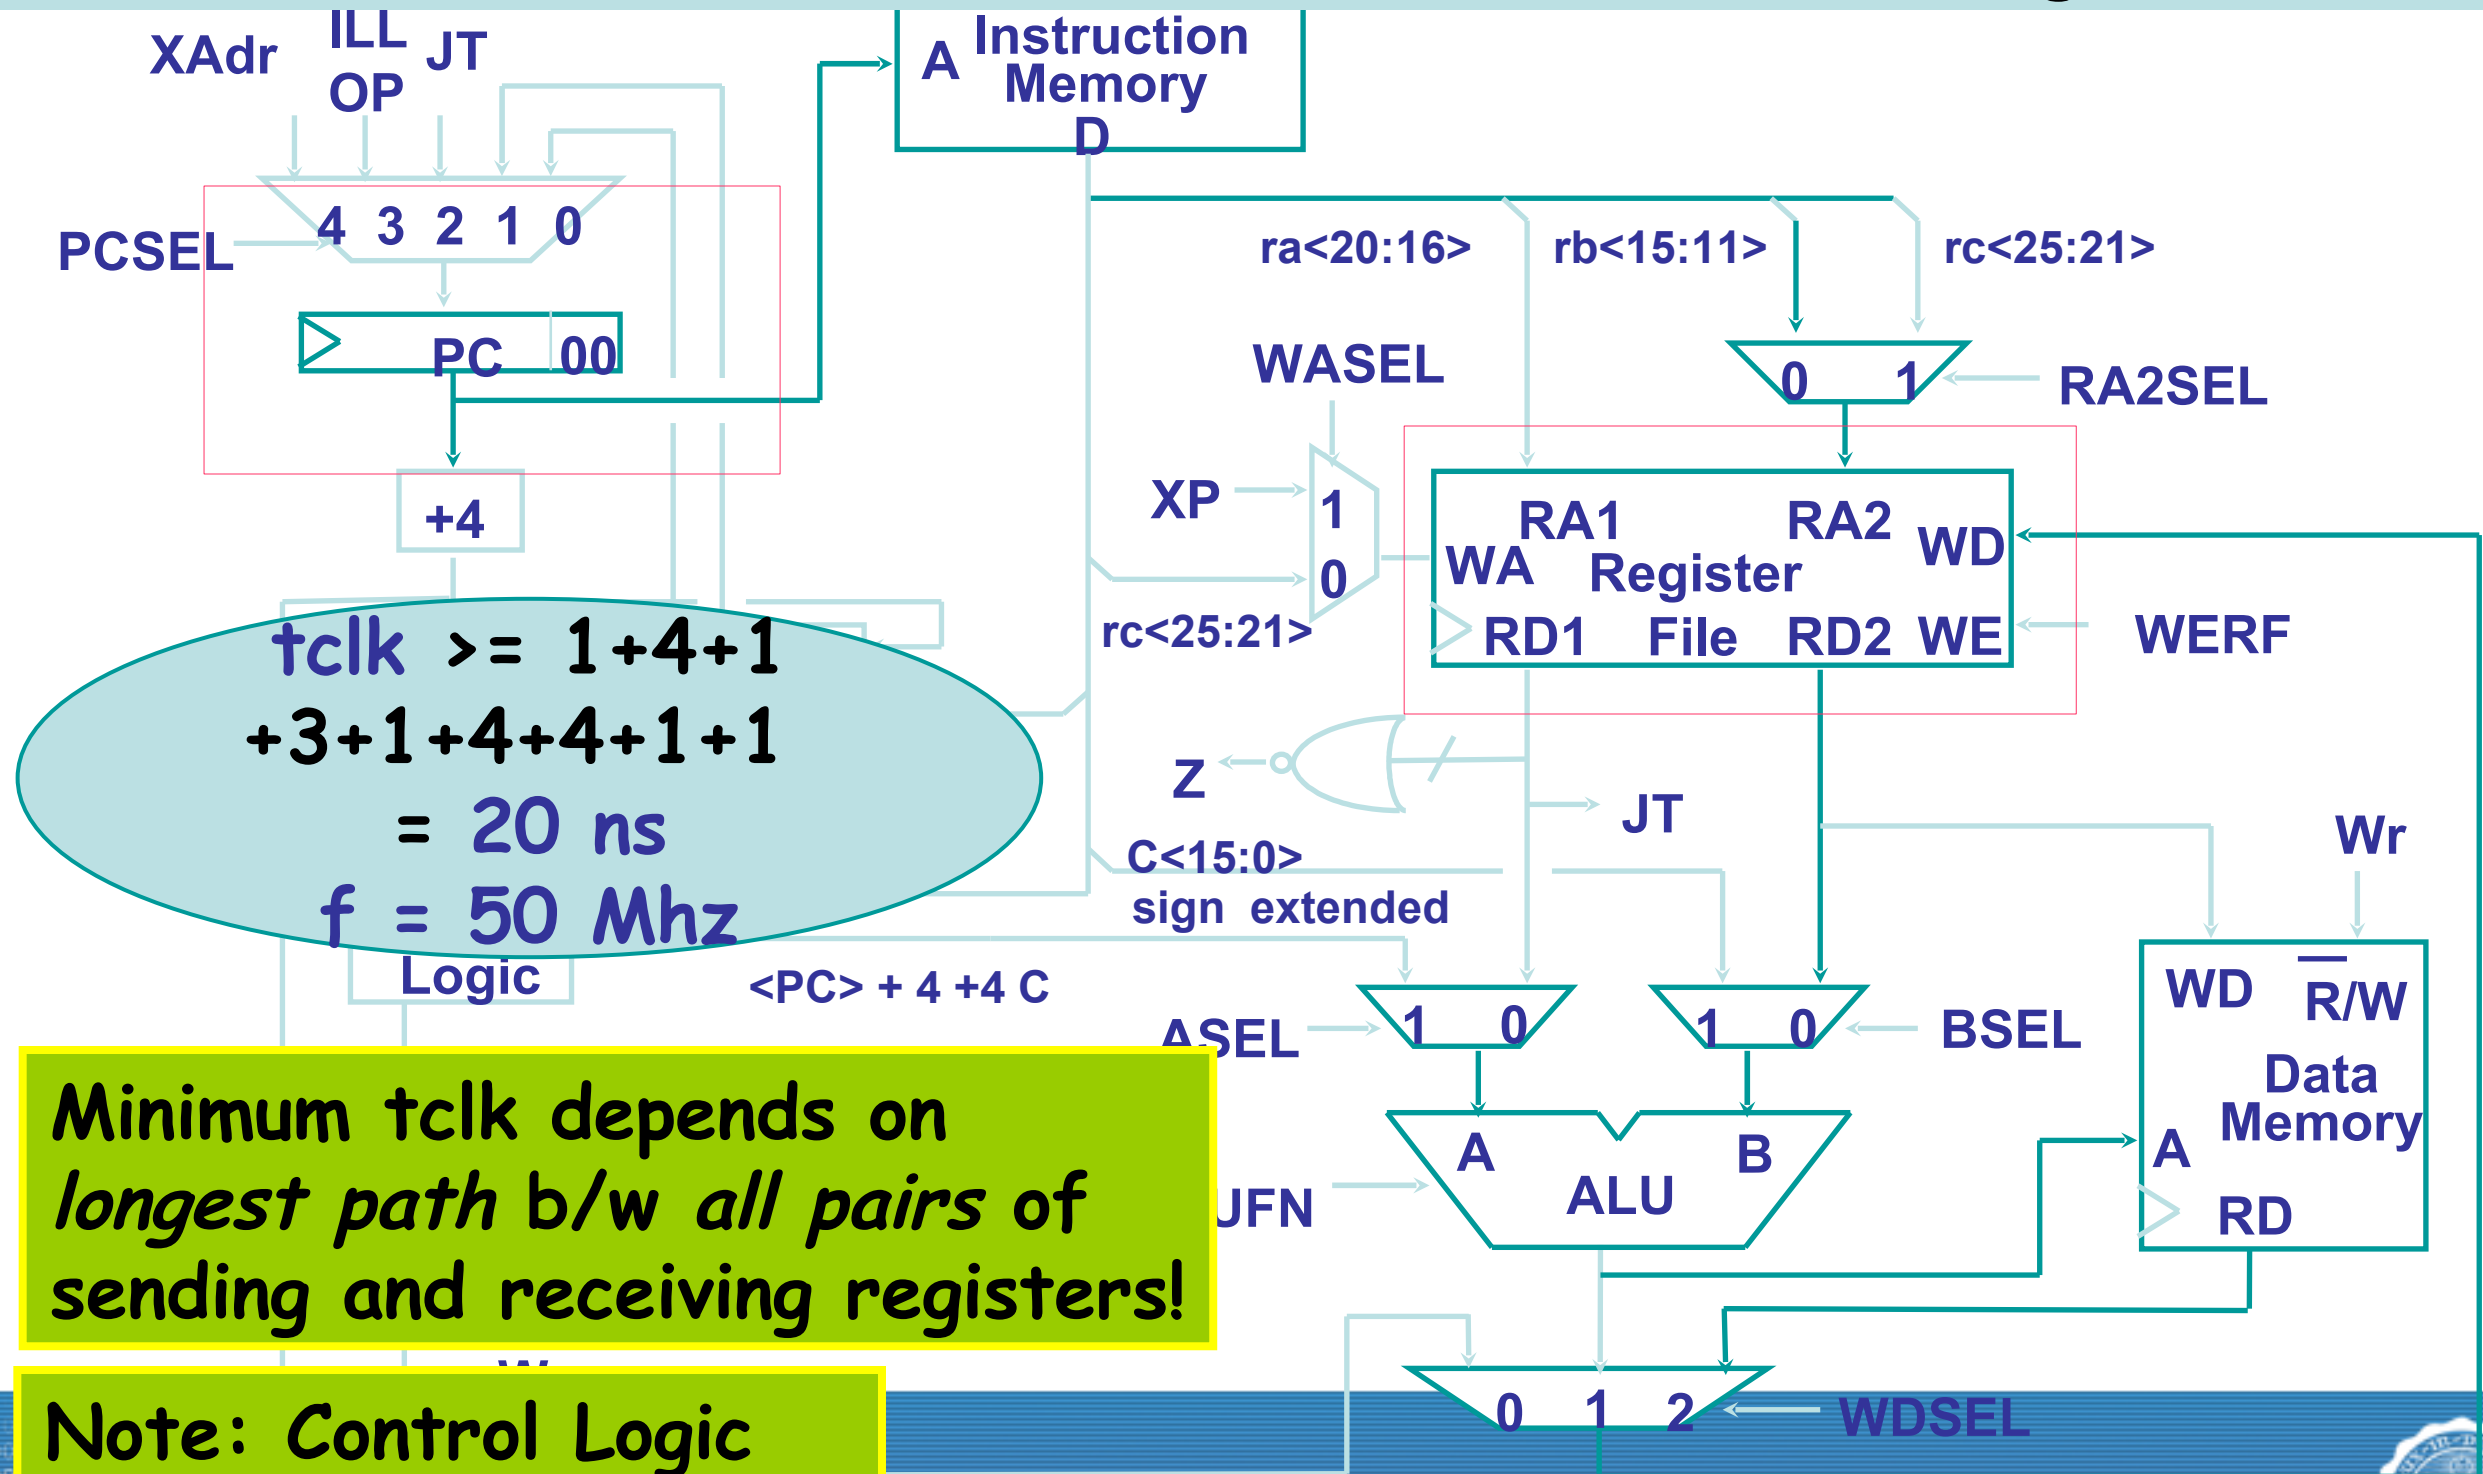
- ▶ How can we increase MIPS?
 - ▶ Decrease CPI
 - ▶ Instruction set simplicity reduces CPI to 1.0!
 - ▶ CISC machines take several cycles per instruction, unlike RISC Beta.
 - ▶ CPI below 1.0 is possible with multiple instruction issue machines.
 - ▶ Increase Frequency
 - ▶ Frequency limited by *longest combinational path* between register outputs to register inputs.
 - ▶ Solution: Pipelining!

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010101010101010101010101010101



CPU Performance

Longest delay here is: PC → Imem → Iword → RA1, RA2 → RegFile → RD1, RD2 → ALU → EA → DMem → Dword → RegFileWD

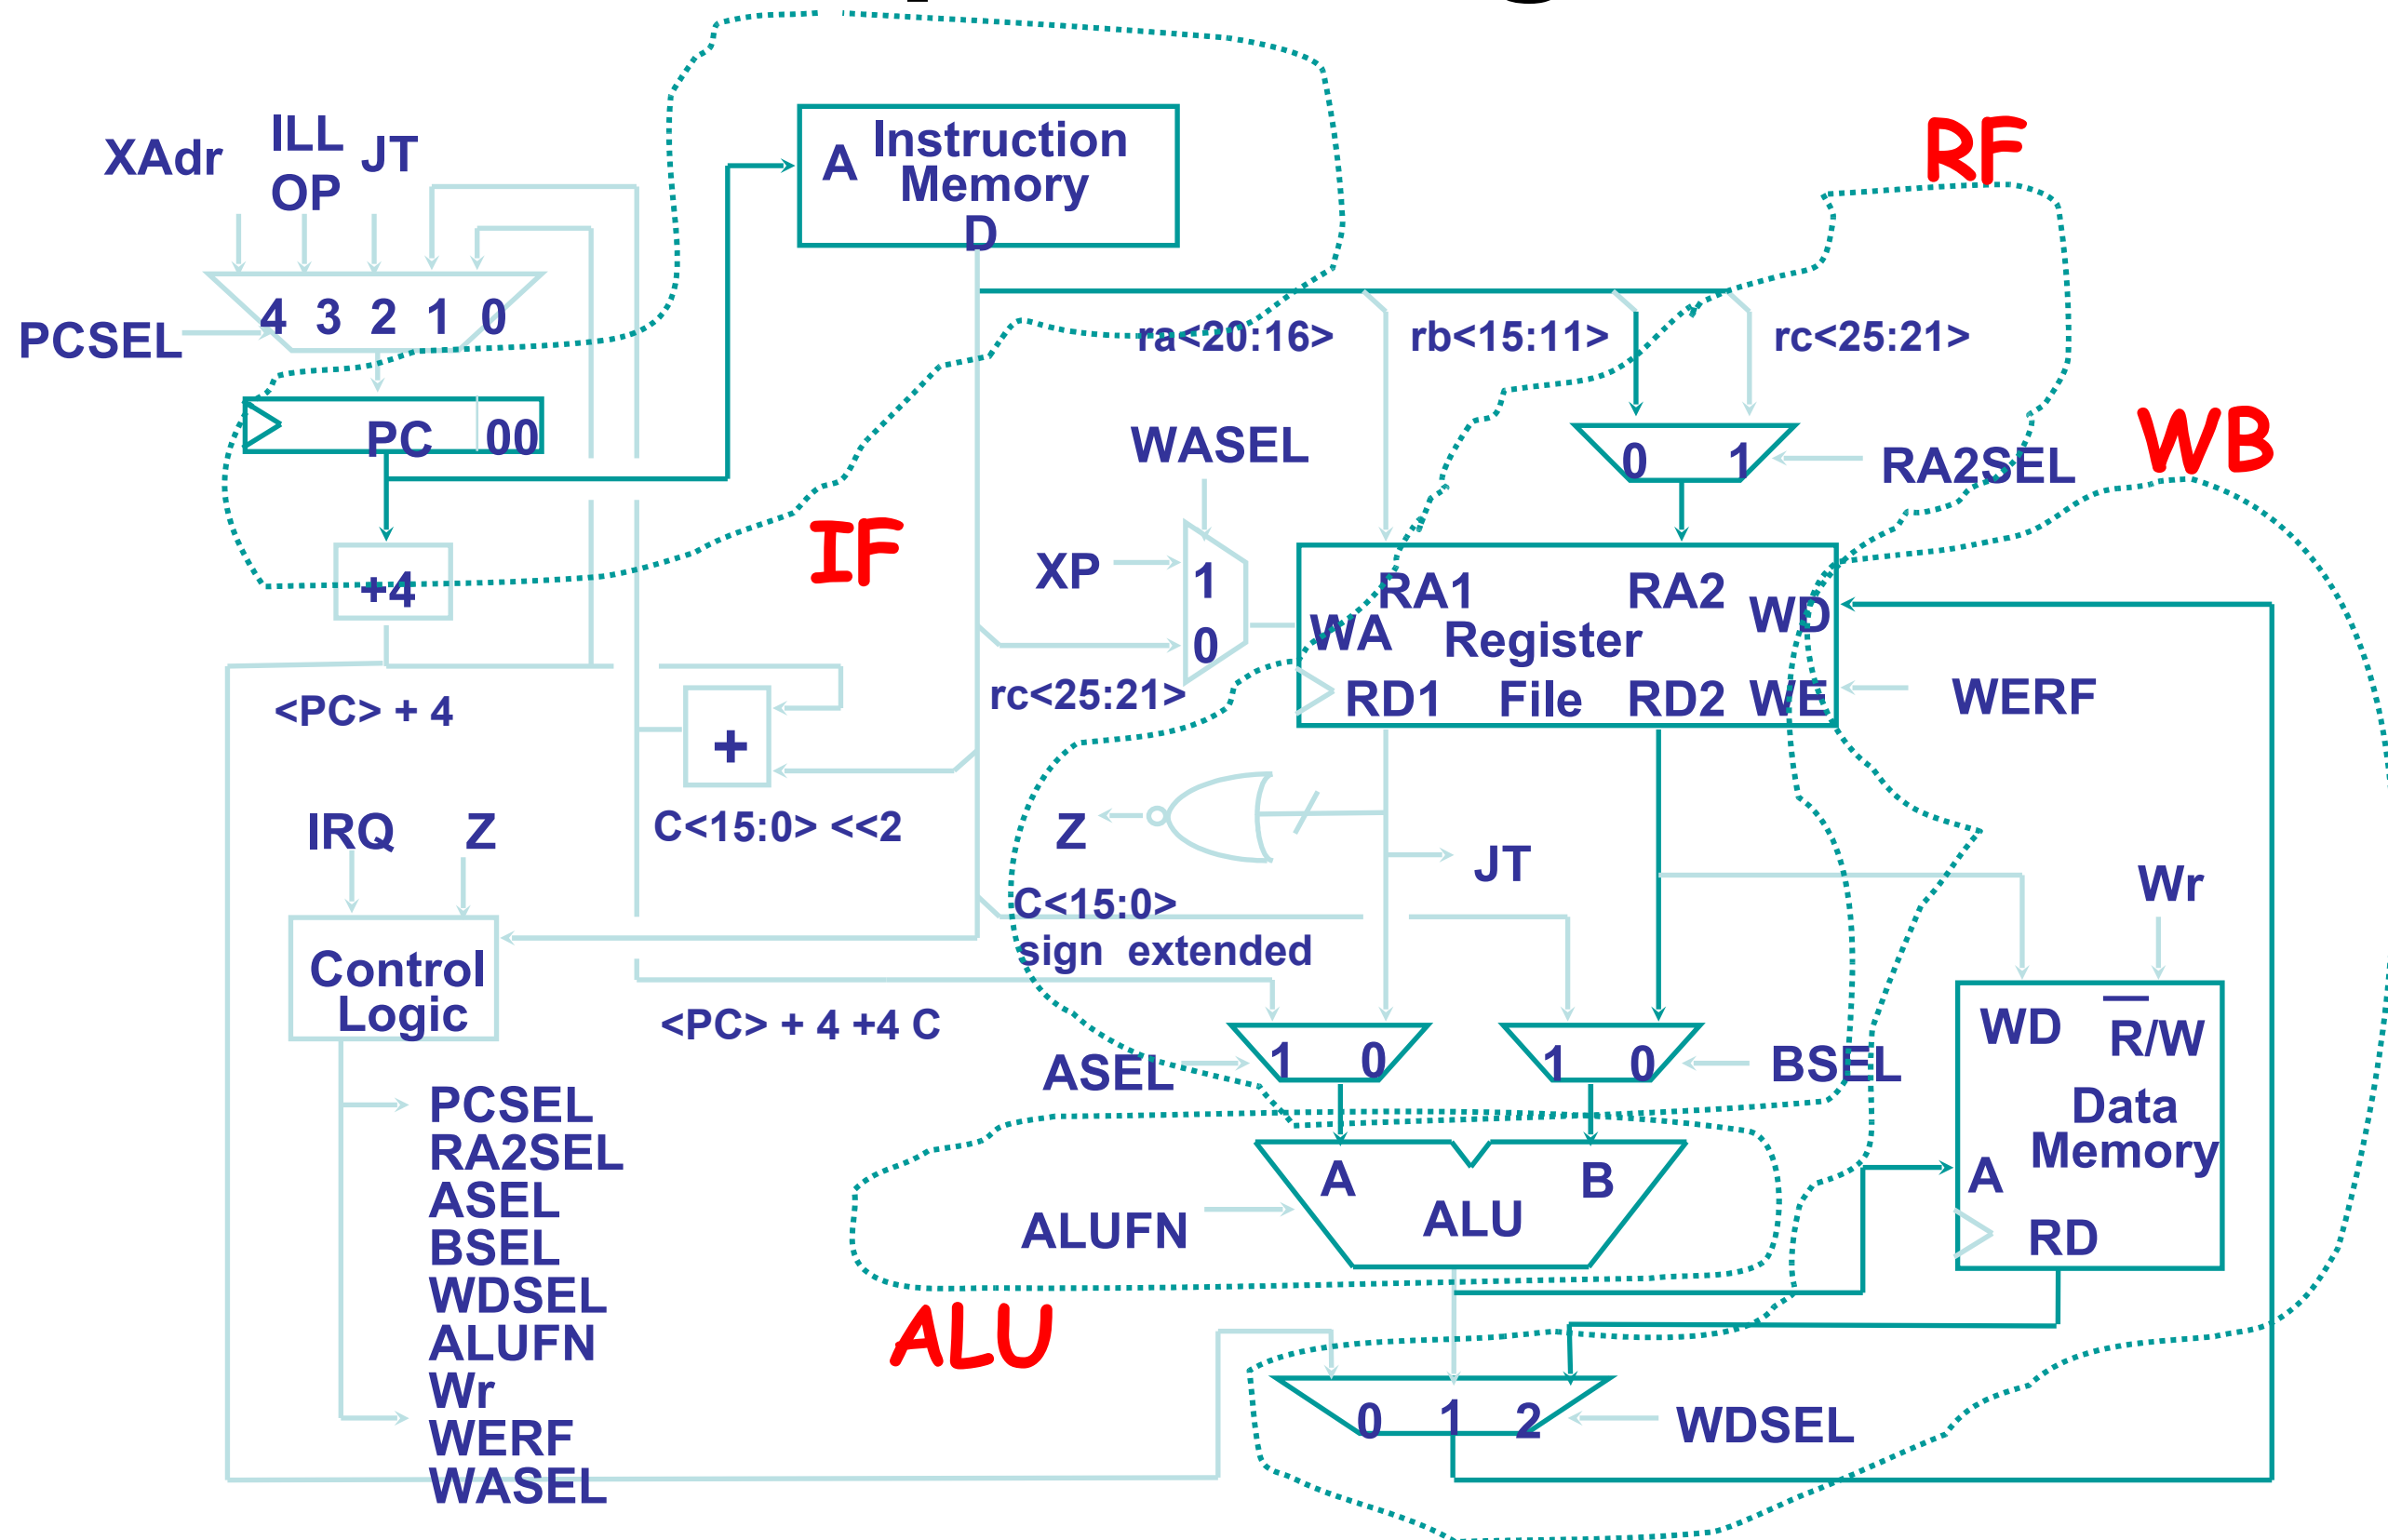


Minimum tclk depends on longest path b/w all pairs of sending and receiving registers!

Note: Control Logic can add delay too!



Datapath Stages



Pipelining the Datapath

- **Goal:** Maintain (nearly) 1.0 CPI, but increase clock speed.

- **Approach:** Structure processor as 4-stage pipeline:

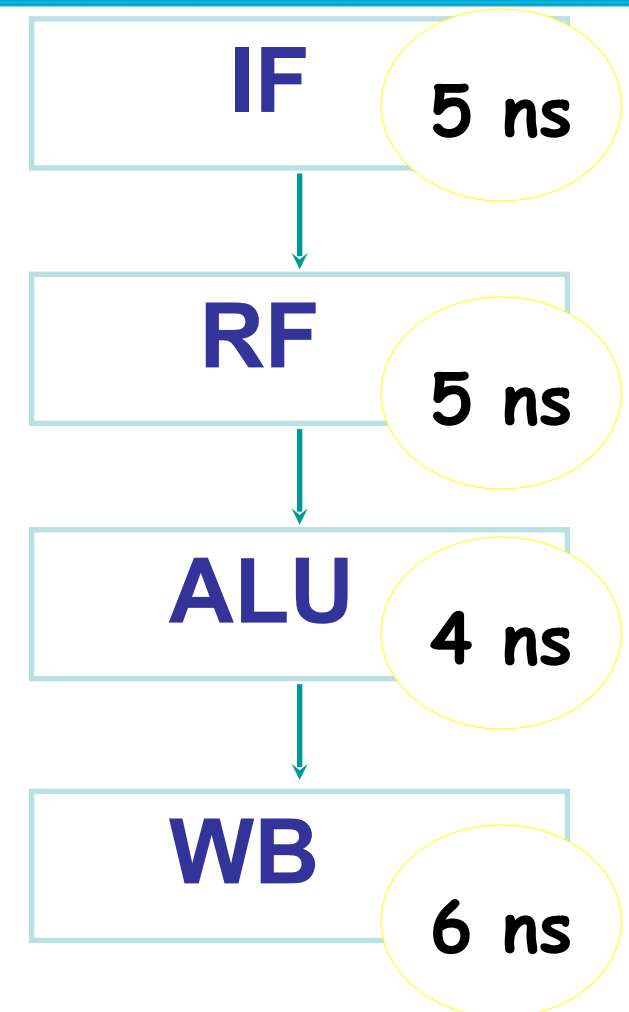
- **Instruction Fetch:** Maintains **PC**. Fetches one instruction per cycle.

- **Register File:** Reads source operands from register file.

- **ALU:** Performs indicated operation.

- **Write-Back:** Writes result back into register file.

For now:
Assume no
control
logic delay




w/o pipelining:

$t_{clk-min} = 20 \text{ ns}$

$f_{max} = 50 \text{ Mhz}$

$MIPS = 50 \text{ MIPS}$

00101010010101000011110100001100
10001100100001111001101010010101
11001010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010101010101010101010101010101

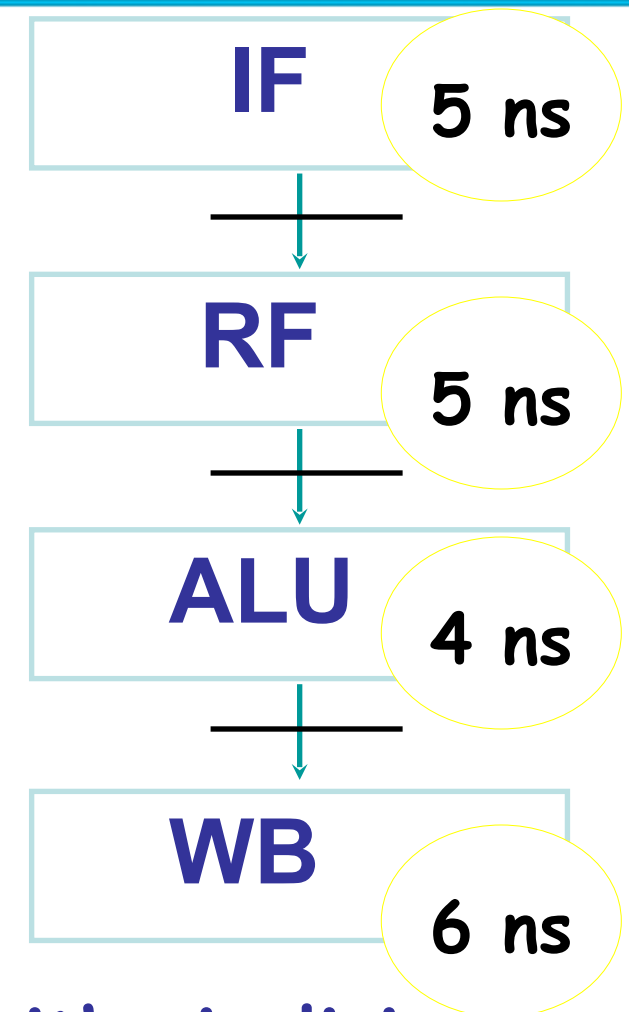


DISCS

Pipelining the Datapath

- **Goal:** Maintain (nearly) 1.0 CPI, but increase clock speed.
- **Approach:** Structure processor as 4-stage pipeline:
 - **Instruction Fetch:** Maintains **PC**. Fetches one instruction per cycle.
 - **Register File:** Reads source operands from register file.
 - **ALU:** Performs indicated operation.
 - **Write-Back:** Writes result back into register file.


Also
assume
ideal
registers



with pipelining:

$t_{clk-min} = 6 \text{ ns}$
 $f_{max} = 166 \text{ MHz}$
 $MIPS = 166 \text{ MIPS}$

00101010010101000011110100001100
10001100100001111001101010010101
110010101010101000010011001010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010101010101010101010101010101



Did You Know?

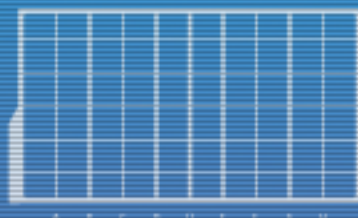
- ▶ **Factoids:**

- ▶ Light travels roughly 1 ft. / ns.
- ▶ Current processors run at around 3.0 GHz,
 - ▶ which means $t_{clk} = 0.33$ ns!

- ▶ **Thus:**

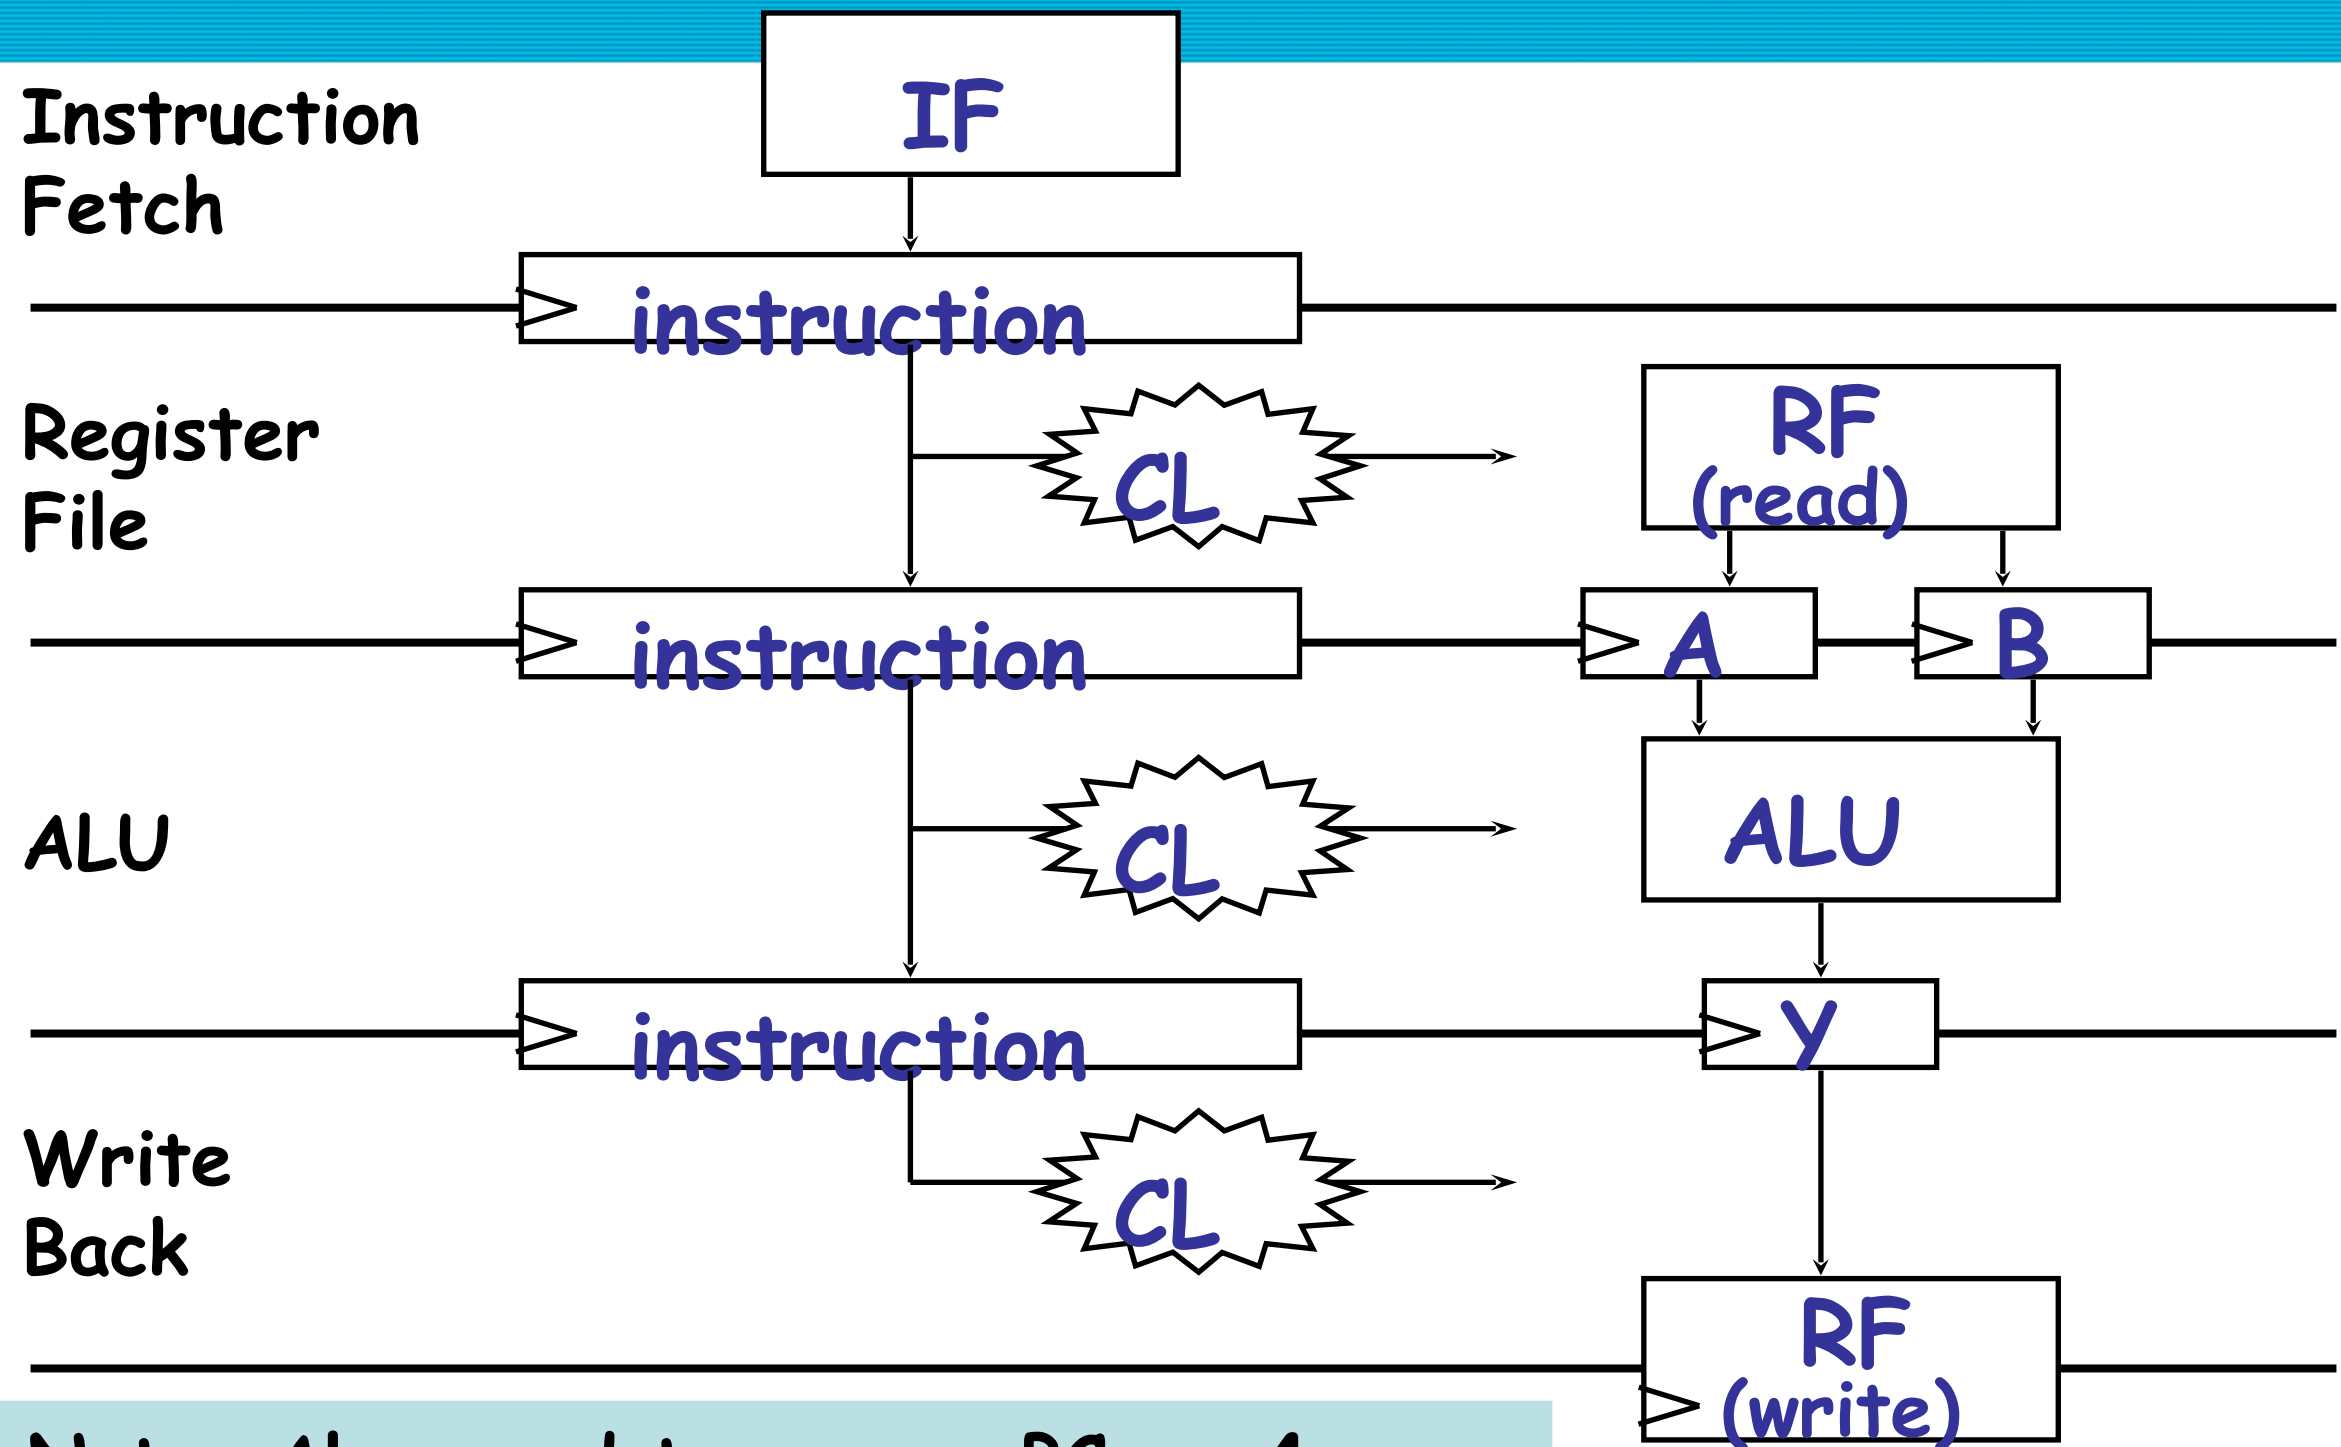
- ▶ Light only gets to travel *4 inches per instruction!*
 - ▶ And that's assuming a CPI of 1.0!
 - ▶ Many of today's processors can do CPIs of 1/2 or 1/4 by running several instructions at a time!

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010101010101010101010101010101



DISCS

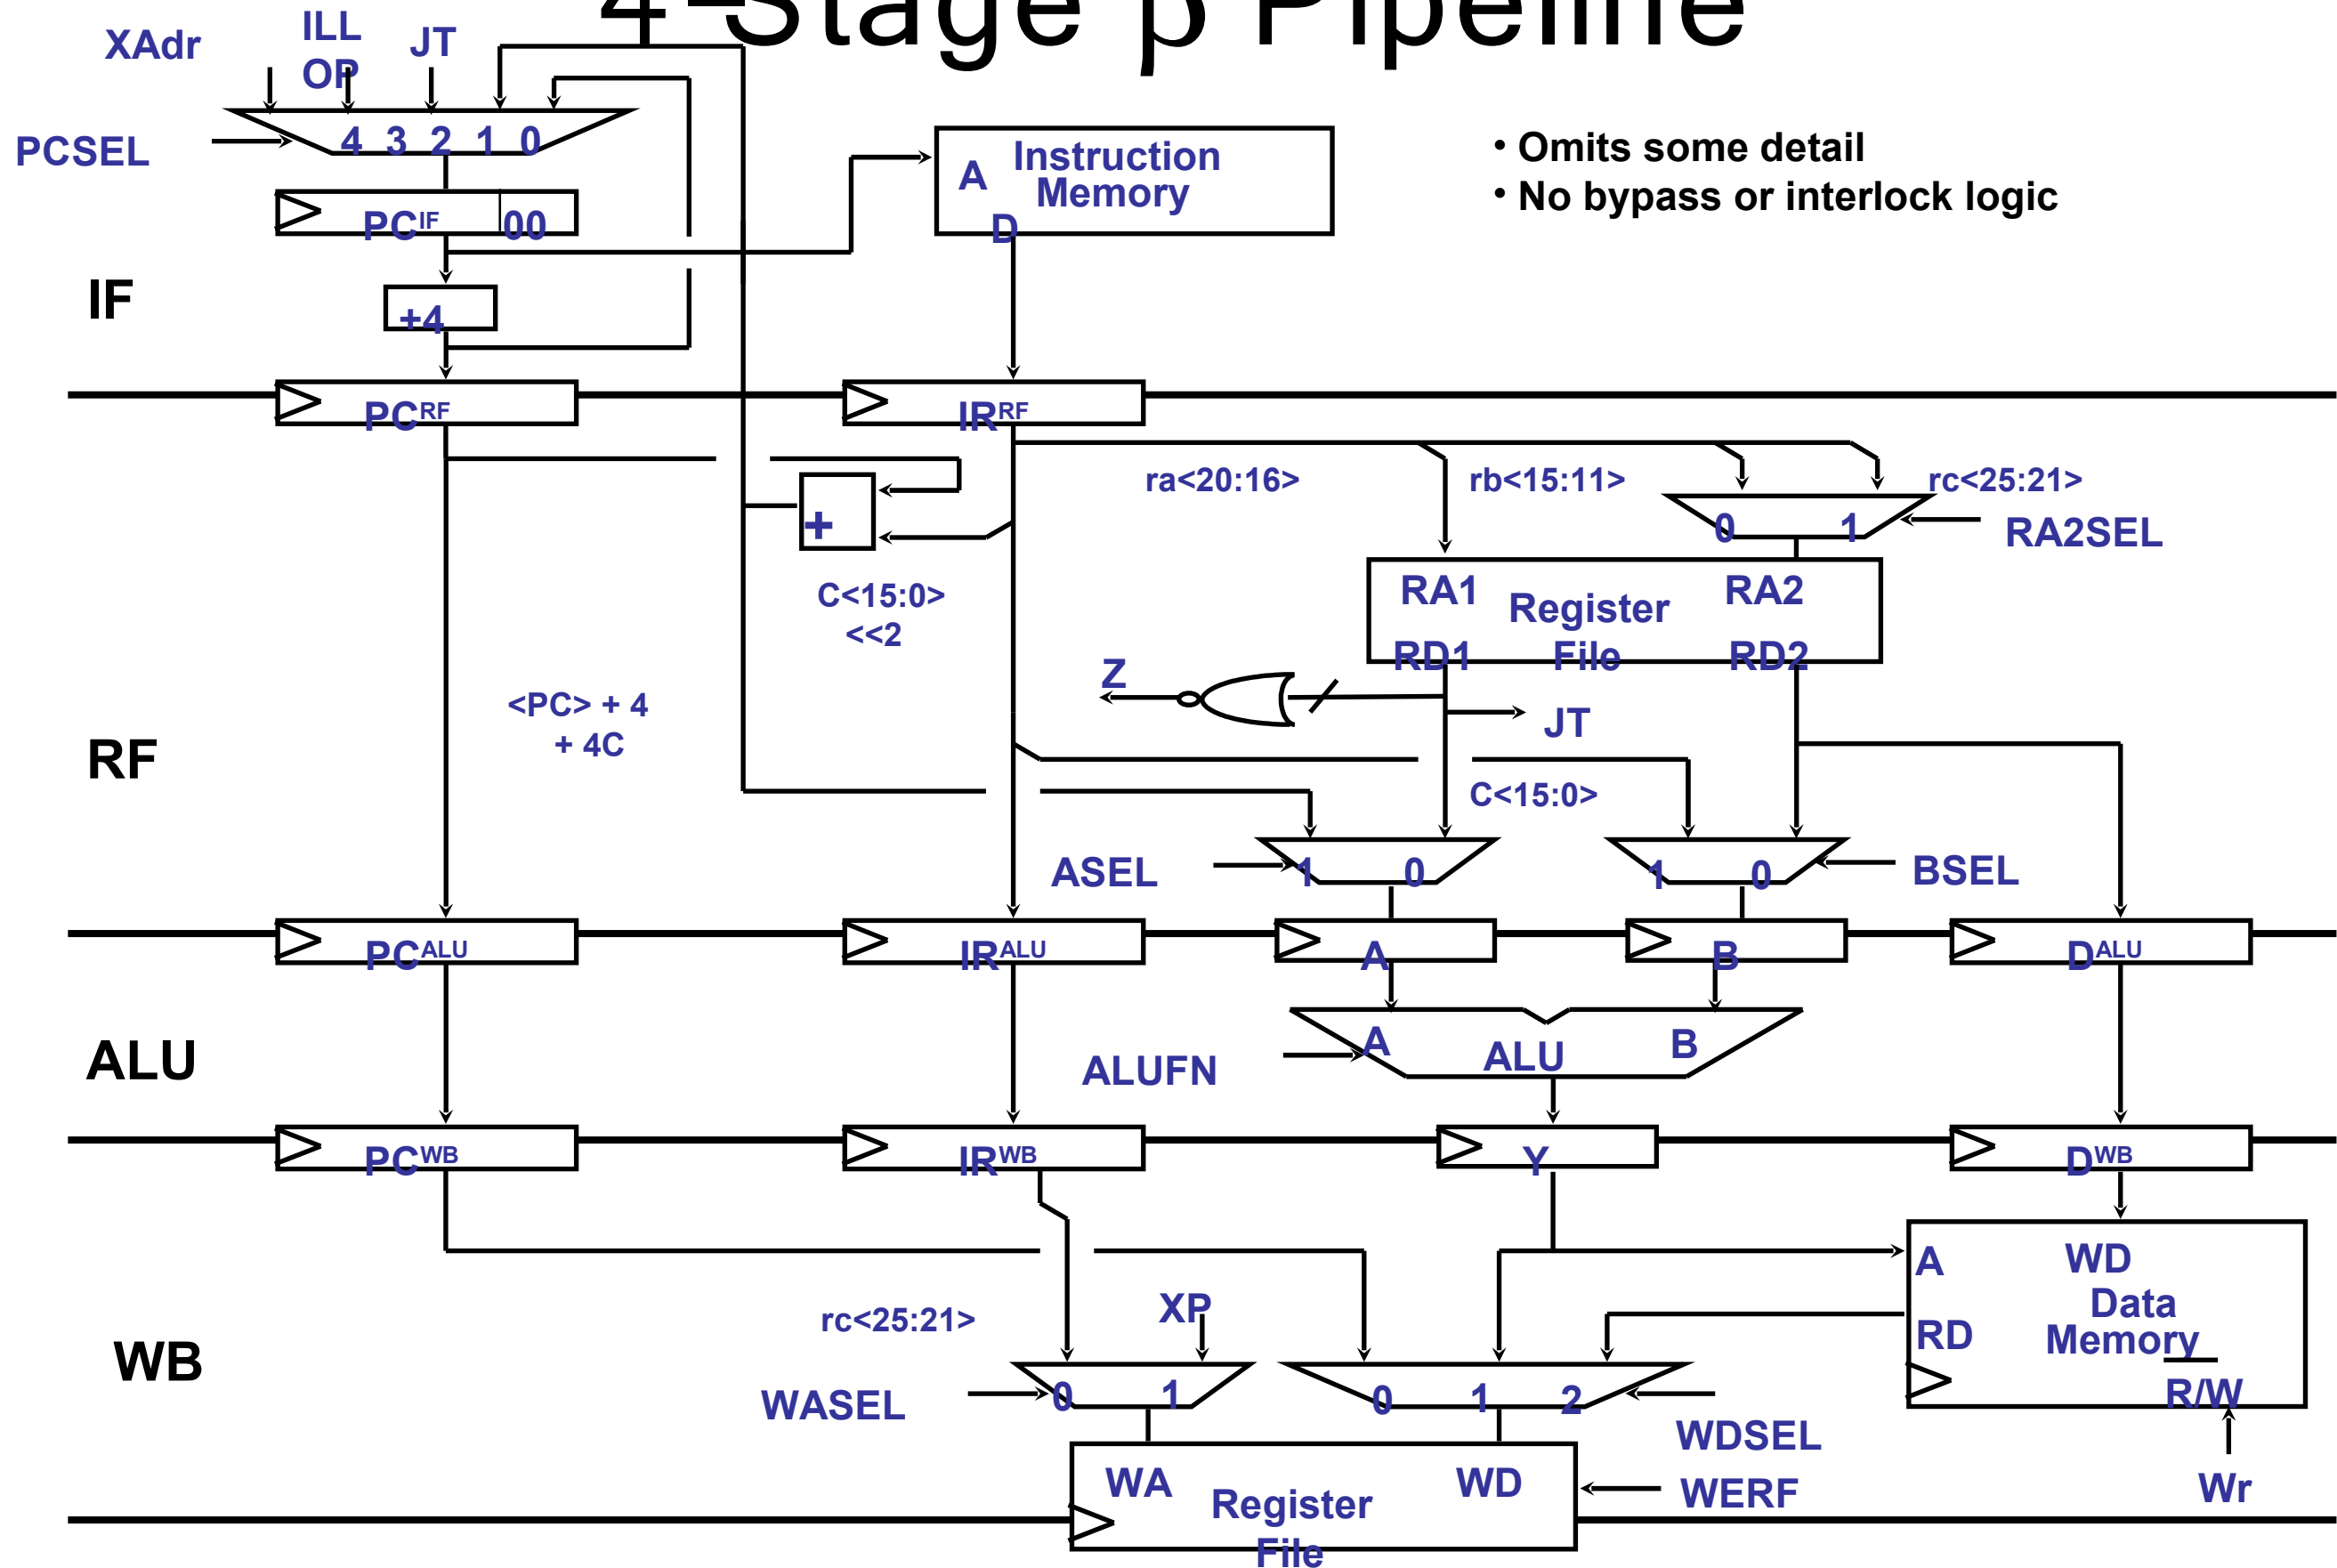
Sketch of 4-Stage Pipeline



Note: Also need to pass $\langle PC \rangle + 4$ through the stages for branch and jump instructions!



4-Stage β Pipeline



- **Omits some detail**
- **No bypass or interlock logic**

4-Pipeline Execution

- Consider a sequence of instructions:

0x100: ADDC(r1, 1, r2)


0x104: SUBC(r1, 1, r3)

0x108: XOR(r1, r5, r1)

0x10C: MUL(r2, r5, r0)

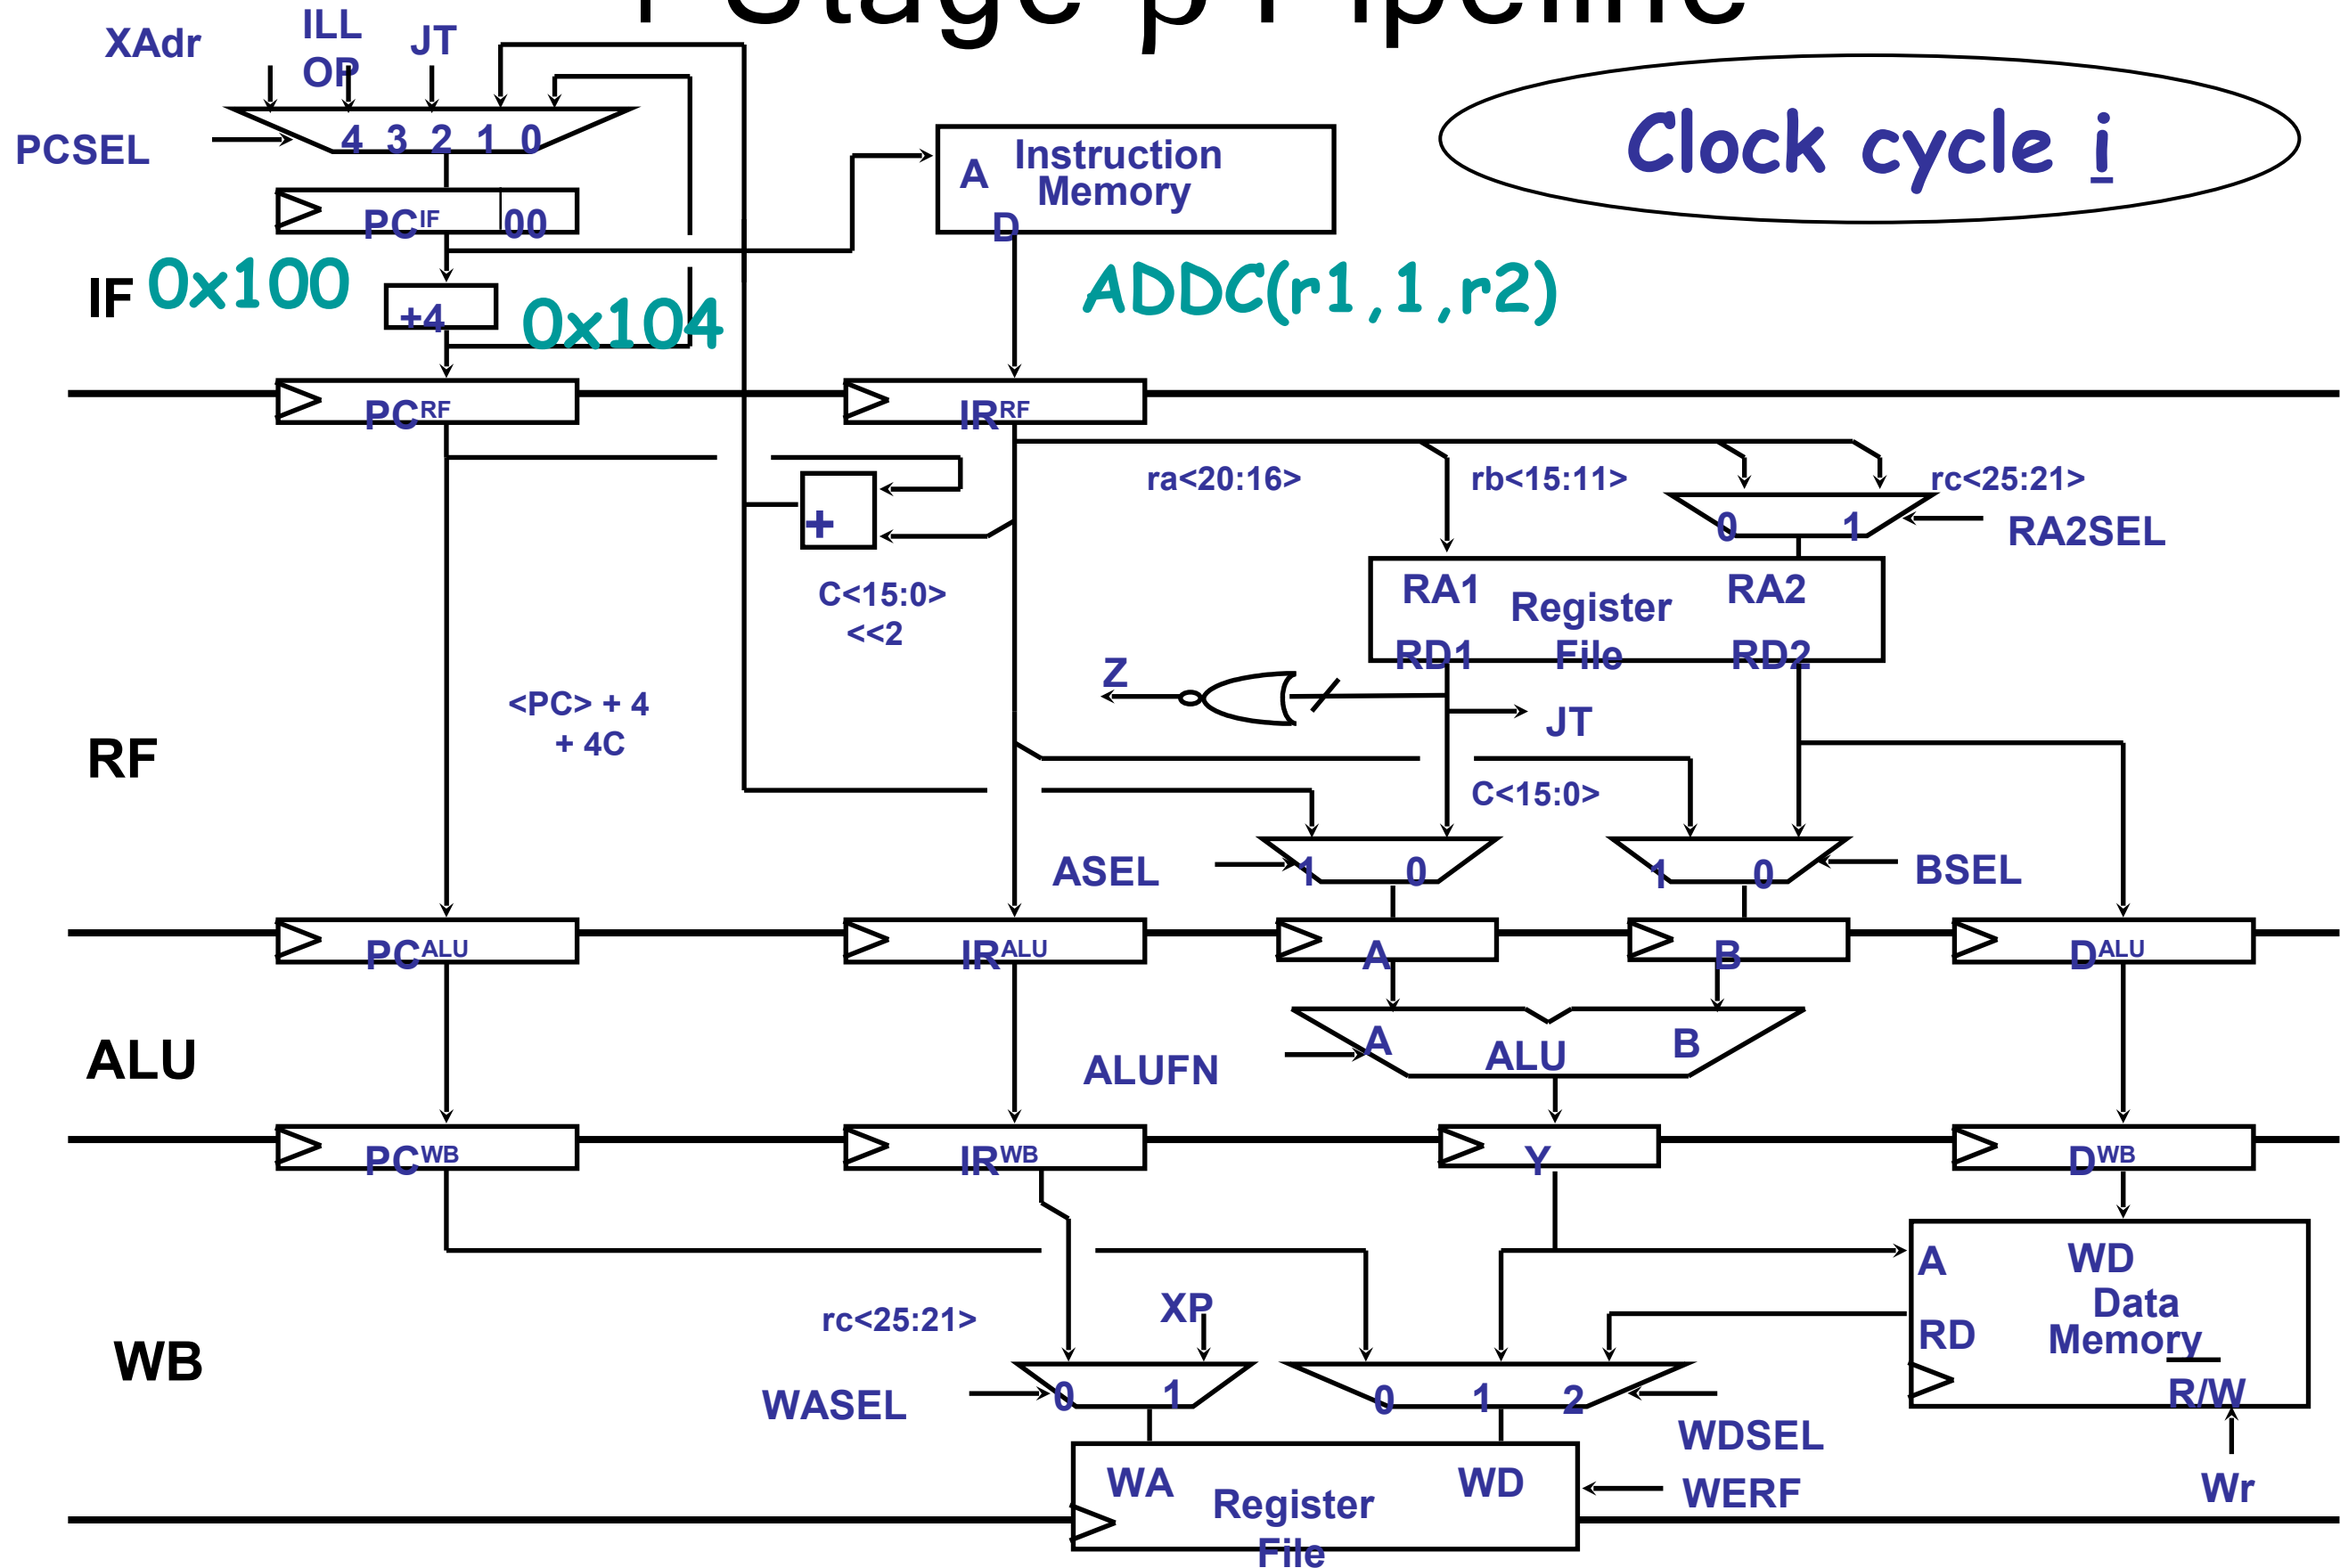
to be executed on the 4-stage pipeline...

00101010010101000011110100001100
10001100100001111001101010010101
11001010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
100101010101001010101010010101

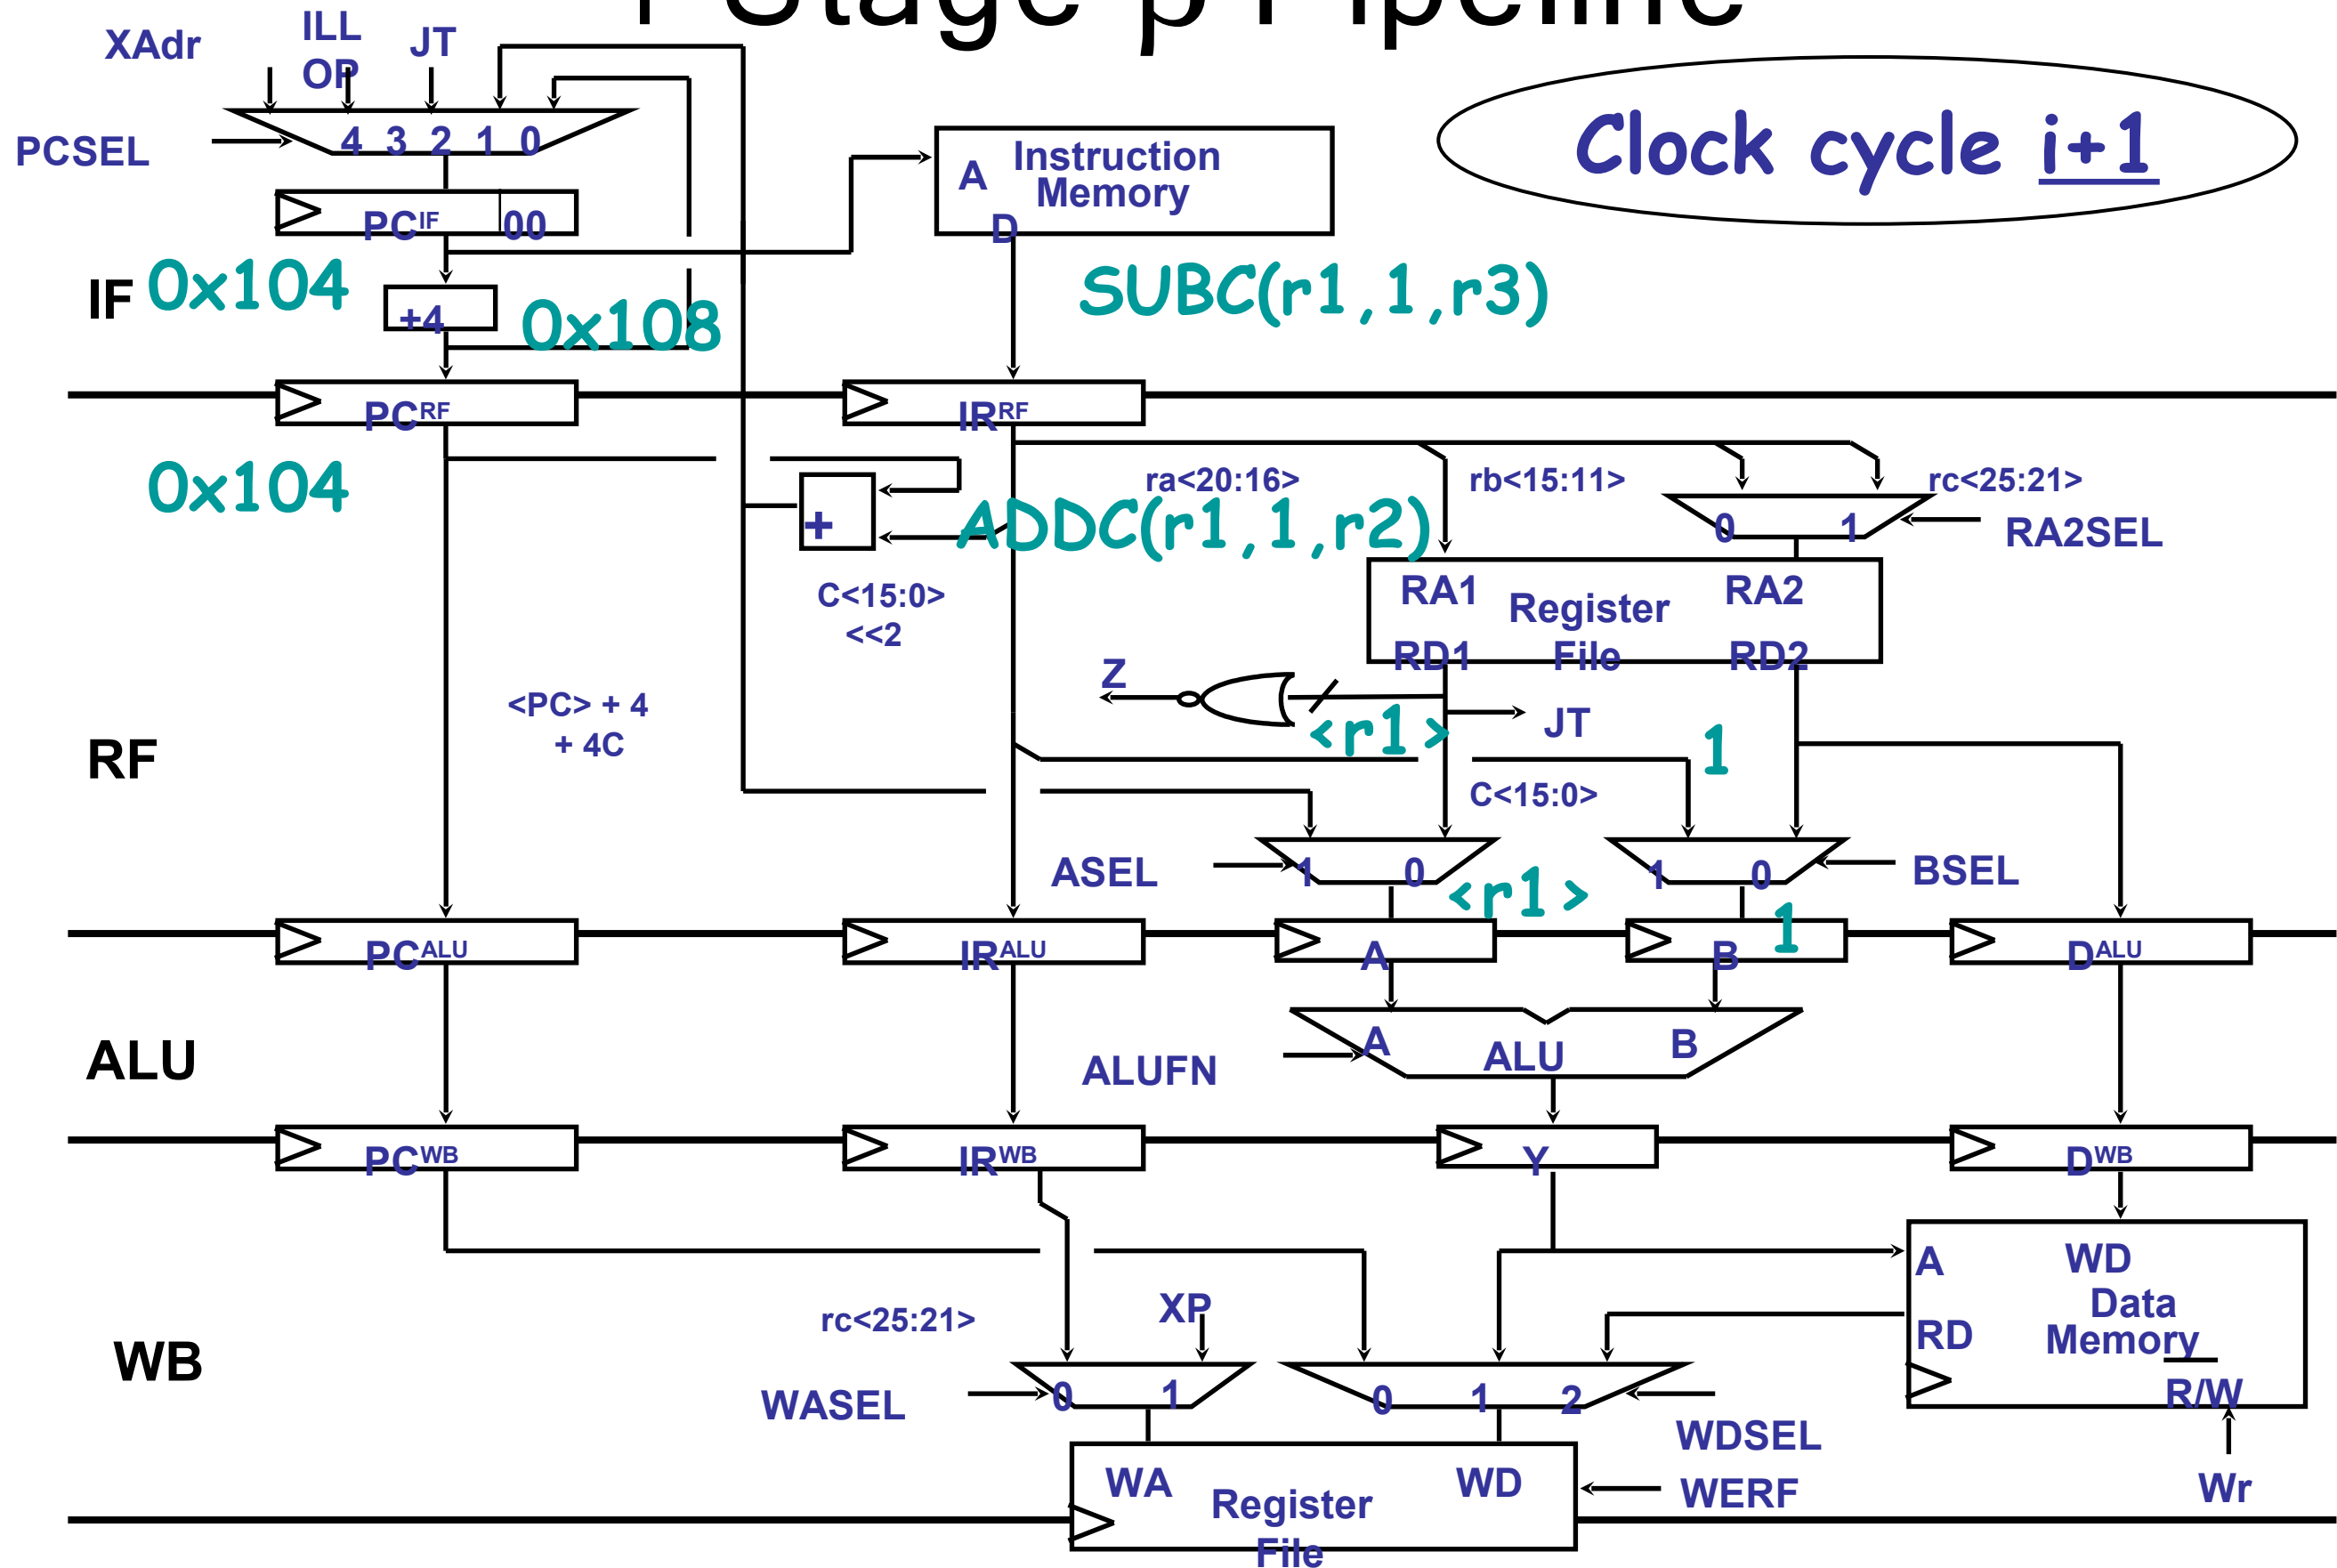


DISCS

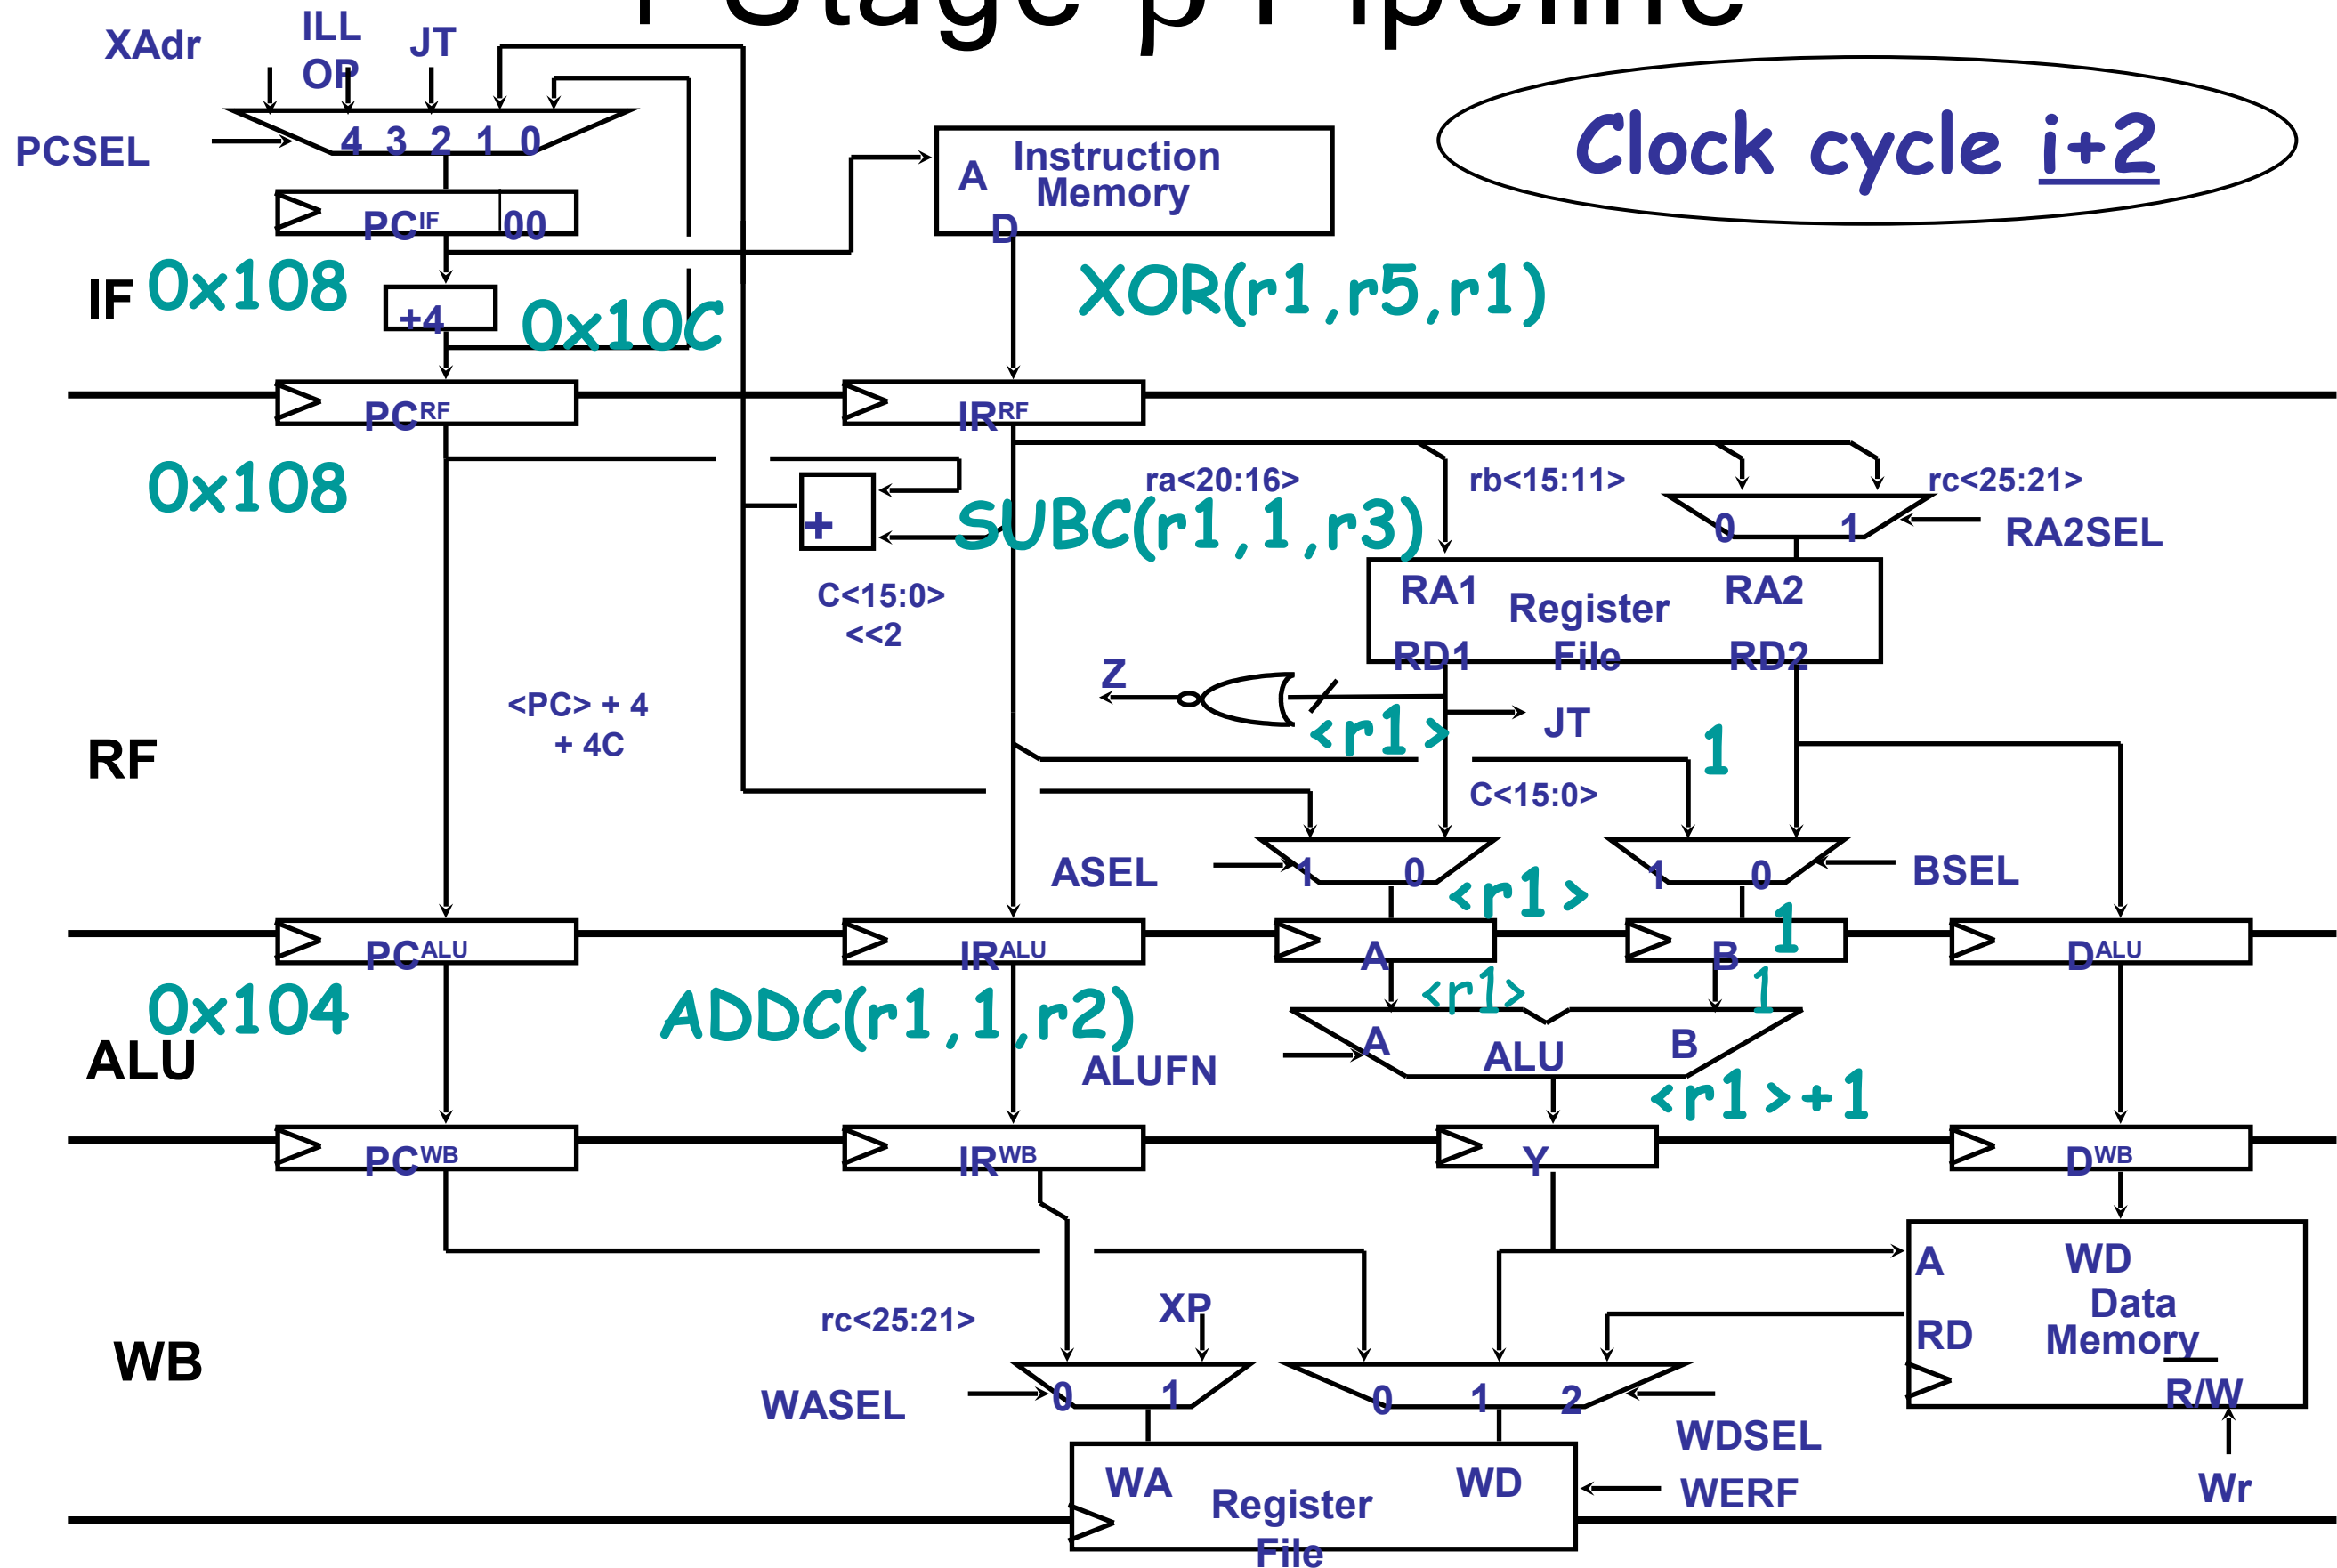
4-Stage β Pipeline



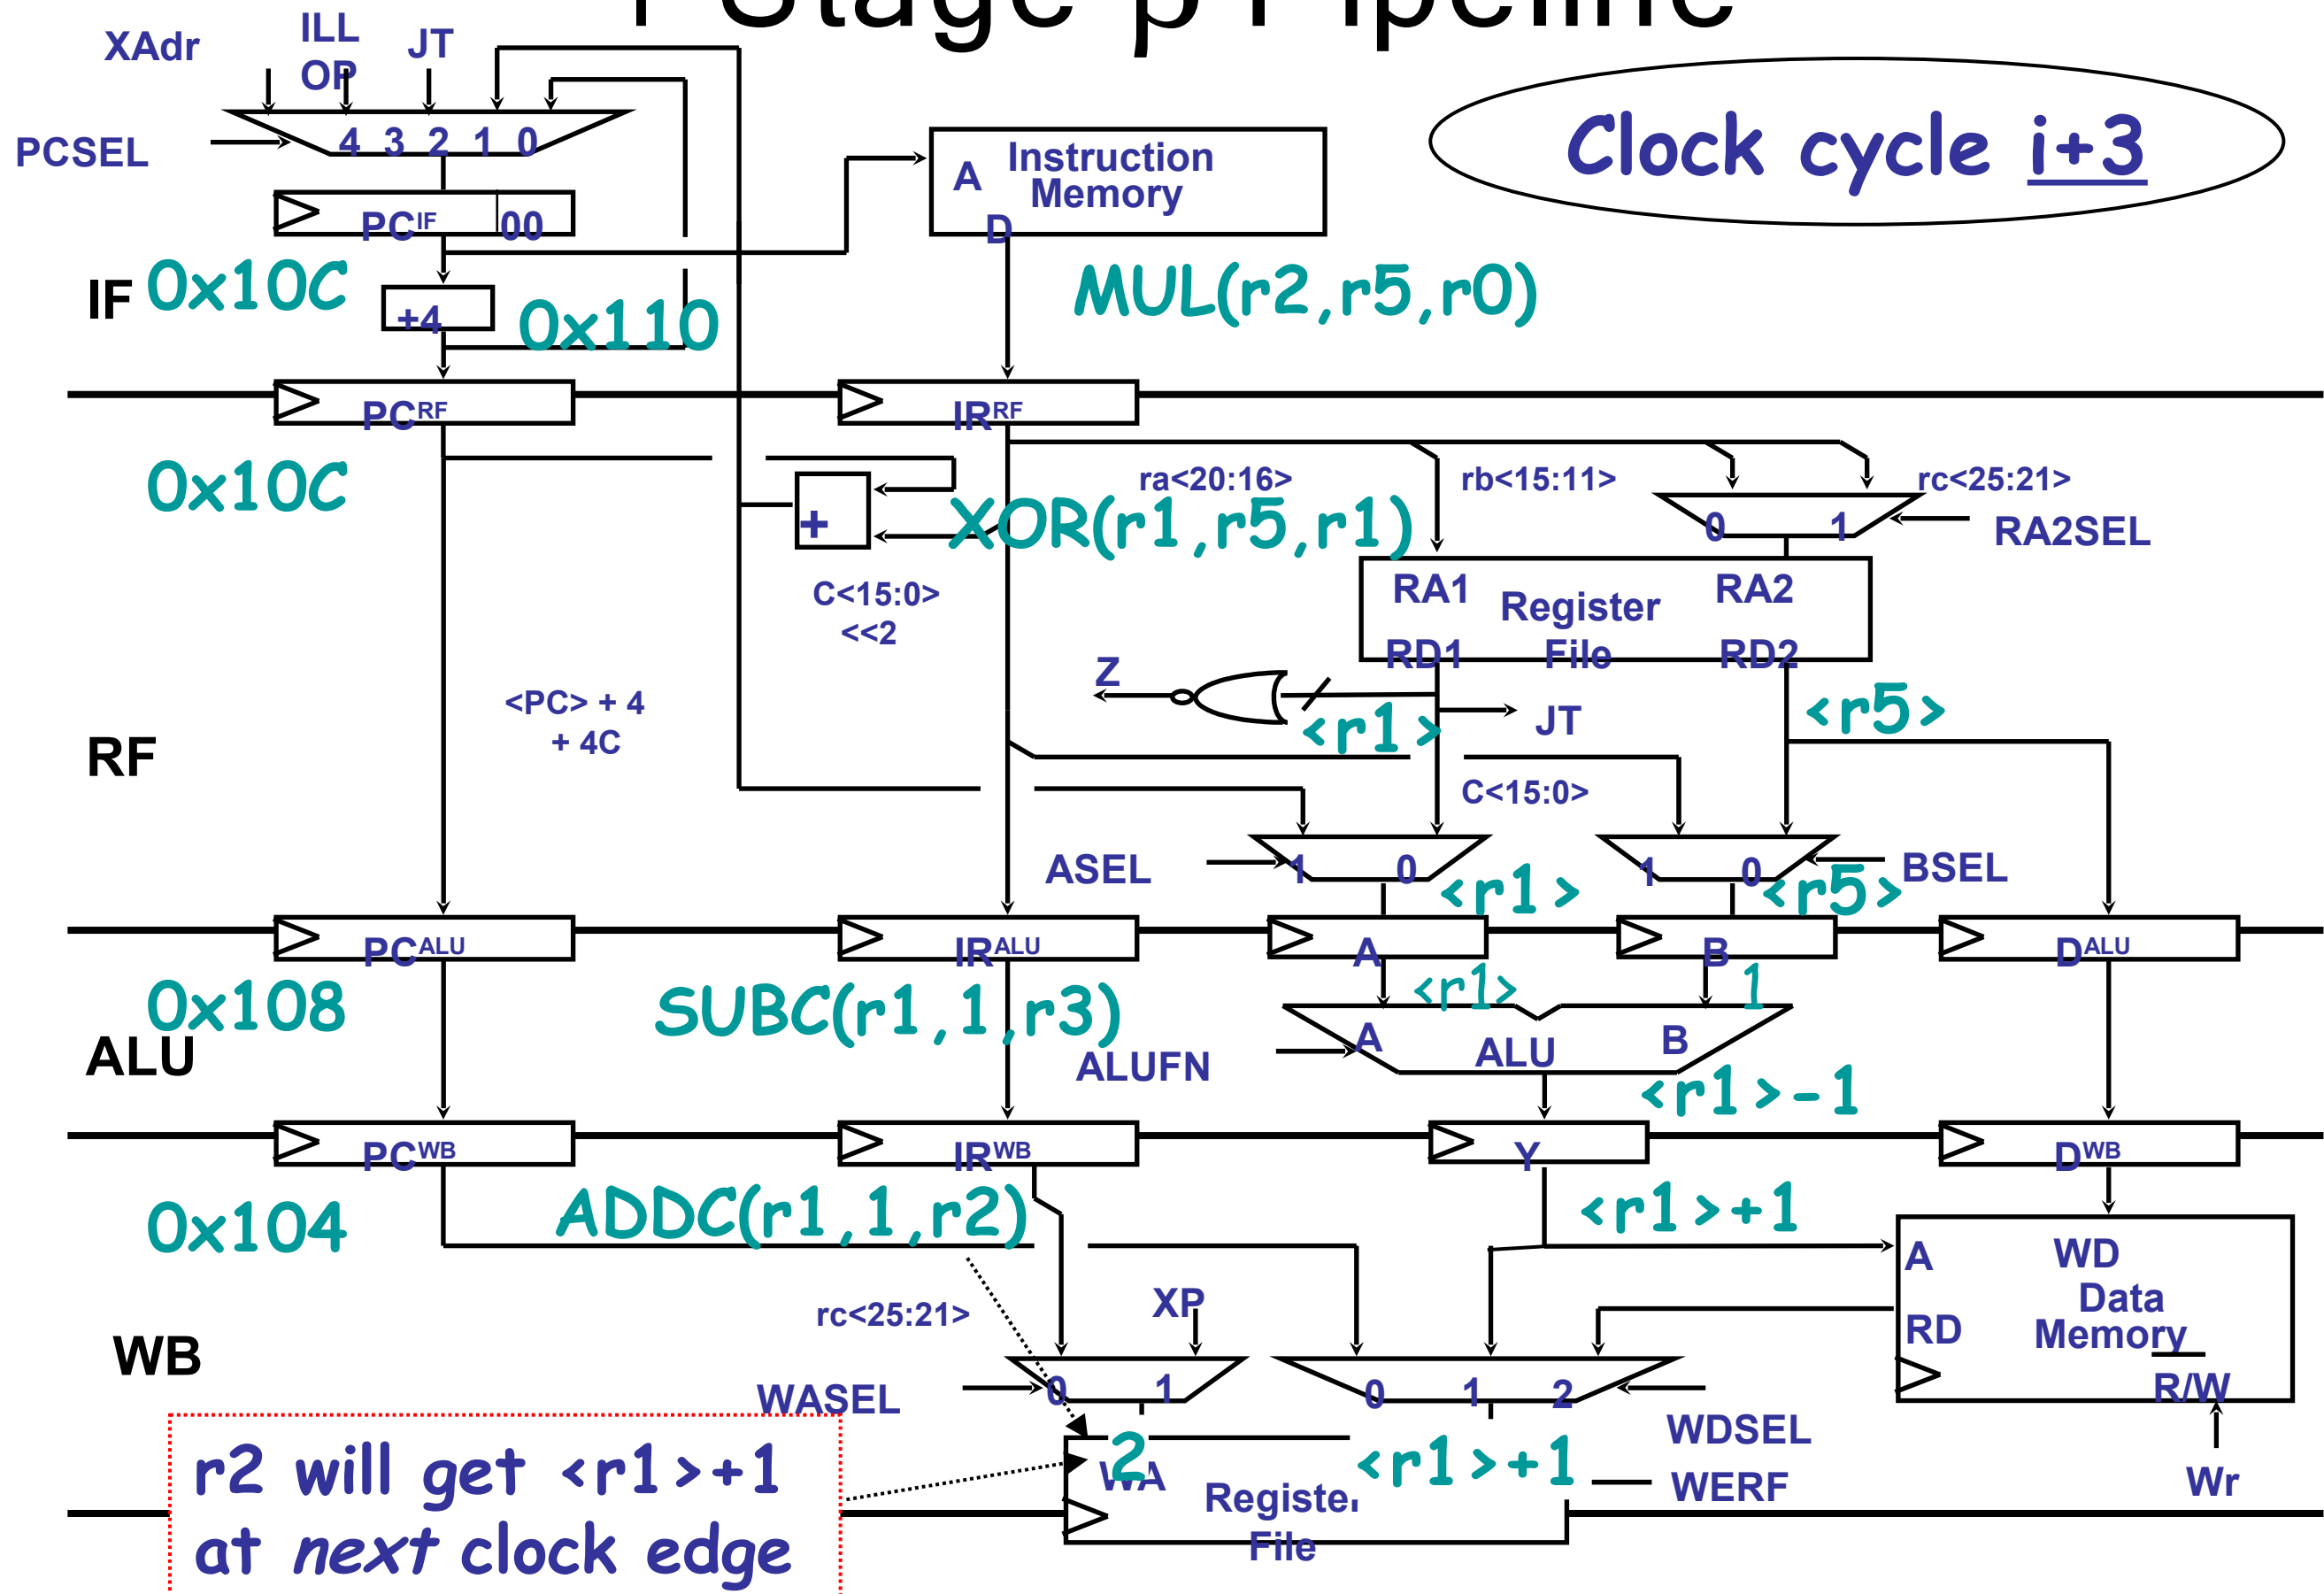
4-Stage β Pipeline



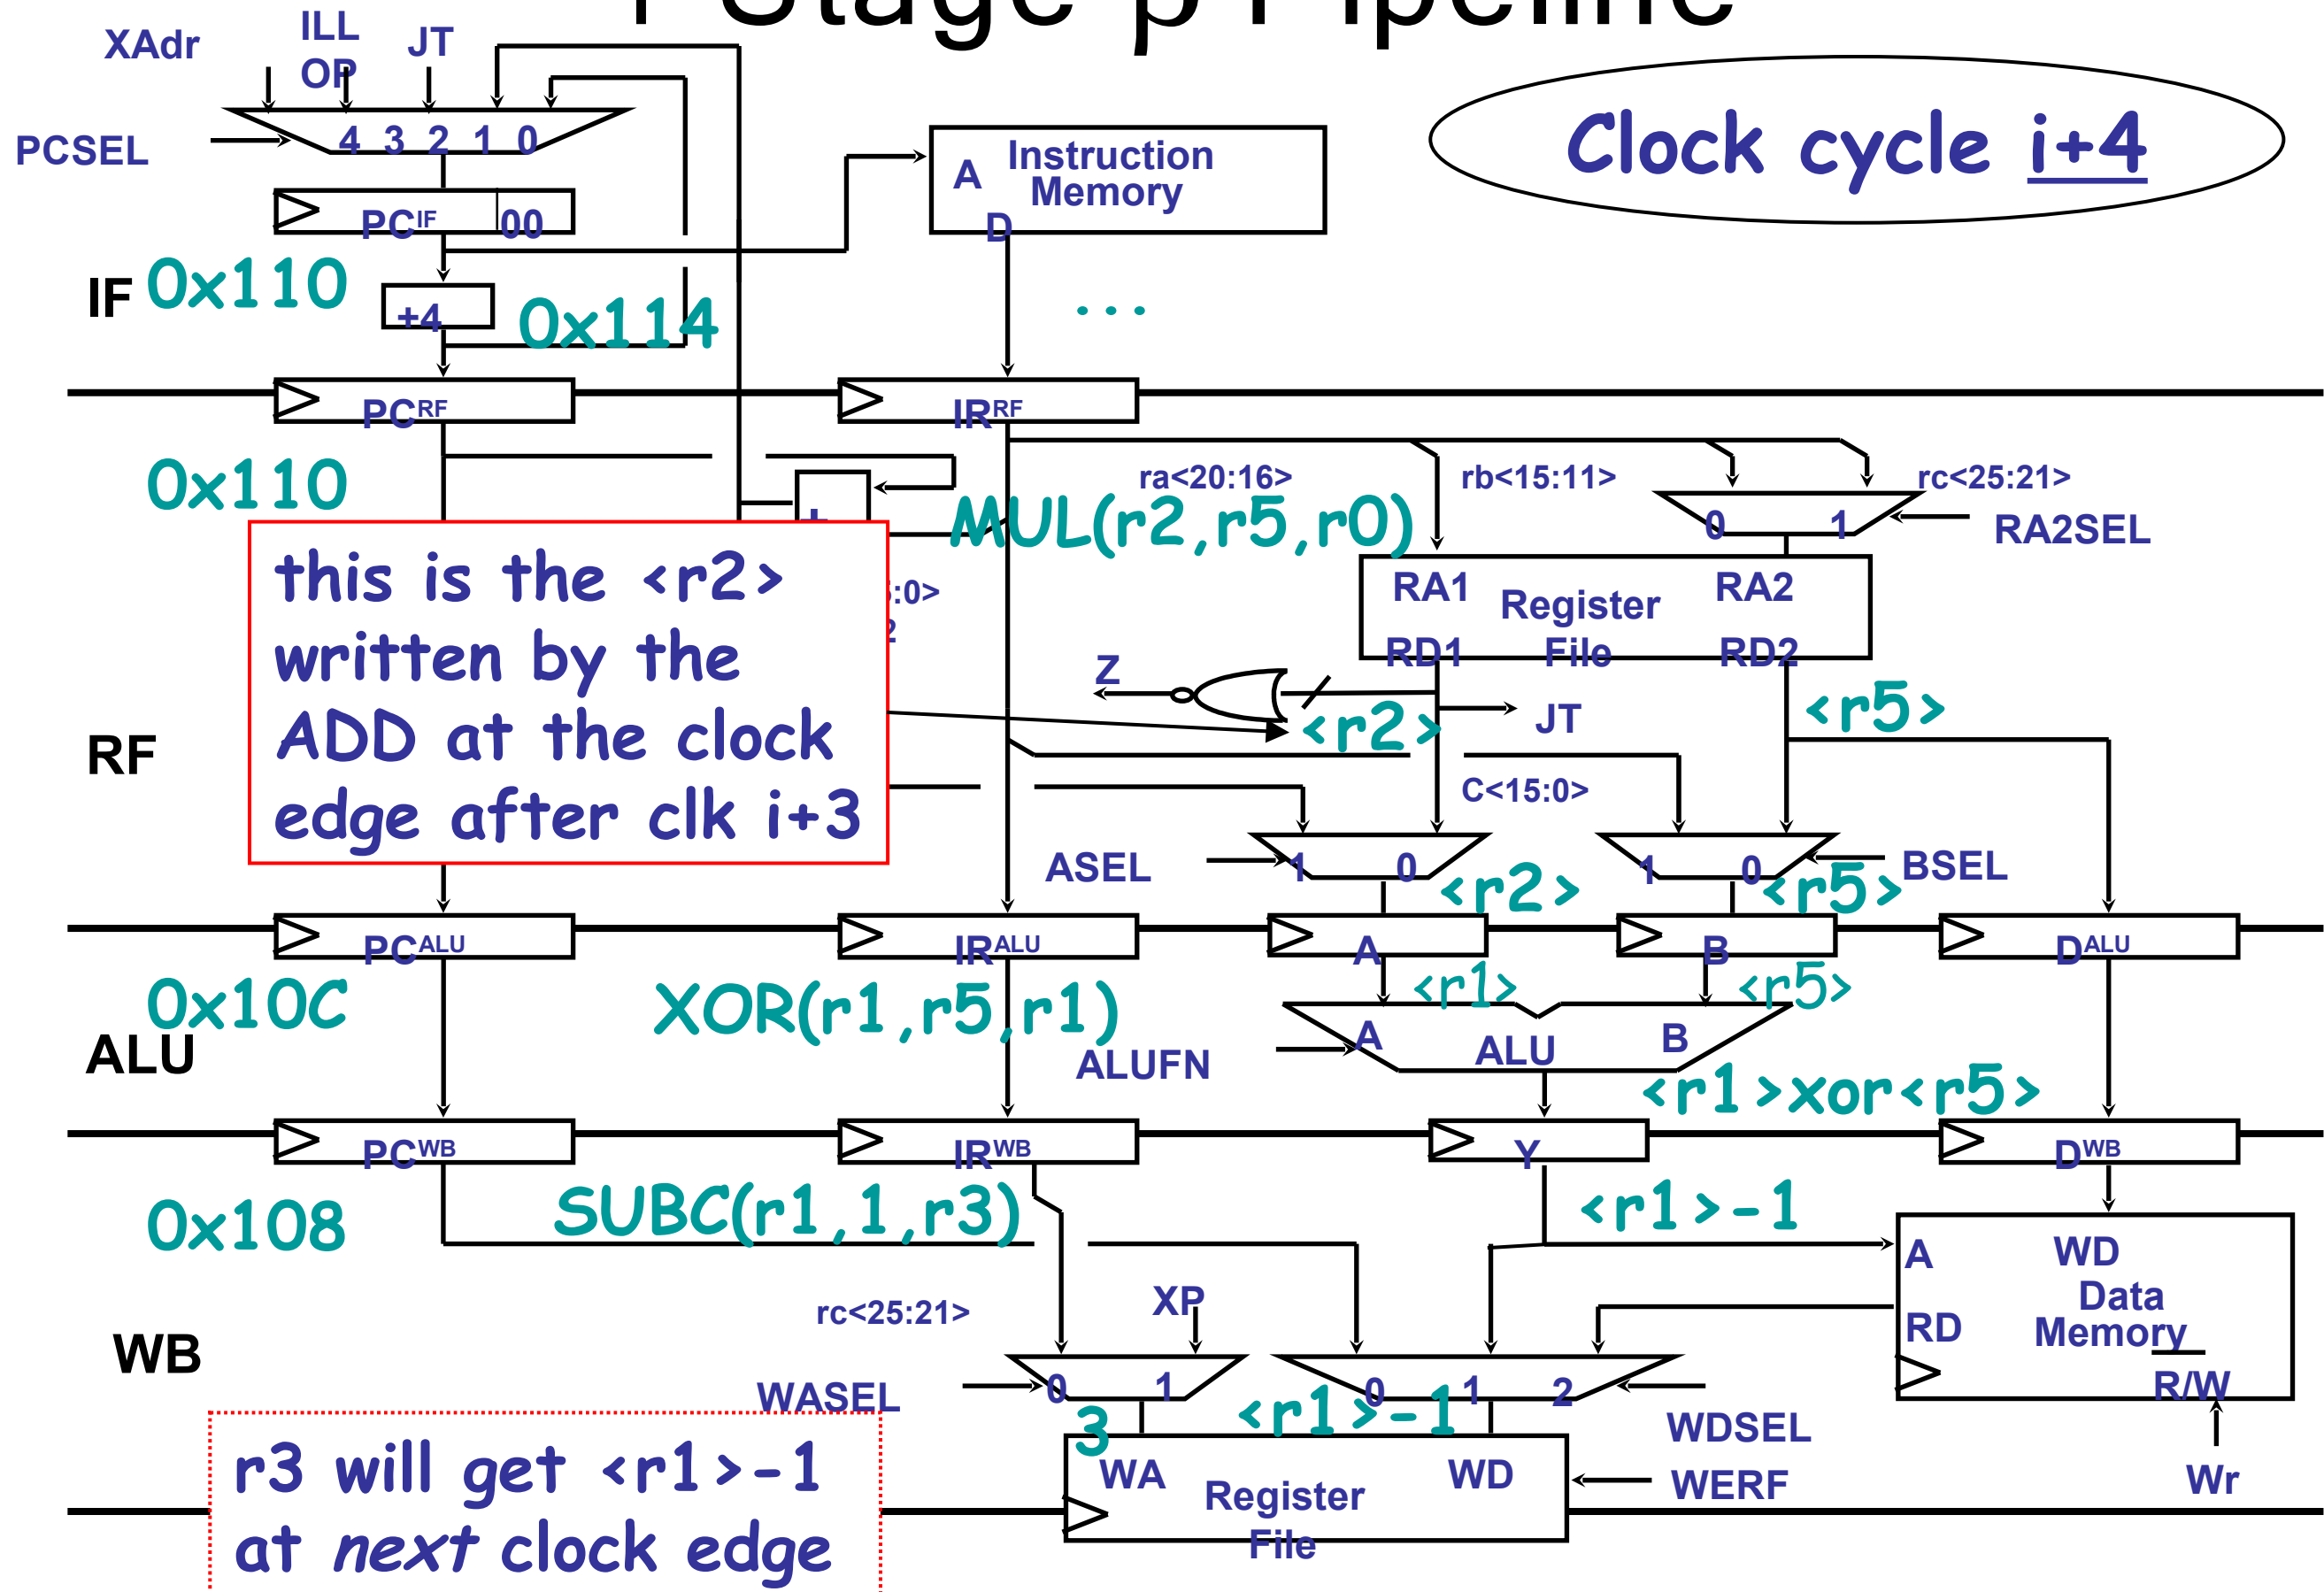
4-Stage β Pipeline



4-Stage β Pipeline



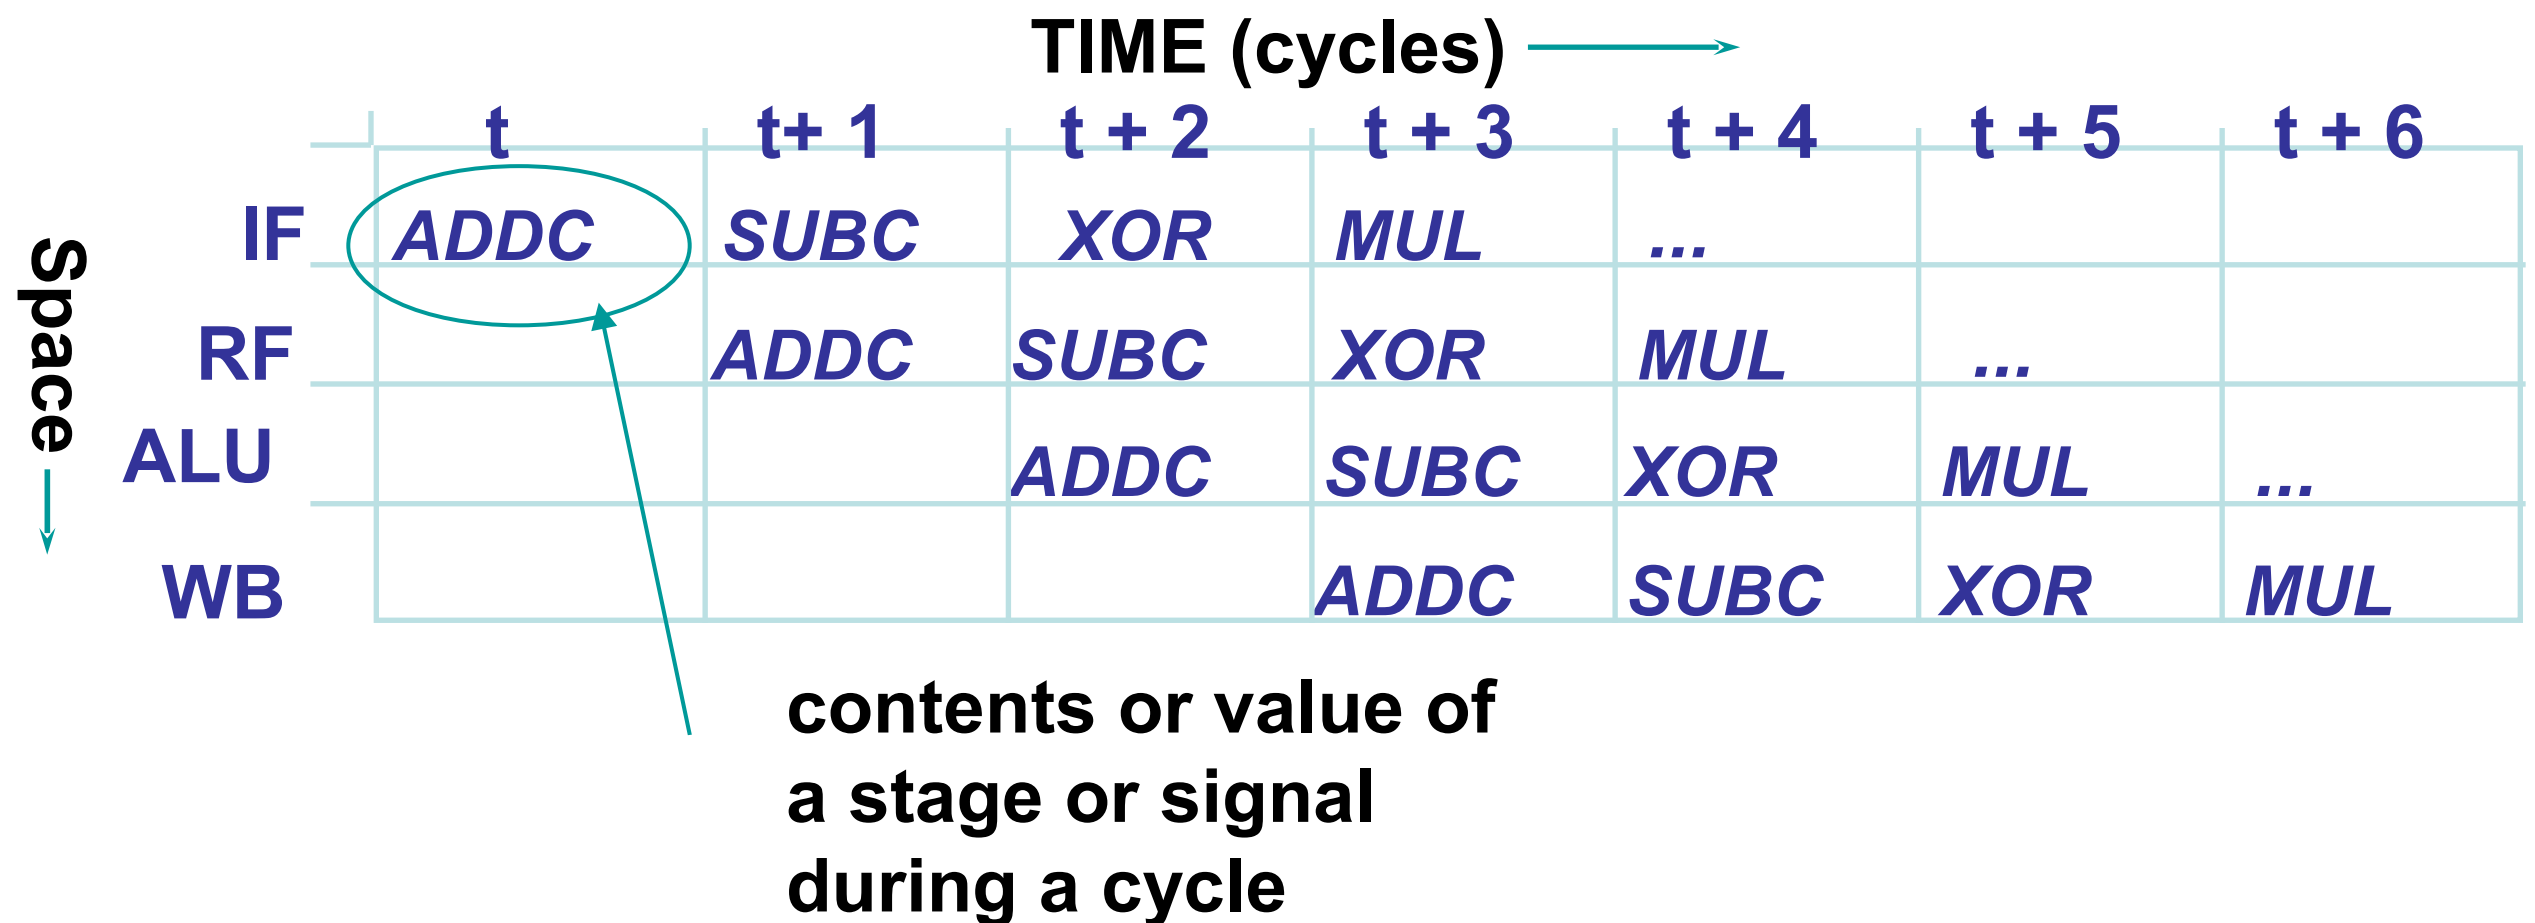
4-Stage β Pipeline



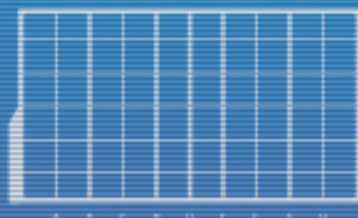
Pipeline Timing Diagram

► Space-Time Diagram

- Shows contents/values of parts/signals across time.
- Shows “What is where?” at time t .
- Can also be used to show register values (see later slides).



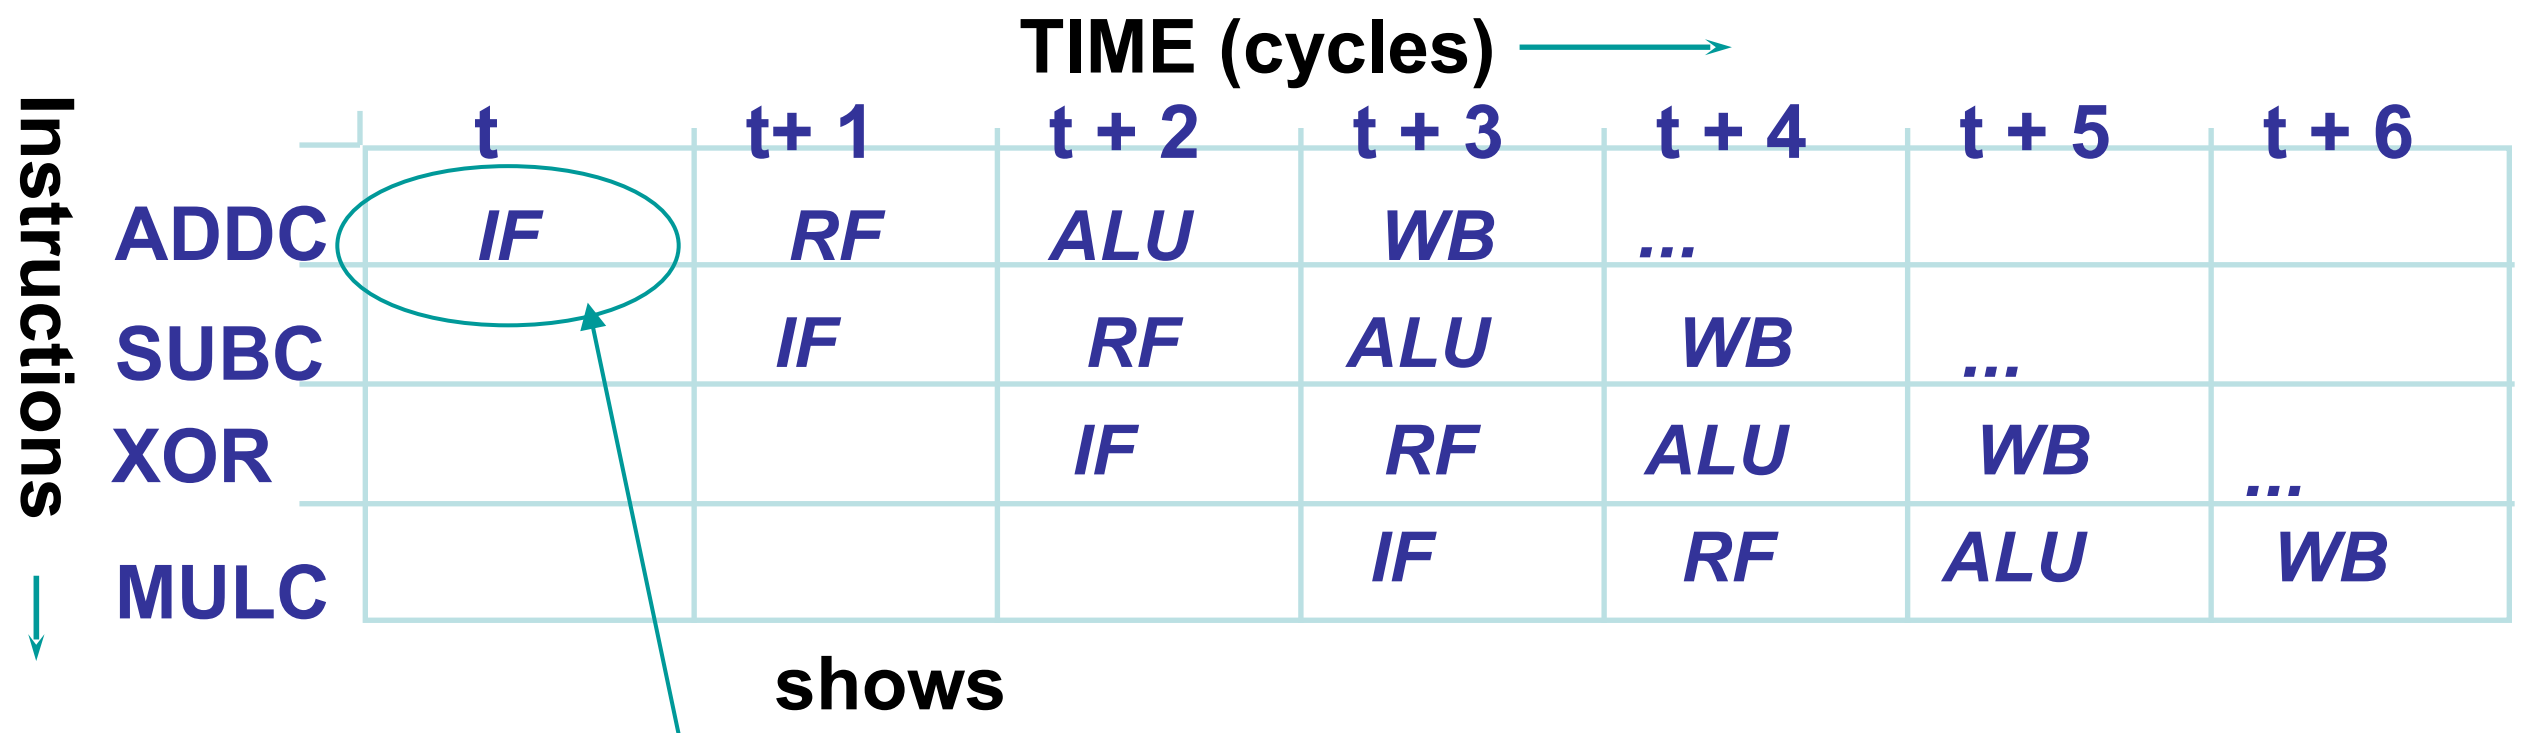
00101010010101000011110100001100
10001100100001111001101010010101
110010101010101000010011001010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010100101010010101



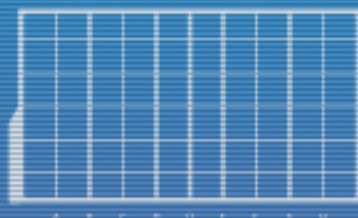
DISCS

Another Way

- ▶ Instruction-Time Diagram
 - ▶ Shows “Where am I?” at time t .
 - ▶ Used in some textbooks and tutorials.



00101010010101000011110100001100
10001100100001111001101010010101
11001010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010100101010010101



DISCS

Is there a Price to be Paid?

► ADDC reads from r1, writes to r2, but...

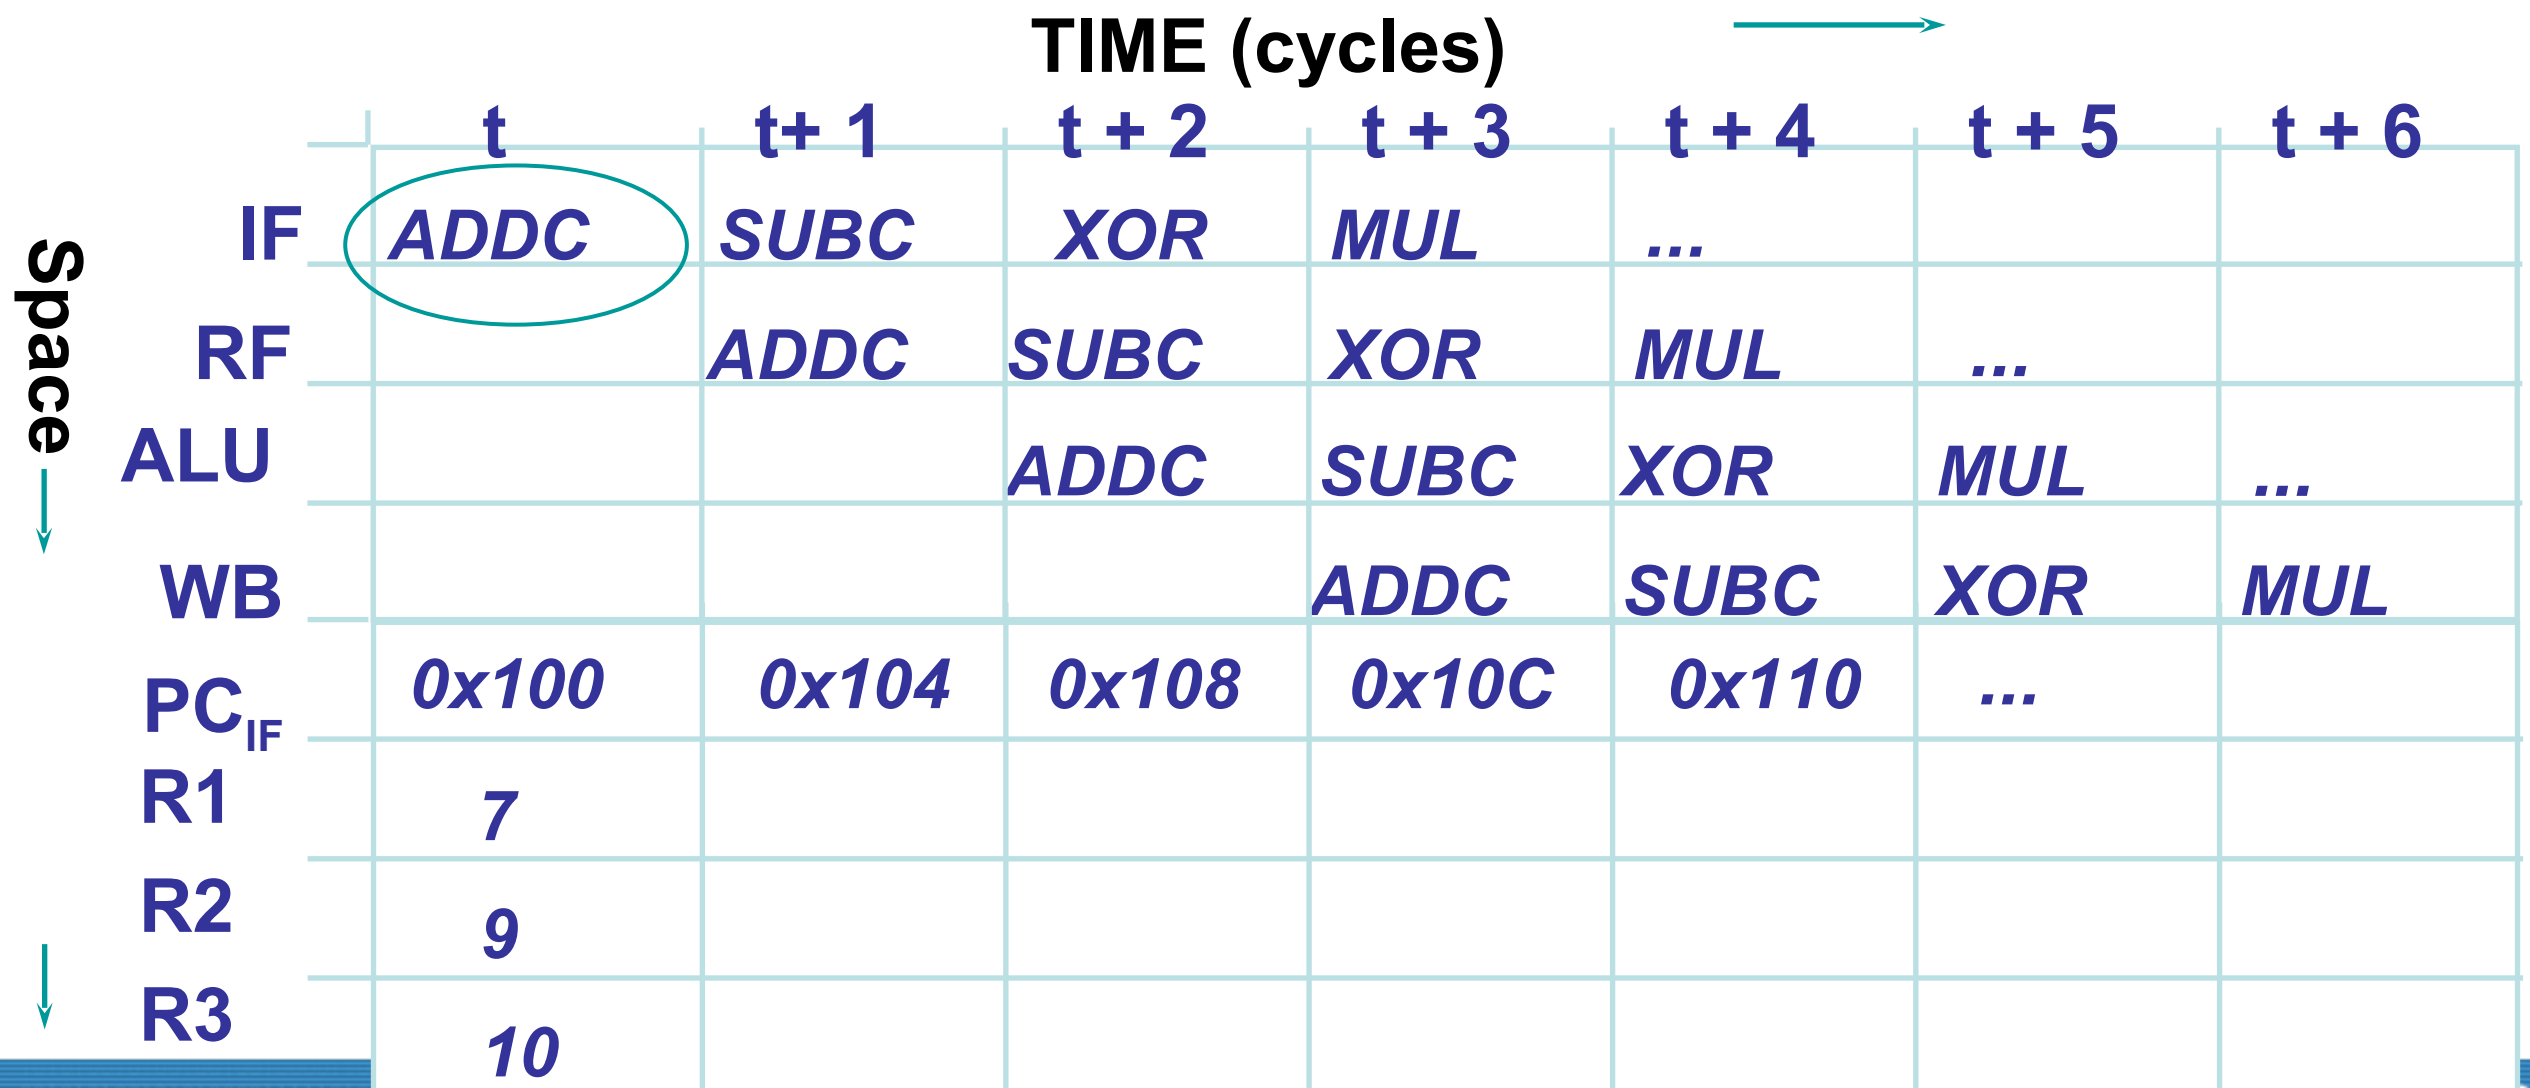
suppose: $r1=7$, $r2=9$, $r3=10$

0x100: ADDC(r1, 1, r2)


0x104: SUBC(r1, 1, r3)

0x108: XOR(r1, r5, r1)

0x10C: MUL(r2, r5, r0)

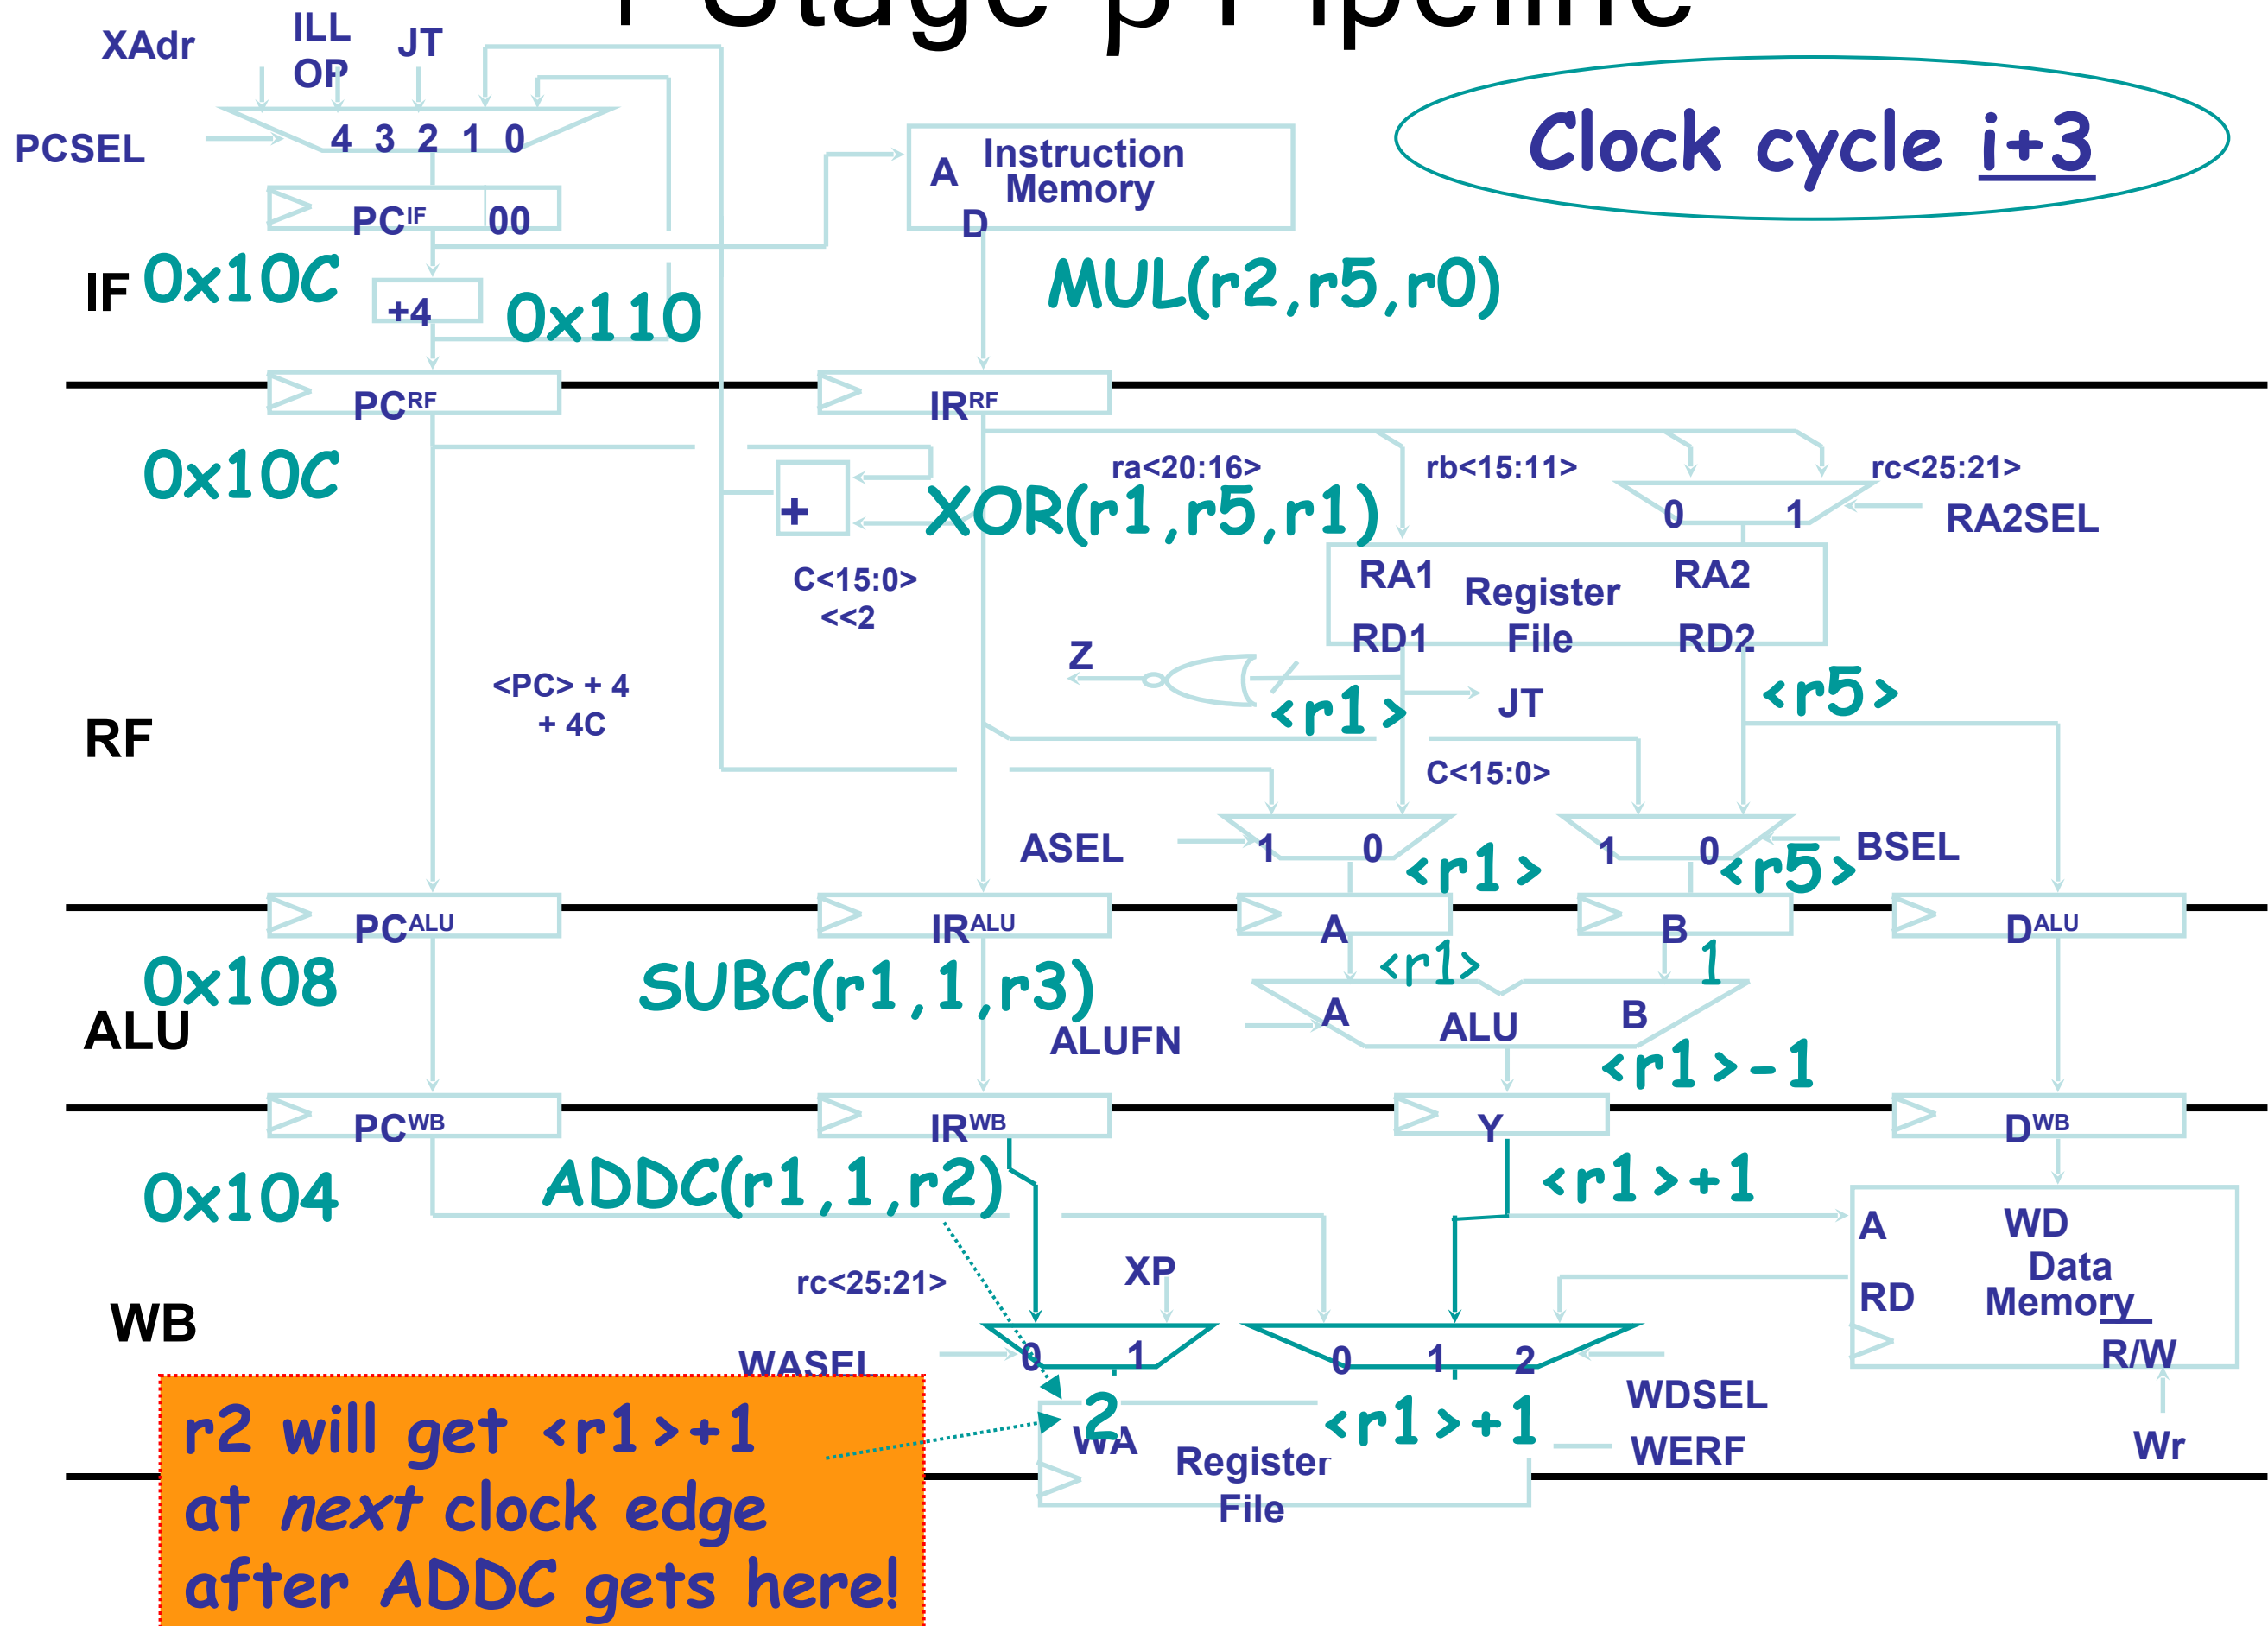


00101010010101000111101000000000
10001100100001111001101010010101
110010101010101000010011001010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010100101010010101

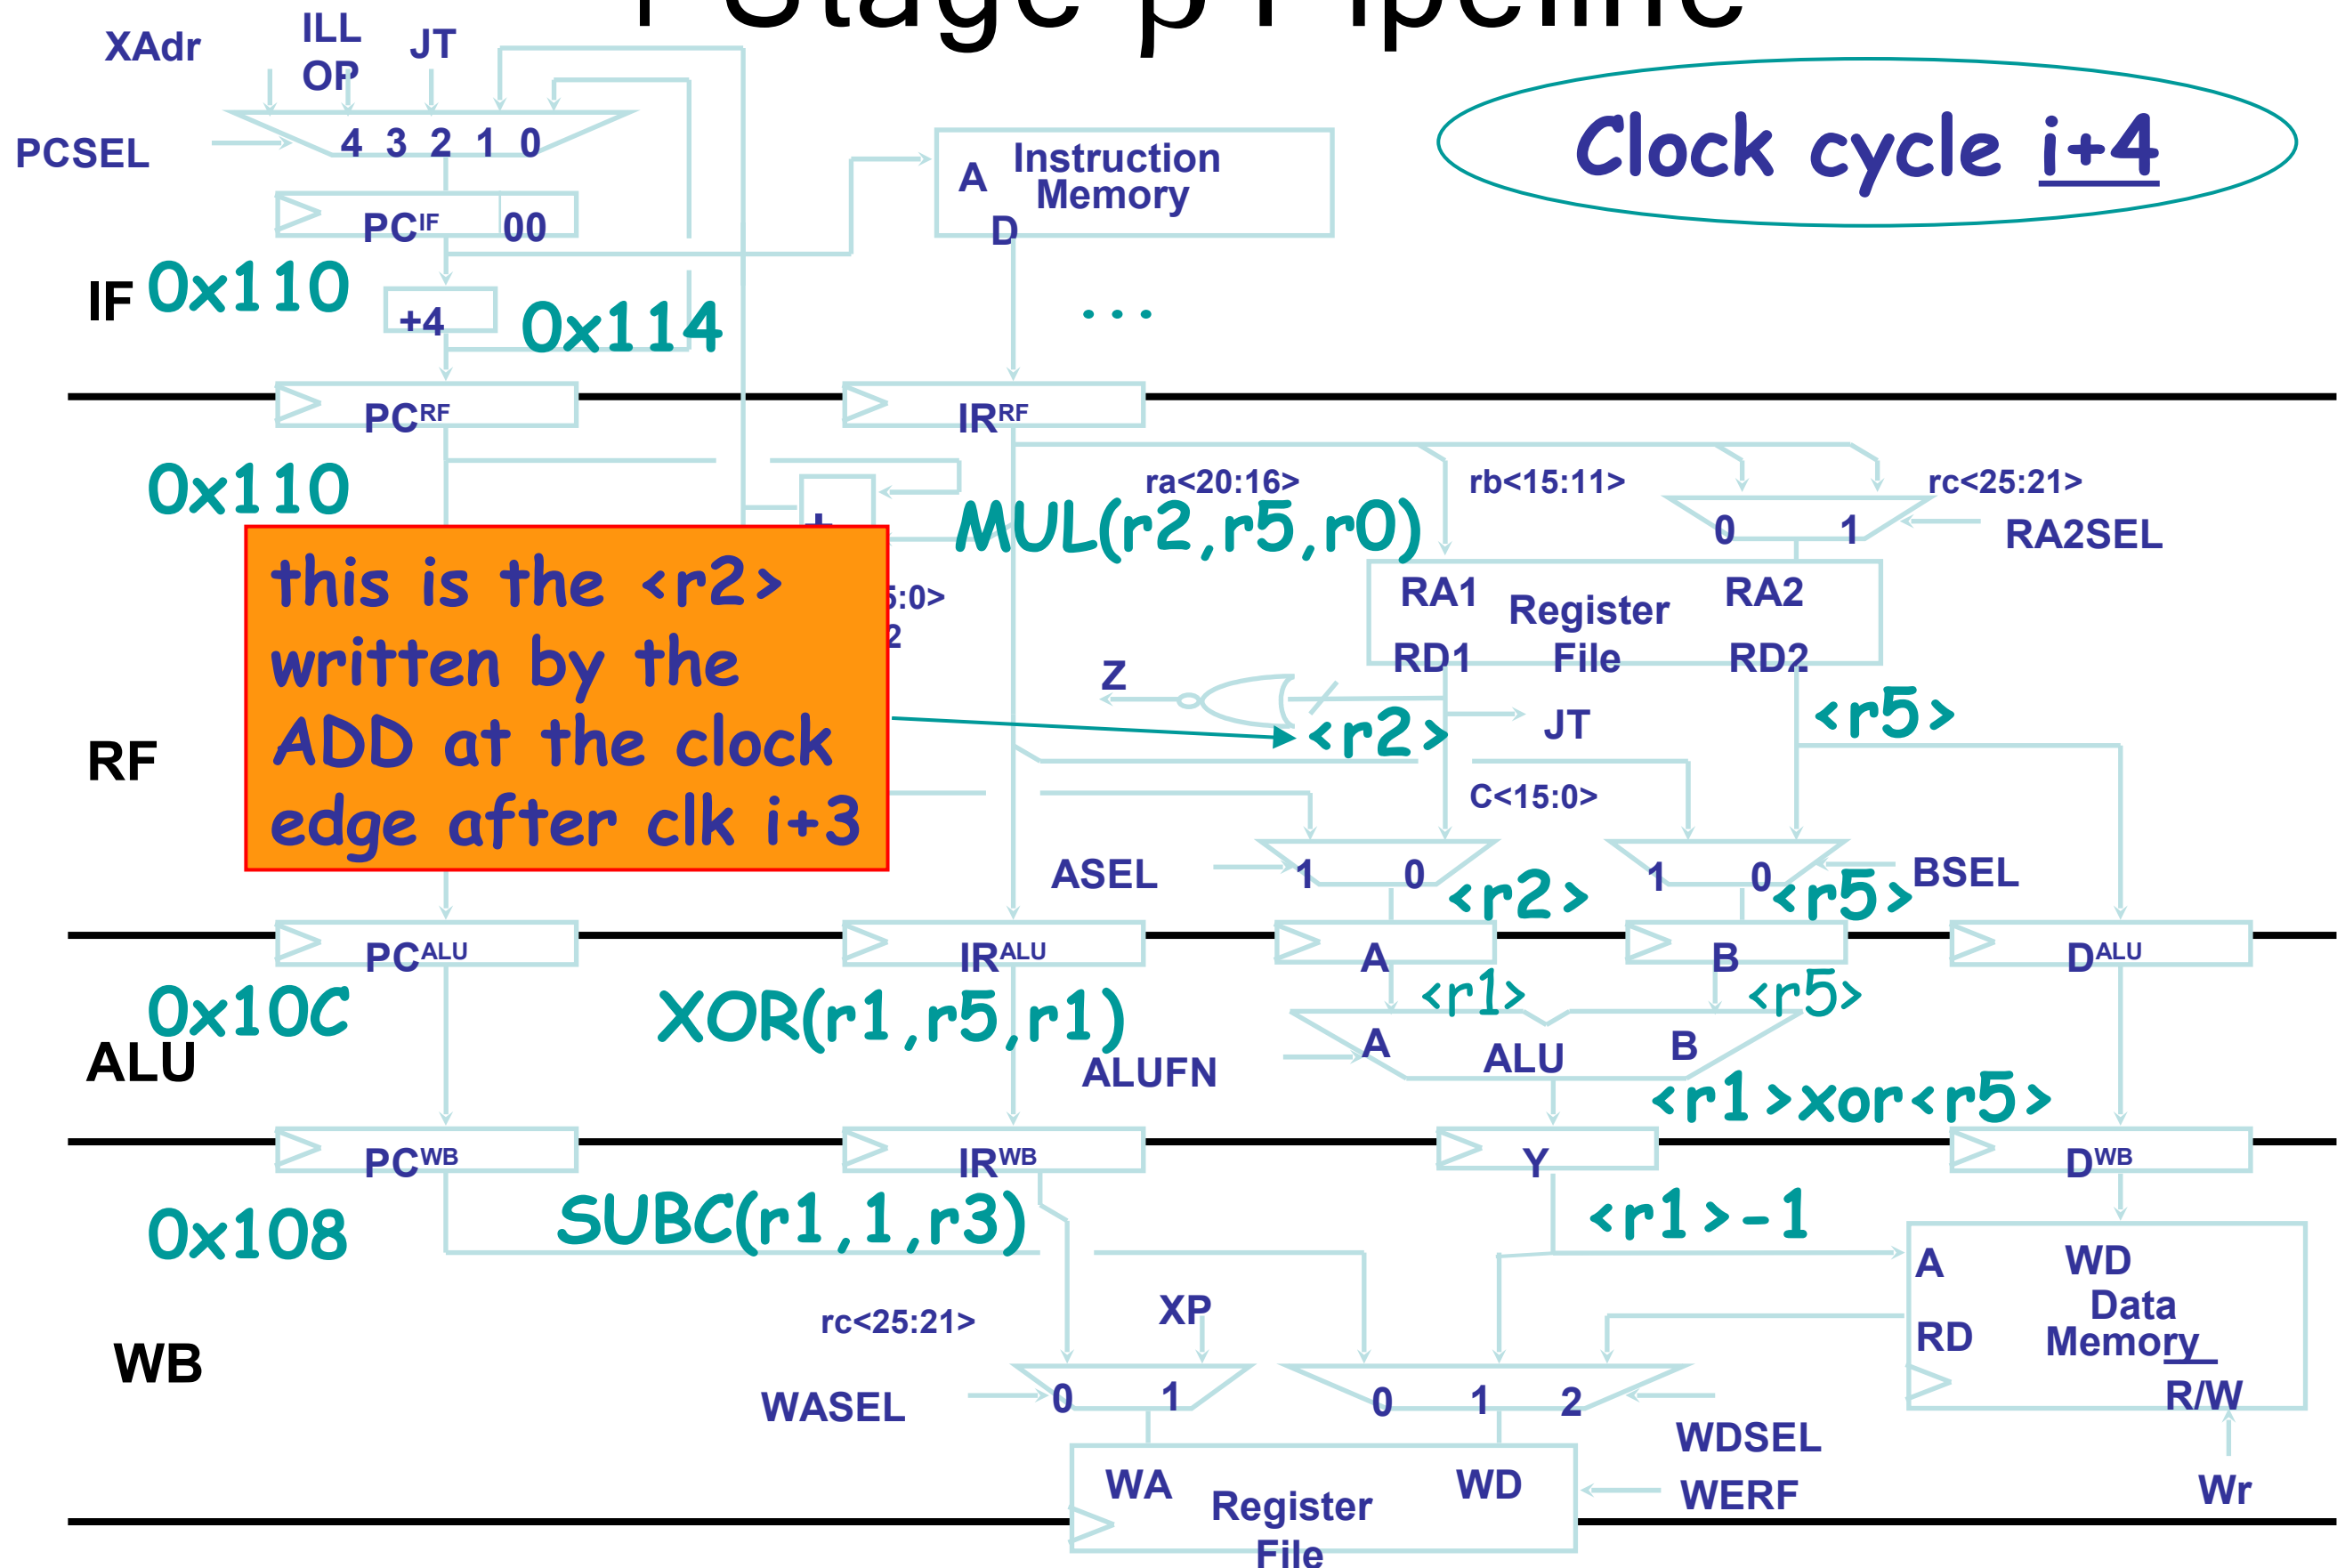


DISCS

4-Stage β Pipeline



4-Stage β Pipeline



Register Write-Back is Delayed

suppose $r1=7, r2=9, r3=10, r5=3$

► Not only for pipeline stages, but for register and wire values as well!

► Looks OK for now, since MUL reads updated R2 value at cycle $i+4$.

0x100: ADDC(r1, 1, r2)

0x104: SUBC(r1, 1, r3)

0x108: XOR(r1, r5, r1)

0x10C: MUL(r2, r5, r0)


TIME (cycles) →

Space ↓

| | t | t+1 | t+2 | t+3 | t+4 | t+5 | t+6 |
|------------------|-------|-------|-------|-------|-------|-----|-----|
| IF | ADDC | SUBC | XOR | MUL | ... | | |
| RF | | ADDC | SUBC | XOR | MUL | ... | |
| ALU | | | ADDC | SUBC | XOR | MUL | ... |
| WB | | | | ADDC | SUBC | XOR | MUL |
| PC _{IF} | 0x100 | 0x104 | 0x108 | 0x10C | 0x110 | ... | |
| R1 | 7 | 7 | 7 | 7 | 7 | 7 | |
| R2 | 9 | 9 | 9 | 9 | 8 | 8 | |
| R3 | 10 | 10 | 10 | 10 | 10 | 6 | |

Diagram illustrating the pipeline stages (IF, RF, ALU, WB) and register values (R1, R2, R3) over time (cycles t to t+6). The instructions are ADDC, SUBC, XOR, and MUL. The register values are shown in the bottom row. A vertical line marks the start of cycle t+4, where the MUL instruction reads the updated R2 value (8) and R3 value (10).

00101010010101000111101010010101
10001100100001111001101010010101
110010101010101000010011001010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010100101010010101



DISCS

What if MUL was Earlier?

- MUL gets 9.
- A stale (and incorrect) value!


0x100: ADDC(r1, 1, r2)
0x104: SUBC(r1, 1, r3)
0x108: MUL(r2, r5, r0)
0x10C: XOR(r1, r5, r1)

TIME (cycles) →

Space ↓

| | t | t+ 1 | t+ 2 | t+ 3 | t+ 4 | t+ 5 | t+ 6 |
|------------------|-------|-------|-------|-------|-------|------|------|
| IF | ADDC | SUBC | MUL | XOR | ... | | |
| RF | | ADDC | SUBC | MUL | XOR | ... | |
| ALU | | | ADDC | SUBC | MUL | XOR | ... |
| WB | | | | ADDC | SUBC | MUL | XOR |
| PC _{IF} | 0x100 | 0x104 | 0x108 | 0x10C | 0x110 | ... | |
| R1 | 7 | 7 | 7 | 7 | 7 | 7 | |
| R2 | 9 | 9 | 9 | 9 | 8 | 8 | |
| R3 | 10 | 10 | 10 | 10 | 10 | 6 | |

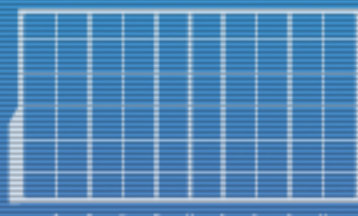
00101010010101000111101010010101
10001100100001111001101010010101
110010101010101000010011001010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010100101010010101



Write-Back Delay Slots

- ▶ **Problem:** Registers are not written back right away, so if later instructions try to read too early, they get “stale” values.
- ▶ **Possible solution:**
- ▶ **Insert NOPs**
 - ▶ `ADD(R31,R31,R31)` | does absolutely nothing!
- ▶ **Reorder Code**
 - ▶ Move non-dependent code into delay slot.

```
00101010010101000011110100001100
10001100100001111001101010010101
110010101010101000010011001010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010101010101010101010101010101
```



Inserting NOPs

► How Many?

0x100: ADDC(r1, 1, r2)

0x104: SUBC(r1, 1, r3)

0x108: MUL(r2, r5, r0)


0x10C: XOR(r1, r5, r1)

TIME (cycles) →

Space ↓

| | t | t+ 1 | t+ 2 | t+ 3 | t+ 4 | t+ 5 | t+ 6 |
|------------------|-------|-------|-------|-------|-------|------|------|
| IF | ADDC | SUBC | | | | | |
| RF | | ADDC | SUBC | | | | |
| ALU | | | ADDC | SUBC | | | |
| WB | | | | ADDC | SUBC | | |
| PC _{IF} | 0x100 | 0x104 | 0x108 | 0x10C | 0x110 | ... | |
| R1 | 7 | 7 | 7 | 7 | 7 | 7 | |
| R2 | 9 | 9 | 9 | 9 | 8 | 8 | |
| R3 | 10 | 10 | 10 | 10 | 10 | 6 | |

00101010010101000011110101000000
10001100100001111001101010010101
11001010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010100101010010101

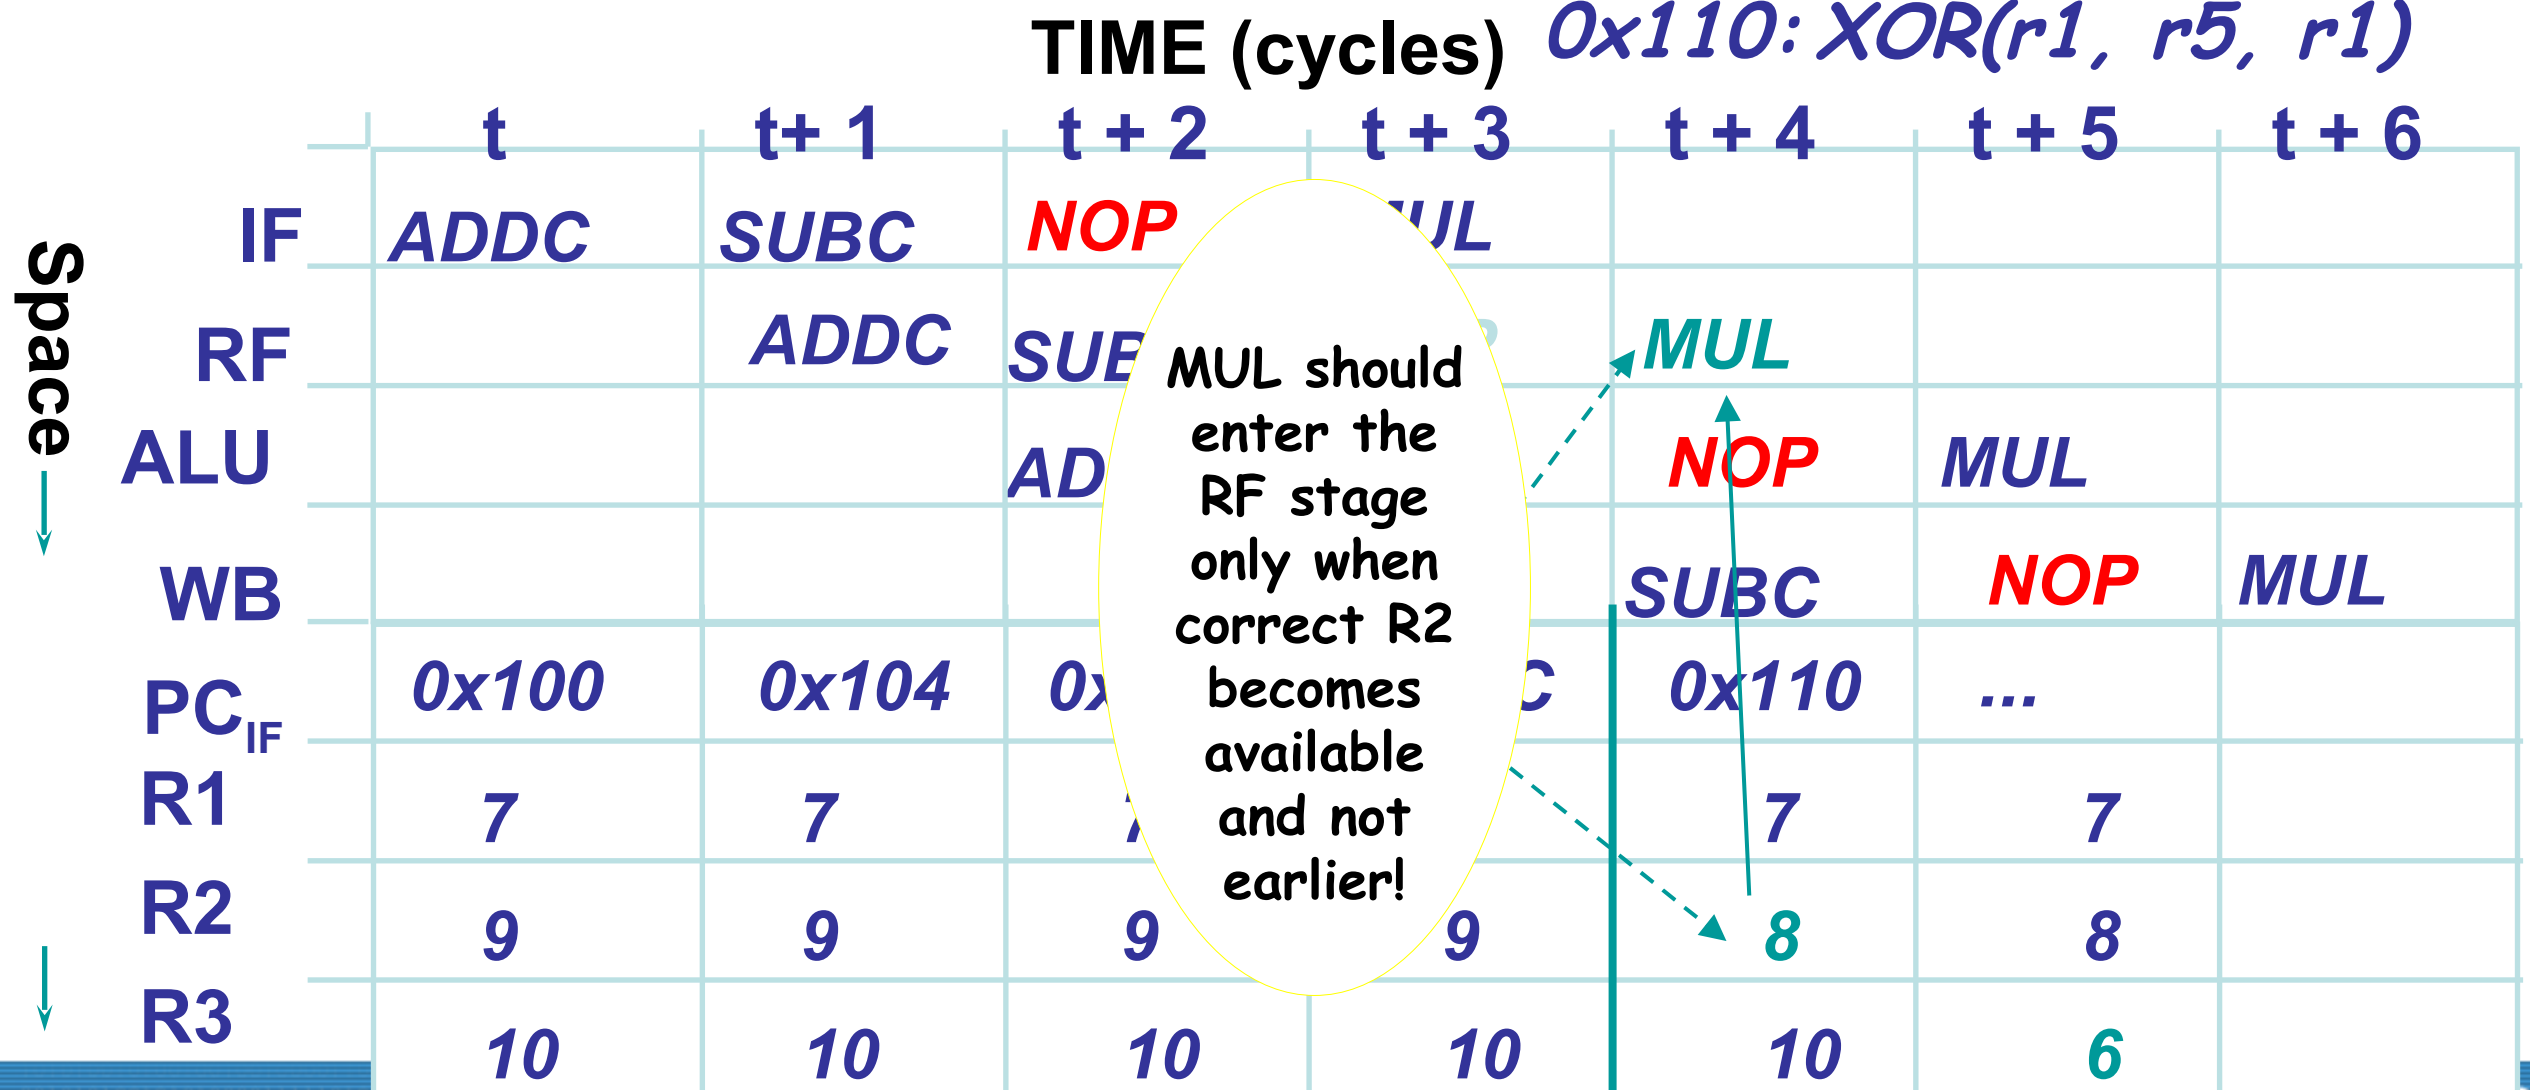


DISCS

Inserting NOPs

► Need to put reading instruction (MUL) in RF stage in the same cycle that correct values of regs become available (must be at least the 3rd instruction after ADDC, which writes to R2.)


0x100: ADDC(r1, 1, r2)
 0x104: SUBC(r1, 1, r3)
 0x108: **NOP()**
 0x10C: MUL(r2, r5, r0)
 0x110: XOR(r1, r5, r1)



Space
↓

↓

00101010010101000011110101010101
 10001100100001111001101010010101
 110010101010101000010011001010100
 100101001001001010101010101010101
 11100001111010110000000111101001
 001001010100101001001010010010110
 10010100100001010100100101001010
 10010100101010010100101010010101
 10010100101010010100101010010101

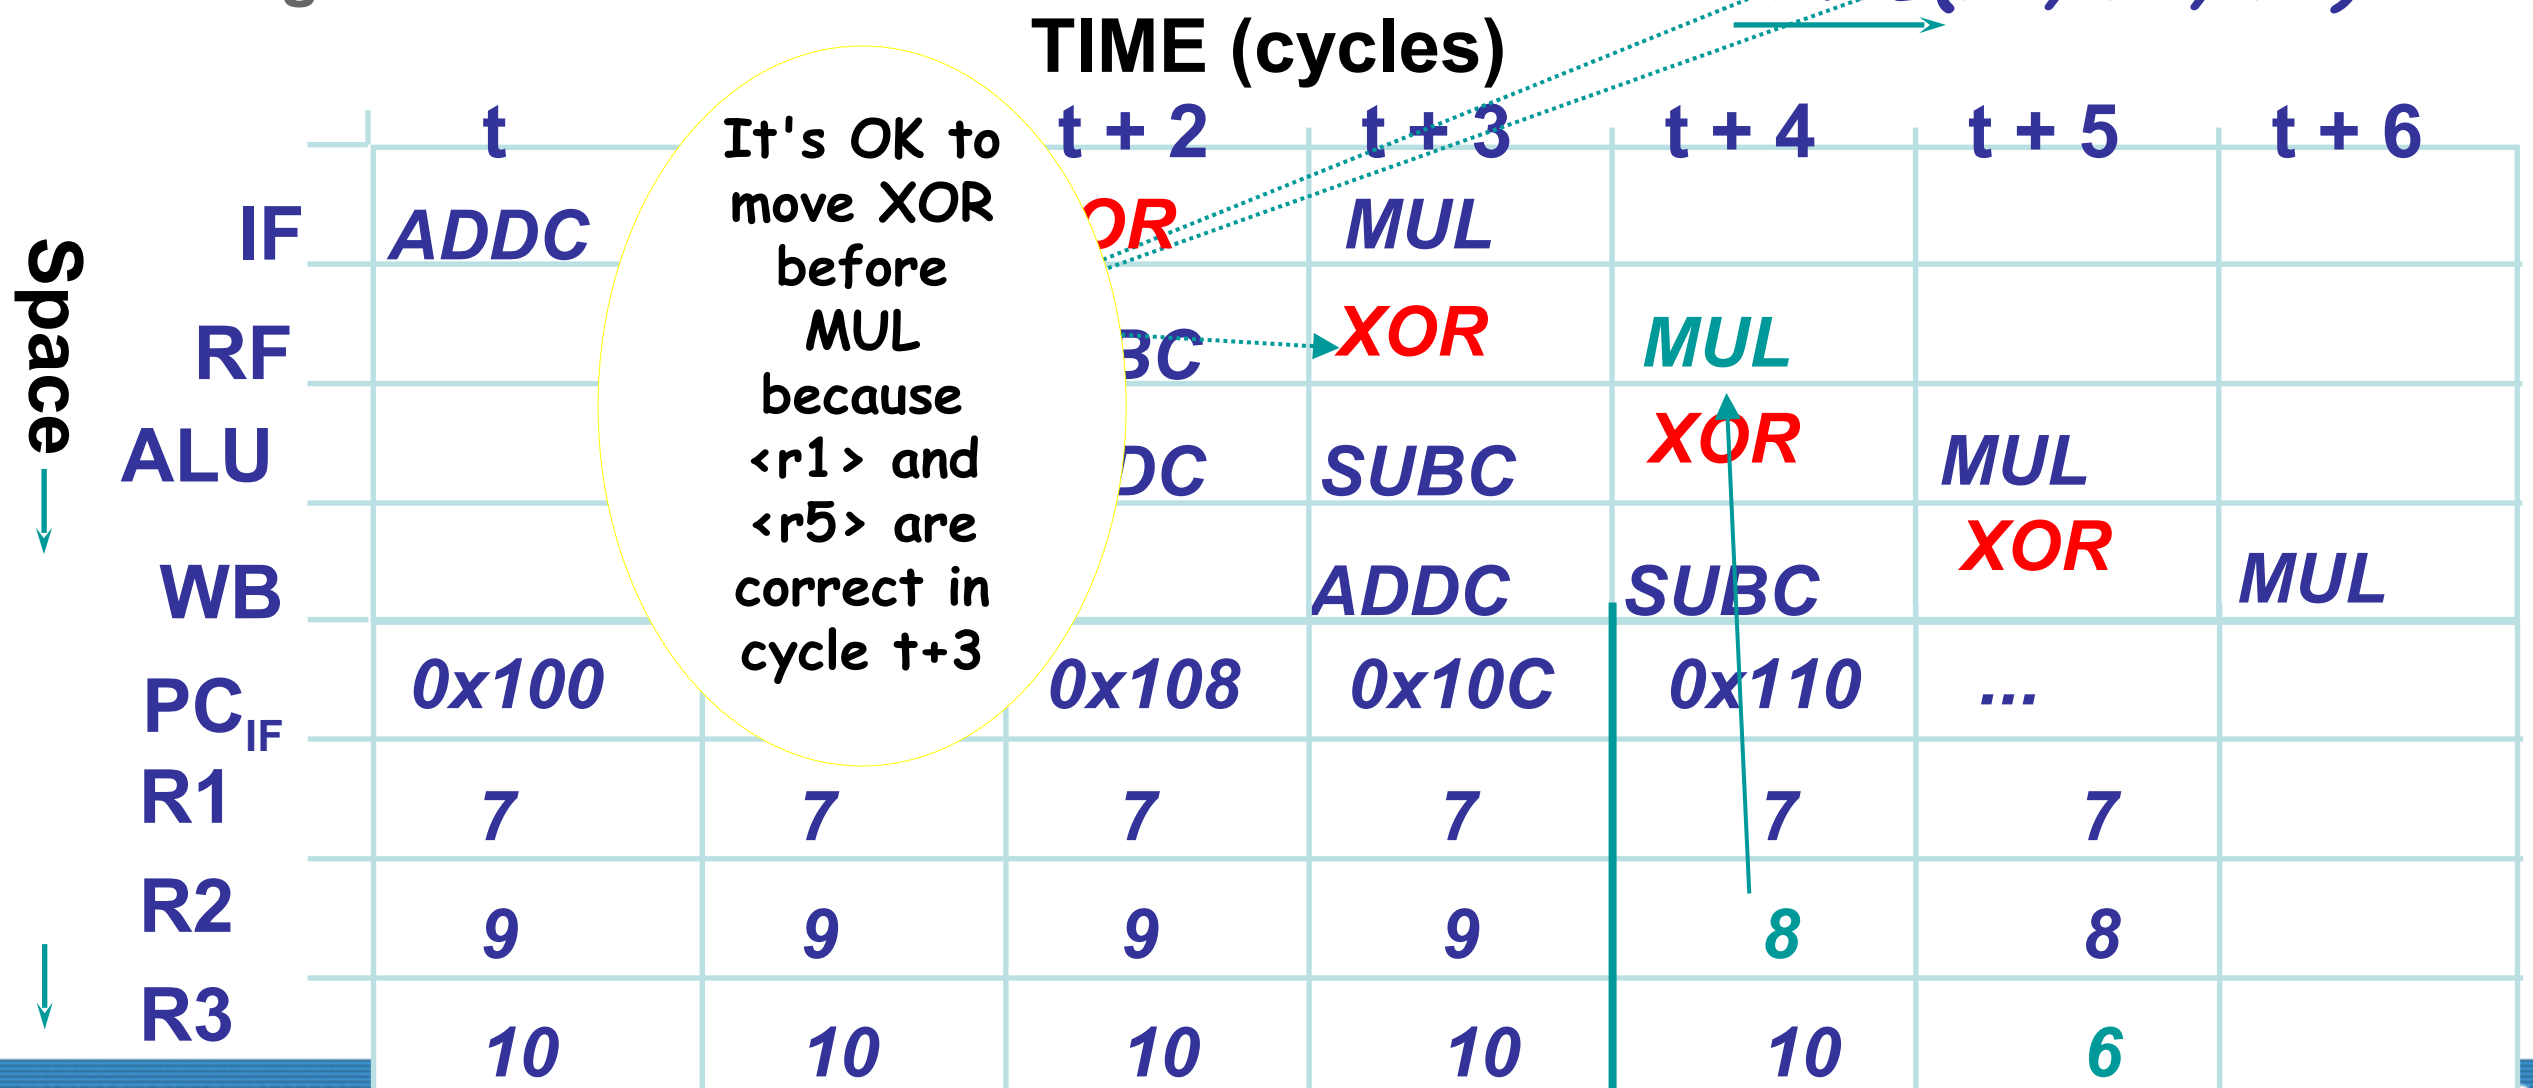


DISCS


Reordering Code

- NOPs waste cycles.
- Instead of NOPs, we can use other instructions.
- But they must not depend on changing regs!

0x100: ADDC(r1, 1, r2)
0x104: SUBC(r1, 1, r3)
0x108: XOR(r1, r5, r1)
0x10C: MUL(r2, r5, r0)



00101010010101000111101010101010
 10001100100001111001101010010101
 110010101010101000010011001010100
 10010100100100101010101010101010
 11100001111010110000000111101001
 001001010100101001001010010010110
 10010100100001010100100101001010
 10010100101010010100101010010101
 10010101010101010101010101010101



Write-Back Delays

- Consider the sequence:

ADD(r1, r2, r3)

CMPLEC(r3, 5, r0)

MUL(r1, r2, r4)

SUB(r1, r2, r5)

| | i | i + 1 | i + 2 | i + 3 | i + 4 | i + 5 | i + 6 |
|-----|-----|-------|-------|-------|-------|-------|-------|
| IF | ADD | CMP | MUL | SUB | | | |
| RF | | ADD | CMP | MUL | SUB | | |
| ALU | | | ADD | CMP | MUL | SUB | |
| WB | | | | ADD | CMP | MUL | SUB |


ADD writes new value in *r3* during _____ cycle,
which is available beginning of _____ cycle.

Value of *r3* read by *CMP* during _____ cycle.

CMP reads _____ value of *r3*.

Fix the
problem:
- insert NOPs
- move code

00101010010101000011110100001100
10001100100001111001101010010101
11001010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
00100101010010100100100100100110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010100101010010101



DISCS

What About Branching/Jumping?

- Consider a different sequence:

LOOP: CMPLC(r3, 100, r0)

ADD(r1, r2, r3)

SUB(r1, r2, r4)

BNE(r0, LOOP)

XOR(r9, r6, r3)


| | i | i + 1 | i + 2 | i + 3 | i + 4 | i + 5 | i + 6 |
|-----|------------|------------|------------|------------|------------|------------|------------|
| IF | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>?</i> | | |
| RF | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | | |
| ALU | | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | |
| WB | | | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> |

What's wrong?

Think: What determines whether to branch or not?


(For JMP: How do we know where to jump?)

0010101001010100011110100001100
10001100100001111001101010010101
110010101010101000010011001010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
1001010010101010101010101010101




DISCS

Branch Execution



| | i | i + 1 | i + 2 | i + 3 | i + 4 | i + 5 | i + 6 |
|------------------|------------|------------|------------|------------|------------|------------|------------|
| IF | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>XOR</i> | <i>CMP</i> | |
| RF | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>XOR</i> | <i>CMP</i> |
| ALU | | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>XOR</i> |
| WB | | | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> |
| PC _{IF} | 0x100 | 0x104 | 0x108 | 0x10C | 0x110 | 0x100 | |
| Z | | | | | 0 | 1 | |
| PCSEL | | | | | | | |



Branch Delay Slots


- ▶ Problem: One (or more) instructions have been fetched by the time a branch can be decided.
- ▶ Software solution: Insert NOPs and/or reorder code!

| | i | i + 1 | i + 2 | i + 3 | i + 4 | i + 5 | i + 6 |
|------------------|------------|------------|------------|------------|------------|------------|------------|
| IF | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>XOR</i> | <i>CMP</i> | |
| RF | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>XOR</i> | <i>CMP</i> |
| ALU | | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>XOR</i> |
| WB | | | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> |
| PC _{IF} | 0x100 | 0x104 | 0x108 | 0x10C | 0x110 | 0x100 | |
| Z | | | | | 0 | | |
| PCSEL | | | | | 1 | | |

```

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010101010101010101010101010101

```



DISCS

Inserting NOPs

LOOP: CMPLC(r3, 100, r0)

ADD(r1, r2, r3)

SUB(r1, r2, r4)

BNE(r0, LOOP)

NOP()

XOR(r9, r6, r3)

"branch
delay
slot"

► Add NOPs up to and including cycle where correct *target PC* is decided.

► What about reordering?

| | i | i + 1 | i + 2 | i + 3 | i + 4 | i + 5 | i + 6 |
|------------------|--------------|--------------|--------------|--------------|--------------|--------------|------------|
| IF | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>NOP</i> | <i>CMP</i> | |
| RF | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>NOP</i> | <i>CMP</i> |
| ALU | | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> | <i>NOP</i> |
| WB | | | | <i>CMP</i> | <i>ADD</i> | <i>SUB</i> | <i>BNE</i> |
| PC _{IF} | <i>0x100</i> | <i>0x104</i> | <i>0x108</i> | <i>0x10C</i> | <i>0x110</i> | <i>0x100</i> | |
| Z | | | | | 0 | 1 | |
| PCSEL | | | | | | | |

