



DEPARTMENT OF INFORMATION SYSTEMS AND COMPUTER SCIENCE



```
00101110101000111010111100100111010101101001000101
1101010110101010000101010101001010101010101010
1010010100100100101010101010101010101010101010101010
1110000111101011000000011110101010101010000010101
11101010111100101000100101111010100010100100111010
1010100101001001001000010101011010101010101010010111
0010101001010100101010000001010101001111101000011001
1000110010000111100110101011000100110101010000101010
1100101010101000010011001010100010010101010101010
1010010100100100101010101010101010101010101010101010
1110000111101011000000011110101010101010000010101
0010010101001010010010100100010101010101001010010
1001010010000101010010010101001010010101010010010
1001010010101001010010101001010010101001001001001
100101010101001010101010100101010101010010101010
```

										01
										02
										03
										04
										05
A	B	C	D	E	F	G	H			

Function Calls

Is Infinite Recursion Possible?

Lecture Time!

- Implementing Functions: Macros
- Implementing Functions: Subroutines with Static Allocation
- Implementing Functions: Using a Stack

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010101010101



DISCS

Writing Functions

- How do we implement a reusable function in Beta assembly?

```
int fact( int n )  
{  
    int f = 1;  
    while( n > 0 )  
    {  
        f *= n;  
        n -= 1;  
    }  
    return f;  
}
```

00101010010101000011110100001100
10001100100001111001101010010101
11001010101010100001001100101010100
100101001001001010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010101010101010010101



Writing Functions

- How do we implement a reusable function in Beta assembly?

```
int fact( int n ) // R1 contains n
{
    int f = 1;      // ADDC( R31, 1, R0 )
    while( n > 0 ) // loop: CMPLT( R31, R1, R2 )
    {                // BF( R2, done, R31 )
        f *= n;      // MUL( R0, R1, R0 )
        n -= 1;      // SUBC( R1, 1, R1 )
    }                // BR( loop )
    return f;        // done: R0 contains f
}
```

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
1001010010100101010101001010101



Using Macros (aka Inlining)

- Specify register mapping: input, output, and temp regs.
 - Input and output regs can be parameters of .macro.
 - *Temporary registers must be documented*, so user knows which regs are “clobbered”.
- Macros are expanded at compile-time.
 - Labels CANNOT be used within .macro in BSim.
 - This results in duplicate names.
 - Need to use relative addressing for branching.
 - Note that it is still converted to word displacement in the final 4-byte representation.

```
00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010101010101
```



Using Macros (aka Inlining)

```
.include beta.uasm                                | duh

macro fact( Ra, Rc ) {                             | Rc = fact(<Ra>)
ADDC( R31, 1, Rc )                                 | Rc=1
MUL( Rc, Ra, Rc )                                  | loop:
SUBC( Ra, 1, Ra )
BT( Ra, -8, R31 )                                  | branch to loop
}                                                    | note: not error-free

| how to use the macro
main:                                                | int main( void )
CMOVE( 6, R1 )                                     | ADDC( R31, 6, R1 )
CMOVE( 3, R2 )
fact( R1, R0 )                                     | R0=fact(6)
fact( R2, R3 )                                     | R3=fact(3)
HALT()
```

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
100101001010010101010010010101



DISCS

Using Macros (aka Inlining)

- Advantages
 - Simple and straightforward
 - Faster than other methods: no overhead
- Disadvantages
 - Compiled binary code grows the more you call macros
 - Did you try running the sample code? What did you notice?
 - Need to be careful about “clobbered” registers
 - What if the macro needed to use a register not included in its arguments for temporary storage?
 - Hard to debug
 - No recursion
- In general, inlining is good for *short* code.

– Often used by C/C++ compilers for optimizing speed.

```
00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010010101010
```



Subroutines with Static Allocation

- Only one copy of function code.
- Linkage Pointer (LP=R28) holds return address:
 - From main program, BR(fact,LP) to call.
 - From subroutine, JMP(LP) to return to where the function was called.
- You still need to specify register mapping: input, output, and temp regs.
 - *Input and output regs are now fixed and must be documented in addition to the temporary regs.*

```
00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010101010101
```



Subroutines with Static Allocation

main:

CMOVE (6, R1)

CMOVE (3, R2)

BR (fact, LP) | call fact (6)

| at this point, R0 contains fact (R1)

MOVE (R2, R1)

BR (fact, LP) | call fact (3)

MOVE (R0, R3)

| at this point, R3 contains fact (R2)

HALT ()

fact: | takes input in R1, puts output in R0

ADDC (R31, 1, R0)

loop: MUL (R0, R1, R0)

SUBC (R1, 1, R1)

BT (R1, loop, R31) | note: not error-free

JMP (LP) | return



DISCS

Subroutines with Static Allocation

- Alternatively, we can put inputs and outputs in main memory instead of registers.
 - Useful when we use more than 32 variables.
 - Or if we don't want to clobber too many registers.
- Just use a few registers, then LD and ST to them as necessary.
 - Let's make code for that now! (and fix the BT issue also)
 - Hint: Documentation should say something like:
“takes input in factN, puts output in factF, clobbers R0 and R1”.

```
00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010101010101
```



Subroutines with Static Allocation

- Advantages:
 - Single copy of code, which saves space and makes it easier to debug.
 - Can use memory for storage.
 - Macros can too, but creates many copies and wastes enough space already.
- Disadvantages:
 - Still clobbers registers.
 - No nested function calls.
 - Still no recursion.
- But we're getting somewhere!

```
00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
1001010010101001010101001010101
```



DISCS


Analyzing Recursive Functions

- Consider the following code:

```
int fact( int n )  
{  
    if( n > 0 )  
        return fact( n-1 ) * n;  
    return 1;  
}
```

- What are the storage requirements?
 - Why can't we use recursion or even nested function calls for subroutines with static allocation?
 - Is there an issue with local variable n? What about the return value?

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010010101010



The Solution: Stacks

- Needs a special pointer, the Stack Pointer (R29).

SP=R29

```
.macro PUSH(rx) {  
  ADDC(SP, 4, SP)  
  ST(rx, -4, SP)  
}
```

```
.macro POP(rx) {  
  LD(SP, -4, rx)  
  SUBC(SP, 4, SP)  
}
```

| assumes SP initially contains a high value

| (is there a problem if it contains 0?)

| and these macros are already in beta.uasm!



DISCS

Possible Issue With Pop?

- Push a few values on to the stack, then pop one.
 - What do the macros actually do?
 - You may want to try it out in BSim.
 - Look at the stack closely. Is there a possible issue?
 - More importantly, will it be an issue?
 - Hint: When formatting your hard drive, what is the difference between quick format and full format?
 - Hint #2: While a Java program is still running, what does garbage collection do?

```
00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010101010101
```



Stack State

- Given a stack's current state:
 - X pushes followed by X pops (assuming # of pops NEVER exceeds # of pushes at any given time) should return the stack to that state.
- Example:
 - Stack contains 2 elements B and T. After 3 pushes, 2 pops, 4 pushes, 1 pop, 1 push, and 5 pops, what does the stack contain?

```
00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
1001010010101001010101010010101
```



Stack Frames

- Storage requirements of each function call is placed in a *stack frame* which contains at least the following:
 - Return address (old LP)
 - Arguments (example: x, y, z in f(x, y, z))
 - Temp/Local variables
 - Original values of clobbered registers
- The stack is used for this storage, of course!
 - Grow (PUSH) stack for each function call.
 - Shrink (POP) stack on return.
 - Each call gets its own “stack frame”.

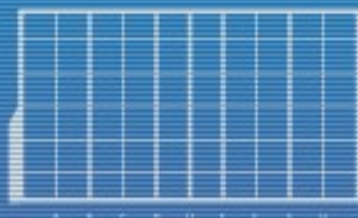
```
00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010101010101
```



Stack State + Stack Frames

- *At the end of a function call, the state of the stack is the same as right before that function call.*
 - In other words, SP should be in the same position before and after the function call.
 - Example using the recursive version of fact()
(go back a few slides):
 - From main, call fact(2) (push stack frame A)
 - From fact(2), call fact(1) (push stack frame B)
 - From fact(1), call fact(0) (push stack frame C)
 - Return fact(0) (pop stack frame C)
 - Return fact(1) (pop stack frame B)
 - Return fact(2) (pop stack frame A)


00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010101010101



So Many Pointers!

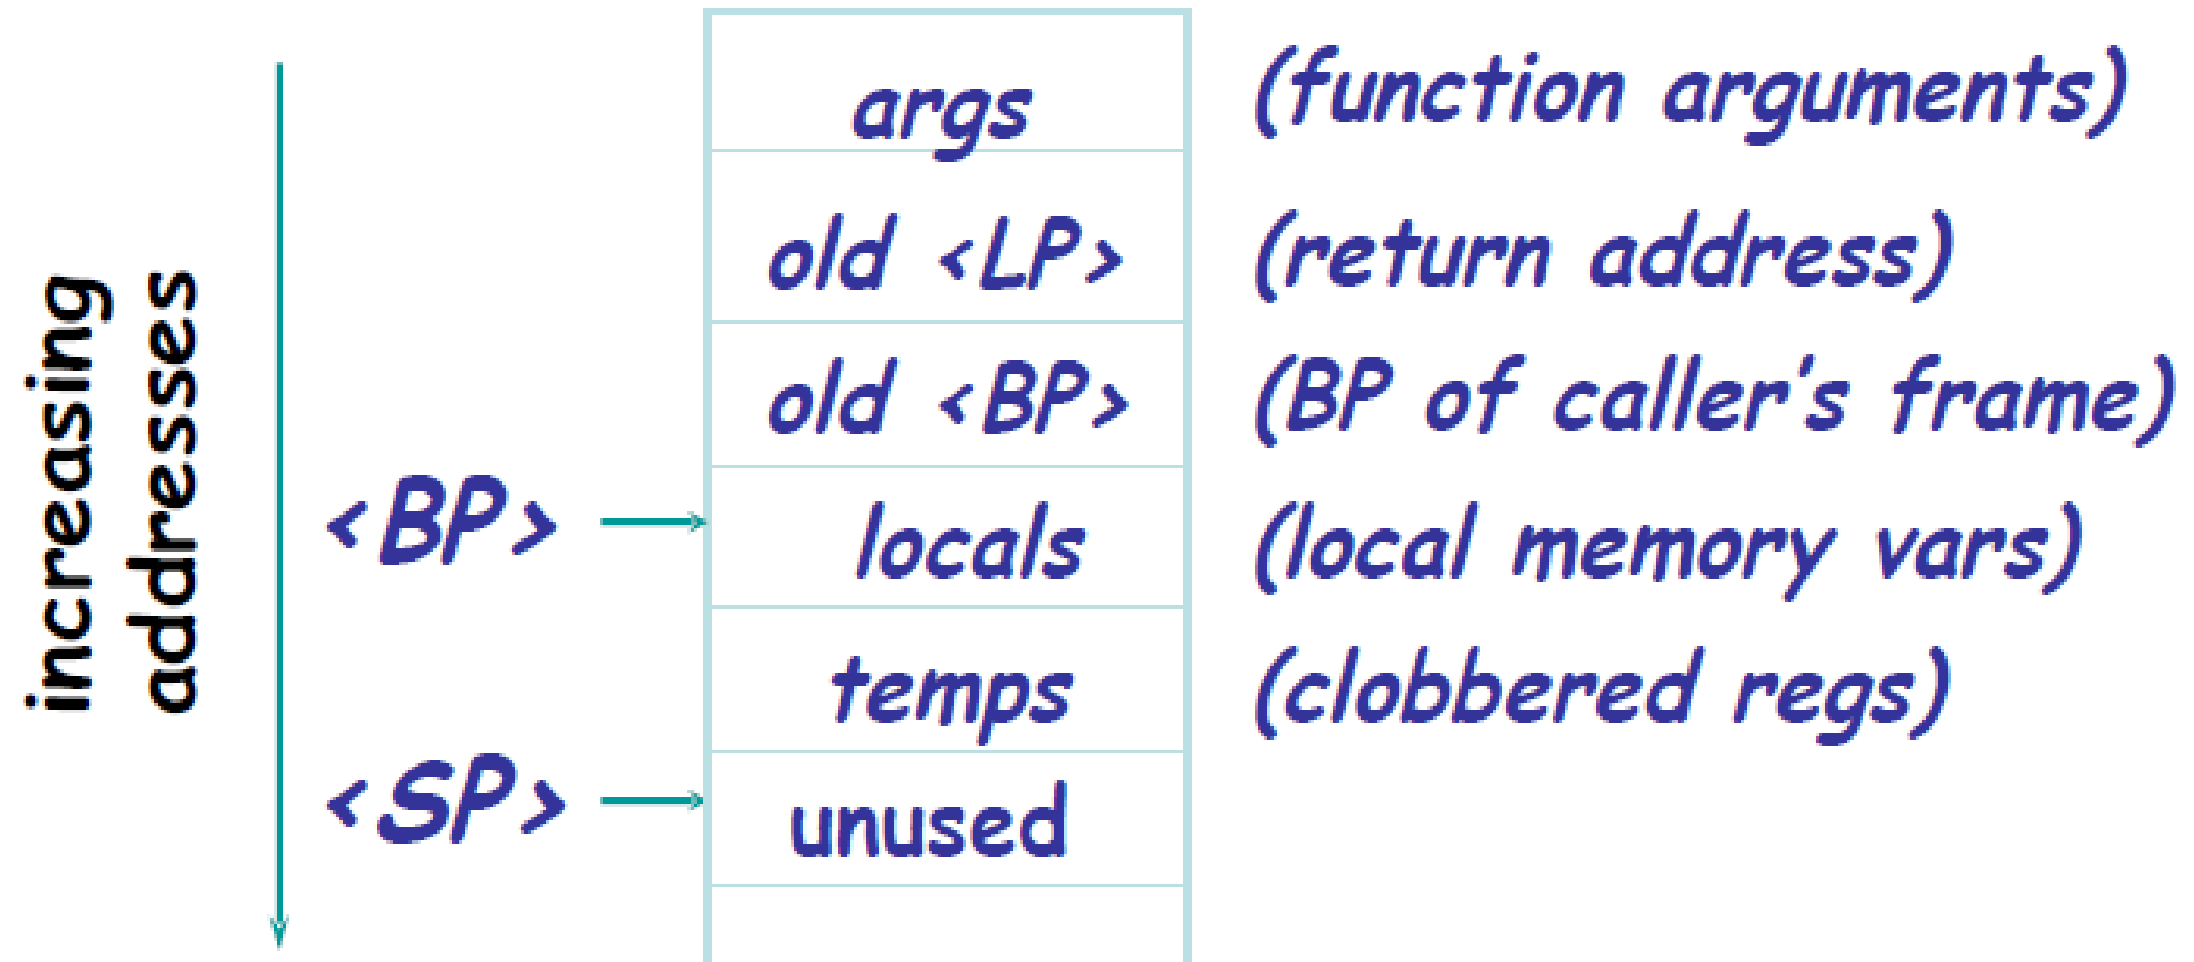
- BP (R27) – Base Pointer, contains the address of the first local variable in a function's stack frame.
- LP (R28) – Linkage Pointer, contains the address of the line to return to when the function is finished.
- SP (R29) – Stack Pointer, contains the address of the line above the top of the stack (or points at the first unused location, if you prefer it that way).
- XP (R30) – Exception Pointer, not yet discussed. Therefore, not yet important.

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
1001010010101010101010101010101




DISCS

Stack Frame Details



00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
1001010010100101010101001010101

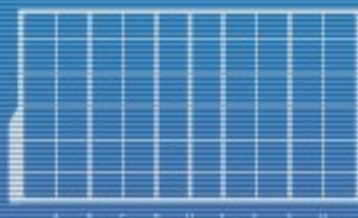


DISCS

Stack Frame Details

- Per stack frame, for BSim:
 - Function arguments appear in reverse order: the LAST arg has the LOWEST address while the FIRST arg has the HIGHEST address.
 - The LP in the frame is the function's LP that will be used for the JMP(LP) at the end.
 - The BP in the frame is NOT the function's BP. It is the BP of whatever called this function.
 - To minimize number of clobbered registers, stack space can be used for local variables.
 - Temps are the original contents of the registers that will be clobbered in this function, and are to be restored before JMP(LP).

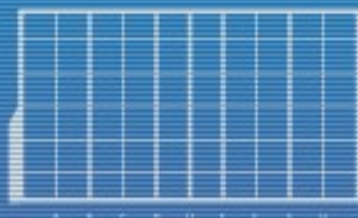
00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010101010101



The Stack Discipline

- Caller's part :
 - Push args in stack in reverse order.
 - BR(callee, LP)
 - Should go back to next line (after branch instruction) after callee's JMP(LP)...
 - ... where it should now remove args from stack (can just subtract $4 * \text{NumberOfArgs}$ from SP).

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
10010100101010010101010101010101



DISCS

The Stack Discipline

- Callee's part:
 - Save LP and BP on stack, then set BP to SP (so BP will point at first local, if any).
 - PUSH(LP) and PUSH(BP), then MOVE(SP, BP).
 - Allocate mem for local vars on stack so each call has its own memory.
 - Perform computation and leave final result in R0.
 - Save and Restore clobbered registers.
 - Restore stack to original state before entry.
 - JMP(LP)

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
1001010010101001010101001010101



DISCS

Accessing the Stack Frame

- BP is base for accessing the current stack frame.
- Within one function call:
 - BP points after args, old LP, and old BP in stack frame.
 - $\langle \text{BP} \rangle$ points to local var #1, $\langle \text{BP} \rangle - 12$ points to arg#1.
- Notice that caller must push arguments in reverse order so arg #1 is always in the same place!
- So, given the function's BP and a register Rx:
 - How do we access argument #N?
 - How do we access local variable #N?
 - Note: For both, we have to be able to read from and write to the variable (hint: memory access).

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
1001010010101001010101001010101



Exercises

- Implement the recursive factorial function in BSim.
- Create a function that averages two integers. Then, use it to get the average of three integers.
 - `average(x, average(y, z));`
- Create a function that returns the square of an integer. Then, use it to replace each element in an array with its square.

00101010010101000011110100001100
10001100100001111001101010010101
110010101010100001001100101010100
1001010010010010101010101010101
11100001111010110000000111101001
001001010100101001001010010010110
10010100100001010100100101001010
10010100101010010100101010010101
1001010010101001010101010010101



DISCS